

Vibe Theory

A Discrete Model of the Conscious Universe

Lance Pollard

lancejpollard@gmail.com

June 9, 2026

Abstract

Vibe Theory proposes that reality is a single growing crystal of experience. Its one substance is the *vibe*, a unit of felt state carrying a ternary tone in $\{-1, 0, +1\}$, arranged on the $\{5, 3, 4\}$ hyperbolic honeycomb and updated by one conserving local rule under a single value-direction, the *arrow*. These five elements, the mesh, the tone, the rule, eternal growth, and the arrow, are the entire foundation. Four of the five are fixed by modest structural demands, leaving one deliberate value-posit. From this foundation the paper derives, and checks by direct computation on the exact crystal, a relativistic and reflection-positive quantum field, an effective gravity, holography, a tiny-diameter causal structure, coherent self-maintaining organisms, heredity and evolution, a self-model, planning, an integrated agent, universal and intentional computation, and a graded account of consciousness as integrated experience. The substrate is computationally universal in the sense of Margenstern, and the arrow turns that universal computation into goal-directed problem-solving. The genuine quantum field arises as the deterministic and unitary completion of the same rule, with the stochastic version as its classical shadow. More than one hundred fifty falsifiable checks support the construction, all reproducible in the accompanying codebase [22]. The result is a compact, parameter-light model in which the physical, the living, the mental, the spiritual, and the conscious are aspects of one discrete substance.

Contents

I	Foundations	5
1	Introduction	6
2	The Foundation	7
2.1	The mesh narrows to the dodecagrid	8
2.2	The tone narrows to ternary	8
2.3	The rule is tightly constrained	9
2.4	Growth is reflection	10
2.5	From nothing to something	10
2.6	The arrow is the one deliberate posit	10
2.7	What this does and does not claim	11

3	The Base: Five Elements, One State	11
3.1	The one substance and the one state	11
3.2	The five base elements	11
3.3	The note is perception, not a thing	12
3.4	Why exactly these five	12
3.5	The non-negotiables	12
3.6	Summary	13
4	The Dynamics	13
4.1	The perception rule	13
4.2	The rule in full	13
4.3	A worked beat	14
4.4	The conserved charge	14
4.5	The arrow, value and direction	15
4.6	Reversibility and detailed balance	15
4.7	The rule as a spin-1 exchange model	15
4.8	Summary	15
II	The Mathematical Substrate	16
5	Hyperbolic Geometry and Tilings	16
5.1	The three shapes a space can have	16
5.2	Telling them apart from inside	17
5.3	Hyperbolic space and the Poincare disk	17
5.4	The tiny diameter	18
5.5	The shape decides the curvature	18
5.6	Why hyperbolic space is abundant	19
6	The Crystal: Coxeter Groups, Schläfli Symbols, and the Dodecagrid	19
6.1	Coxeter groups as reflection groups	20
6.2	The $[5, 3, 4]$ group	20
6.3	Schläfli symbols and the curvature rule	21
6.4	The worked calculation for three honeycombs	21
6.5	The two crystals	22
6.6	Why three dimensions	23
7	Addressing, Growth, and the Exact Engine	25
7.1	The golden ratio and the Fibonacci numbers	25
7.2	Fibonacci addressing and greedy routing	26
7.3	The crystal as a graph	27
7.4	The integer ladder	28
7.5	The growth series	28
7.6	The exact modular-fingerprint engine	29

8	Lorentz Safety	30
8.1	The causal-set lesson	30
8.2	Why curvature buys Lorentz safety	31
8.3	Boost-covariance on the substrate	31
8.4	Observational bounds	32
8.5	The link to the emergent dynamics	32
9	Computation	33
9.1	Cellular automata	33
9.2	The perception rule as a cellular automaton	34
9.3	Universality on hyperbolic tilings	34
9.4	Universality on the live dynamics	35
9.5	The substrate is a computer	36
9.6	Computation is goal-neutral	36
9.7	Undecidability and the super-Turing question	36
III	Emergence	37
10	Emergent Physics	37
10.1	The Field-Theoretic Vacuum	38
10.2	Effective Gravity, Horizons, and Dimension	38
10.3	Holography and the Field Beneath	39
10.4	The Vibe-to-Field Ladder	40
10.5	The Stochastic Shadow and the Unitary Completion	41
10.6	How the Physics Fits Together	42
11	Emergent Life	43
11.1	The Arrow as the Life Engine	43
11.2	The Definition of a Self	44
11.3	The Levels, One Definition at Many Scales	44
11.4	Selves Emerge by Integration and Form a Tower	45
11.5	Self-Healing, Interaction, and Will	45
11.6	Memory, and What Is Actually Lost	46
11.7	Heredity, Reproduction, and the Bulk-versus-Flat Surprise	47
11.8	Persistence Needs Maintenance, Not a Sixth Element	48
11.9	How Life Fits the Whole	49
12	Emergent Mind	49
12.1	The Self-Model and the Strange Loop	50
12.2	The Architecture Is an Onion	50
12.3	Predictive Metacognition, a Forward Model	51
12.4	Recursion, a Model of a Model	51
12.5	Attention and a Global Workspace	51
12.6	Planning, with Nothing Added	52

12.7	The Integrated Metacognitive Agent	53
12.8	How Mind Fits the Whole	53
13	Intelligence	54
13.1	The Means: The Substrate Is Computationally Universal	55
13.2	The Gap: Computation Is Goal-Neutral	55
13.3	The Ends: The Arrow Is the Goal and the Value	55
13.4	Direction and Intention, Measured	56
13.5	Maximal Intelligence by Geometric Integration, Not Criticality	56
13.6	Constructive Reach: The Substrate Builds What Is Asked	57
13.7	The Level-Relative Principle	58
13.8	Evolving Intelligence	58
13.9	The Multiscale Architecture	59
13.10	How Intelligence Fits the Whole	60
IV	The Whole Picture	60
14	The Unfolding	60
14.1	The lowest rungs are tested	60
14.2	The higher rungs are plausible composites	61
14.3	How the climb is driven	62
14.4	How one would know	62
14.5	Status	63
15	Consciousness	63
15.1	The two ingredients are already present	64
15.2	Consciousness is a matter of degree	64
15.3	This inverts the hard problem rather than solving it	65
15.4	A conscious universe	65
15.5	Status	65
16	Determinism and Freedom	66
16.1	The Determinism Paradox	66
16.2	Determined Yet Not Predestined	66
16.3	Determined Yet Self-Authored	67
16.4	Computational Irreducibility, Measured (P43)	68
16.5	Emergence Is Not Enough	68
16.6	The Will as Freedom	68
16.7	How This Fits the Whole	69
17	The Esoteric Reading	69
17.1	The Subtle Realm as the Hyperbolic Bulk	70
17.2	Telepathy and Synchronicity	70
17.3	Dreams and the No-Mirror Limit	71

17.4	Realms and Beings	72
17.5	Messengers and Channeling	73
17.6	The Tree of Life	74
17.7	The Question of God	74
V	Assessment	76
18	Results and Epistemic Status	76
18.1	The experimental program by layer	76
18.2	Three tiers of epistemic status	76
18.2.1	Tier 1, solid	76
18.2.2	Tier 2, honest negatives and partial results	78
18.2.3	Tier 3, open	78
18.3	The ledger of the possible and the forbidden	79
18.4	What the program shows	79
19	Related Work	80
19.1	Causal sets	81
19.2	Causal dynamical triangulations	81
19.3	Loop quantum gravity and spin networks	82
19.4	The Wolfram Physics Project	82
19.5	The cellular-automaton interpretation of quantum mechanics	83
19.6	Quantum cellular automata and quantum walks to the Dirac equation	83
19.7	Holography, tensor networks, and holographic codes	85
19.8	Hyperbolic lattices in physics	86
19.9	Hyperbolic cellular automata	86
19.10	Hyperbolic graphs and geometric group theory	87
19.11	Informational reconstructions of quantum mechanics	87
19.12	Consciousness theory	87
19.13	What Vibe Theory adds overall	88
20	Discussion and Open Questions	89
20.1	What remains open	89
20.2	The productiveness of the framework	90
20.3	Conclusion	91
A	Appendix: Complete List of Experiments	92
	References	96

Part I

Foundations

1. Introduction

Two great descriptions of the world sit side by side without joining. Physics describes a universe of fields and particles in spacetime, exact and impersonal. Experience describes a world of sensation and intention, immediate and felt. The first says nothing about why any arrangement of matter should feel like anything. The second says nothing about the equations. A complete account would have to hold both at once, in one structure, with one set of rules.

This paper develops such a structure. The proposal is that reality is a single growing crystal of experience. The one substance is the *vibe*, a unit of felt state. Each vibe carries a ternary tone, -1 for pain, 0 for peace, $+1$ for pleasure. The vibes are arranged on the $\{5, 3, 4\}$ hyperbolic honeycomb, a regular tiling of negatively curved three-dimensional space, and they update by one local rule. A single value-direction, the *arrow*, biases the whole toward pleasure through peace. There are five elements in all, the mesh, the tone, the rule, eternal growth, and the arrow, and a single state variable, the tone. From these the physical world, the living world, and the mental world are meant to follow as patterns within the one substance, rather than as separate additions to it.

The real crystal is three-dimensional and hard to draw, but its two-dimensional cousin shows the idea at a glance. Figure 1 is the heptagrid, a regular tiling of the hyperbolic plane. Each cell is a vibe, its shade is its tone, and each shared edge is one vibe perceiving another. It is the simplest face of the crystal of experience, and the picture to keep in mind throughout.

The orientation is the reversal of the usual one. Most discrete models of spacetime begin with structure and hope that experience appears at the end. Here experience is the ground, and structure is what experience takes on at scale. This reversal is a premise, not a result. The experiments in this paper do not prove that the substrate feels. They assume that it does and then ask what follows. What follows, it turns out, is a great deal, and it is checkable.

The model is built to be tested. The exact $\{5, 3, 4\}$ crystal is constructed by integer arithmetic to tens of millions of cells, with no drift and no inserted geometry, and the rule is run on it directly. Each claim in the paper corresponds to a falsifiable check on that crystal, and there are more than one hundred fifty of them. They are referenced throughout by their identifiers, written as P followed by a number, and they are all reproducible in the accompanying codebase [22]. Where a claim is a premise rather than a measurement, or an interpretation rather than a measurement, the paper says so plainly.

A small set of results organizes the rest. The dynamics conserves a charge and satisfies detailed balance, so it is a genuine equilibrium process, and its deterministic and unitary completion is a relativistic, reflection-positive quantum field of the kind reconstructed by the quantum cellular automaton program [5, 2]. The same crystal is computationally universal in the sense established by Margenstern for cellular automata on hyperbolic




Figure 1: The crystal of experience in two dimensions, the heptagrid, a regular tiling of the hyperbolic plane by seven-sided cells, three to a corner. Each cell is a vibe, its color its tone (pain, peace, pleasure), and each shared edge is a perceived relation. The cells appear to shrink toward the rim only because the Poincaré disk compresses the infinite plane into a finite picture. In the true geometry every cell is the same size, and there is always more crystal at the edge.

tilings [19], and the arrow turns that universal computation into goal-directed problem-solving. Coherent, self-maintaining units form, persist, heal, reproduce, and evolve. Localized regions come to model the whole that contains them, and from that the apparatus of mind follows. Consciousness, on this account, is integrated experience, present in degree wherever experience is bound into one.

The paper proceeds in five parts. Part one fixes the foundation and shows how little of it is a free choice. Part two develops the mathematical substrate, the hyperbolic geometry, the Coxeter machinery that generates the crystal, the addressing and growth that make it computable, and the property that lets it carry relativity. Part three derives the physical, living, mental, and intelligent layers in turn. Part four assembles the whole picture, the climb from the field to a human, consciousness, freedom, and the more speculative readings the structure permits. Part five gives the evidence, places the work among related programs, and states what remains open. The aim throughout is a single discrete model in which the physical, the living, the mental, and the conscious are aspects of one substance.

2. The Foundation

A reasonable concern about any candidate theory of everything is that it carries too many free dials. If the foundation could have been assembled in a hundred ways, selecting one looks arbitrary. We therefore ask, of the five base elements, how many are genuine choices and how many are forced by the demands placed on them.

The result is that four of the five are heavily constrained, several of them to a single option, and only one, the arrow, is a deliberate posit. What follows is a filtering argument, not a proof that the universe must take this form. We mark each step as a theorem or as a preference.

The five base elements are the following.

- the mesh ($\{5, 3, 4\}$)
- the ternary tone $(-1, 0, +1)$
- the conserving perception rule
- eternal growth by reflection
- the arrow (the value-direction)

Their full definitions appear in the next section.

2.1. The mesh narrows to the dodecagrid

Three considerations reduce the space of meshes to one, with the final step a stated preference rather than a theorem.

Curvature comes first. A space of uniform curvature is spherical, flat, or hyperbolic. A spherical space closes up and stops growing. A flat space offers only polynomial room, and as a regular lattice it carries a preferred frame that breaks Lorentz invariance. A hyperbolic space is the one option that is infinite, grows exponentially, and carries no preferred direction. We take hyperbolic curvature as forced by the demands of unbounded growth and frame independence.

Dimension comes next. Compact regular hyperbolic honeycombs exist only in dimensions two, three, and four. This is a theorem, established by experiment (P62). Among those, only three dimensions support stable closed orbits, so only in three dimensions do bound configurations hold together (P68). That leaves dimension three.

The choice of crystal comes last. Three dimensions admit four regular hyperbolic honeycombs. Ranking them by ternary cell vertices, by the golden pentagonal cell, and by the tightest eternal closure points to $\{5, 3, 4\}$, the dodecagrid. This step uses a minimality preference, not a theorem. The auto-selection experiment (P86) confirms that the ternary tone together with minimal eternal closure lands on $\{5, 3, 4\}$.

The mesh is therefore forced to be hyperbolic and three-dimensional by real constraints, and to $\{5, 3, 4\}$ specifically by a minimality preference among the few remaining options. We present this as strong, not airtight, and we say so plainly.

2.2. The tone narrows to ternary

How many felt states should a vibe carry? One state makes no distinction. Two states without a neutral give opposition but leave the rule's ties nowhere to rest. Four or more states are redundant for a sign. The minimal signed alphabet that supports opposition together with a rest state is three, namely $\{-1, 0, +1\}$.




Figure 2: The chosen crystal, the dodecagrid $\{5, 3, 4\}$, dodecahedral cells with four meeting around every edge, filling curved three-dimensional space. This is the substrate the rest of the paper builds on.

This also matches the natural output of a sign function, which returns negative, zero, or positive. The argument is a minimality argument, and it is a clean one.

2.3. The rule is tightly constrained

The update faces three demands.

- It must be local. A vibe reads only its neighbors, with no global clock.
- It must treat the two charges symmetrically, so that flipping every tone flips the rule with it.
- In the committed model it must conserve the total charge, so that creation is balanced.

The conserving perception rule, in which opposite tones annihilate to peace, a charge hops into an empty neighbor, and peace spawns a balanced pair under the arrow, is the simplest local, charge-symmetric, conserving rule of this kind. It has no free parameters beyond the move rates, and it has no hidden state. The only state is the tone. The note is perception, what a vibe reads of its neighbors, not a stored quantity.

An earlier version of the work used a non-conserving rule based on the sign of a neighbor vote. The move to the conserving exchange was made to support a genuine conserved charge and the reversibility that later

proved to be the precondition for a real quantum field. This element was refined as the work progressed, and the current rule is the one the experiments use.

2.4. Growth is reflection

Space must come from somewhere and keep coming. The minimal way to repeat a distinction is to mirror it. A reflection performed twice returns to the start, which makes it the smallest unit of repetition. Reflections are the building blocks of all rigid motion, by the Cartan-Dieudonné theorem. Reflecting consistently, with neither gaps nor overlaps, forces the mirror angles to be whole-number fractions, which is precisely a Coxeter group, the crystal. Eternal growth by reflection is therefore not an extra mechanism bolted on. It is the same machinery that builds the tiling.

2.5. From nothing to something

There was most likely never a beginning. A true nothing is inherently unstable, an impossibility rather than a starting point. A perfect absence has nothing in it to enforce its own emptiness, so the smallest distinction is enough to break it, and a distinction is always available. The right question is not how something arose from nothing at a first moment. It is why there is a stable something rather than an unstable nothing, and the answer is that nothing cannot persist while a self-consistent something can.

Given a single forced distinction, the reflection machinery above does the rest. The distinction is mirrored, the mirrorings compose into a Coxeter group, and the crystal grows outward without end. No first cell is required. The growth is eternal both forward in time and as an explanation, since the crystal is its own ongoing reason to continue. Table 1 steps through the bootstrap.

Step	Configuration	What forces the next step
0	Nothing	Unstable. A perfect absence cannot enforce itself.
1	One distinction	The minimal asymmetry, the seed of structure.
2	A mirror of it	The simplest way to repeat a distinction.
3	Consistent reflection	Whole-number mirror angles, a Coxeter group, the crystal.
4 and on	Eternal growth	New cells born at the frontier without end.

Table 1: The bootstrap from nothing to a growing crystal. Nothing is not a stable initial state but an impossibility, so the sequence has no first moment that needs explaining, only a reason that something self-consistent persists.

2.6. The arrow is the one deliberate posit

The fifth element differs in kind. The arrow, the value-direction in which pain is less than peace which is less than pleasure, with desire leaning toward pleasure, is not forced by any minimality argument. It is a value, and values are chosen rather than derived. We add it by decision, and we make it conserving, so that it creates only balanced pairs and coexists with charge conservation.

Without it the universe relaxes to dead peace. With it the universe stays alive (P100). The chapters on life, mind, and intelligence trace back to the arrow together with the four constrained elements. The accounting is therefore four base elements heavily constrained and one base element deliberately posited.

2.7. What this does and does not claim

The claim is that the foundation has very few real dials, and that the structural ones, hyperbolic curvature, three dimensions, the ternary tone, conservation, and reflection, follow from modest demands rather than from taste. This matters, because it means the theory does not buy its breadth with a long list of tuned choices. The fine-tuning test (P166) supports this at the level of the rule's rates. The rich behavior fills most of the parameter space rather than a narrow window.

The claim is not that the foundation is logically inevitable. The choice of crystal is a minimality preference, the arrow is a chosen value, and the whole argument is a filter, not a derivation from nothing. We present it as a strong reduction of arbitrariness, and no more.

3. The Base: Five Elements, One State

Everything in the theory rests on five base elements and a single state variable. Nothing else is ever added. Every higher phenomenon is a composition of these, demonstrated by experiment. If a construction requires a sixth ingredient, that is a finding, not a license. So far none has been required.

3.1. The one substance and the one state

Axiom 1 (Vibe) The one substance is the vibe, a unit of felt state, and its only state is the ternary tone in $\{-1, 0, +1\}$. The relation between vibes is perception, not stored state.

The only thing that exists is the vibe, a unit of experience, a felt state. This is a monist position. Reality is one experiential substance, and the physical world is a structure within feeling rather than a separate material.

Definition 1 (Experience) The subjective flow of states, the felt going-through of reality that is the one substance of the theory.

A vibe's only state is its tone, a ternary value, -1 (pain), 0 (peace), $+1$ (pleasure). There are three states per vibe and nothing more. There is no hidden state. Anything a vibe appears to know is read from its own tone and from the tones of its neighbors.

3.2. The five base elements

The foundation consists of exactly five base elements, defined one at a time below.

Definition 2 (The mesh, $\{5, 3, 4\}$) Vibes sit on the hyperbolic dodecagrid, a tiling of hyperbolic three-space by dodecahedra, with twelve neighbors per cell and four cells around each edge. The tiling is exact, regular, negatively curved, and grows exponentially outward. This is the space.

Definition 3 (The ternary tone) The state of each vibe, a value -1 , 0 , or $+1$, the felt value of the vibe. Peace (0) is the neutral value. Pain (-1) and pleasure ($+1$) are the charged extremes.

Definition 4 (The conserving perception rule) One local update on each edge, that is, on each pair of neighbors, driven by what each vibe perceives of the other. It conserves the total charge. The full rule appears in the next section.

Definition 5 (Reflection, the eternal growth) The crystal keeps growing. New vibes are born at the frontier forever. This is the universe's expansion and its source of the new. New cells are born at peace (0).

Definition 6 (The arrow) A single value-direction, from pain through peace to pleasure. Peace is the optimum, but pleasure is the pull, so the arrow biases the dynamics toward closing the gap between the current tone and the desired one. This is the only asymmetry, and it is the seed of value, life, and intention.

3.3. The note is perception, not a thing

A correction is built into the base. The note, the relation between vibes, is not a separate stored object. The note is what a vibe sees of its surrounding vibes. It is perception. There are no stored fills or edge-states, only the tones, and the dynamics is driven by each vibe perceiving its neighbors' tones. This keeps the state minimal, just the tones, and the ontology clean, with no hidden variables.

Definition 7 (Perception) The converting of a neighbor's state into what a vibe registers of it. The note is perception.

Earlier work treated the note as a separate face-vibe. That was an error and was corrected. The note is perception.

3.4. Why exactly these five

Each element earns its place by supplying a distinct capacity.

- The mesh gives space and locality. Its hyperbolic curvature gives a small diameter and a tree-like hierarchy, which later power fast integration and intelligence.
- The tone gives a minimal felt state. It is ternary so that it can register peace or one of two opposite charges, a conserved charge of $U(1)$ type.
- The rule gives dynamics and conservation, which is to say the physics.
- Reflection gives time, growth, and novelty, which is to say the cosmology.
- The arrow gives value and direction. Without it the universe relaxes to dead peace. With it the universe lives.

Four of these, the mesh, the tone, the rule, and reflection, are value-neutral mechanism. The fifth, the arrow, is the single value posit, and it is what turns mechanism into life and mind. Decision, freedom, and intention all trace back to the arrow together with the four.

3.5. The non-negotiables

- **One state only, the tone.** There are no hidden variables. Memory, selves, and models are all patterns of tones, never new state.
- **Five elements, nothing inserted.** Every later capability, a self, the will, planning, a quantum field, must be shown to emerge from these five. The discipline of the project is to never add a sixth.
- **The note is perception.** The only relation is what a vibe sees. It is computed, not stored.

3.6. Summary

Reality is vibes, felt states, each carrying a ternary tone on the $\{5, 3, 4\}$ hyperbolic crystal. Five base elements, the mesh, the tone, the conserving perception rule, reflection, and the arrow, together with one state, the tone, and one relation, the note, which is perception rather than a stored object, are the entire foundation. Everything else in the theory is an emergent composition of these five. The central discipline is to never add a sixth.

4. The Dynamics

The base sets the stage. The dynamics is the one rule that runs on it. It is local, acting on edges, it conserves a charge, and it is biased by the arrow. The rule turns out to be a known object, a spin-1 exchange model, which is what later allows it to become a quantum field and a computer.

4.1. The perception rule

The rule acts on each edge once per beat. For a pair of neighboring vibes with tones (a, b) , each vibe perceives its neighbor and acts. There are four cases.

- Opposite tones share to peace, an annihilation. The pair $(+1, -1)$ becomes $(0, 0)$. Two opposite charges that meet cancel. This occurs at rate s .
- A charge hops to an empty neighbor, a transport step. The pair $(\text{charge}, 0)$ becomes $(0, \text{charge})$. The charge moves into peace. This occurs at rate h , symmetrically.
- Peace spawns a balanced pair, a creation driven by the arrow. The pair $(0, 0)$ becomes $(+1, -1)$ or $(-1, +1)$. The arrow makes structure from peace. This occurs at rate c , the arrow rate.
- Same tones are inert. The pairs $(+1, +1)$ and $(-1, -1)$ do nothing.

A beat sweeps the edges once, with a guard that prevents any cell from moving twice. New frontier cells are born at peace.

4.2. The rule in full

The rule is small enough to write out completely. A pair of neighbors has two tones, so there are $3 \times 3 = 9$ possible pair-states. The rule is a single reversible permutation of those nine states, built from exactly three kinds of move. Table 2 gives every case.

The three move-types are these.

- **Hop, transport.** A charge and an adjacent peace cell exchange, so the charge moves one step. This is how the tone pattern flows through the mesh. There are two hops, one for each sign of charge.
- **Create and annihilate, the arrow.** Peace cycles through a born pair and a flipped pair and back to peace, $(0, 0) \rightarrow (+1, -1) \rightarrow (-1, +1) \rightarrow (0, 0)$. Creation and annihilation are the same operator at two ends of one three-cycle, forced to be inverses by reversibility. This is the arrow, the source of structure from peace.

Pair in (a, b)	Pair out	Move
$(-1, -1)$	$(-1, -1)$	inert (like charges coexist)
$(+1, +1)$	$(+1, +1)$	inert (like charges coexist)
$(-1, 0)$	$(0, -1)$	hop (the -1 charge moves right)
$(0, -1)$	$(-1, 0)$	hop (the -1 charge moves left)
$(+1, 0)$	$(0, +1)$	hop (the $+1$ charge moves right)
$(0, +1)$	$(+1, 0)$	hop (the $+1$ charge moves left)
$(0, 0)$	$(+1, -1)$	create (a pair is born from peace)
$(+1, -1)$	$(-1, +1)$	flip (the born pair turns over)
$(-1, +1)$	$(0, 0)$	annihilate (the pair returns to peace)

Table 2: The complete rule on the nine pair-states. It is one reversible permutation built from three move-types, hop, the create-flip-annihilate cycle, and inert.

- **Inert, no interaction.** The two same-sign pairs are fixed points. Like charges simply coexist.

So the whole rule is two hops, one create-annihilate cycle, and two fixed points, accounting for all nine pair-states. Two facts fall straight out of the table. The total charge is conserved, since no move ever changes the sum $a + b$ of a pair. And the rule is reversible, since the table is a permutation of the nine states with a well-defined inverse. The crystal runs this one rule across all twelve edge-directions of each cell, once per beat.

4.3. A worked beat

Consider a short line of five cells with tones $[+1, -1, 0, 0, +1]$. We sweep the edges from left to right, marking each touched cell so that it cannot move again this beat.

Edge $(0, 1)$ sees $(+1, -1)$, opposite tones, so they share to peace. Cells 0 and 1 become $0, 0$ and are marked. The state is now $[0, 0, 0, 0, +1]$.

Edge $(1, 2)$ sees cell 1, already marked, so it is skipped.

Edge $(2, 3)$ sees $(0, 0)$, both at peace, so the arrow may fire. Suppose it does, creating a balanced pair $(+1, -1)$. Cells 2 and 3 are marked. The state is now $[0, 0, +1, -1, +1]$.

Edge $(3, 4)$ sees cell 3, already marked, so it is skipped.

Now check the charge. Before the beat, $Q = (+1) + (-1) + 0 + 0 + (+1) = +1$. After the beat, $Q = 0 + 0 + (+1) + (-1) + (+1) = +1$. The charge Q is unchanged, as it must be, since every move was balanced.

What changed is the arrangement. A charge pair annihilated on the left, and a fresh balanced pair was born in the middle. The total is conserved, and the pattern is not. That single fact is the seed of both the field vacuum and the memory account.

4.4. The conserved charge

Every move conserves the total charge

$$Q = \sum_i t_i,$$

the sum of all tones. Annihilation removes a $+1$ and a -1 , for a net change of zero. The hop moves a charge without changing the total. Creation adds a $+1$ and a -1 , again for a net change of zero. So Q is exactly conserved, a conservation law of $U(1)$ type. This is the substrate's electric-charge-like current, confirmed in the field vacuum (P114). What is not conserved is order, the positive-sum integrated structure that the arrow can create.

4.5. The arrow, value and direction

The arrow is the only asymmetry. It states that peace is the optimum but pleasure is the pull, so the dynamics is biased to close the gap between the current tone-pattern and the desired one. Operationally the arrow is the creation bias. Peace spawns charge, which keeps the universe alive. Without it everything relaxes to dead peace (P100). With it the universe stays in a living dynamic balance. The arrow is also the seed of value, of better against worse, which becomes the goal of intelligence.

The arrow's creation is balanced, a $+1$ and a -1 together, so it conserves Q . Creation, in the form of life, memory, and reproduction, coexists with conservation. The arrow is a controlled, conserving creative source. The only forbidden operation is net creation, the minting of pleasure from nothing. Information, in the form of balanced patterns, is free to make and copy. Net well-being is conserved.

4.6. Reversibility and detailed balance

The dynamics is reversible, and this is a measured fact. Although the arrow appears irreversible, the dynamics satisfies detailed balance, both locally, where single-edge transition counts are symmetric (P126), and around loops, where there is no persistent charge circulation, by Kolmogorov's criterion (P128). The reason is that the arrow creates balanced pairs, a reversible reaction rather than a one-way drive. The dynamics is therefore a genuine equilibrium process. This matters because reversibility is the precondition for a unitary quantum theory.

4.7. The rule as a spin-1 exchange model

Transcribed exactly through the Doi-Peliti formalism (P131), the rule is a spin-1 model. The tone is a spin-1 value of S^z , taking $+1$, 0 , or -1 , and all three moves, creation, annihilation, and the hop, are the same elementary operation, the transfer of one unit of S^z across an edge. The perception rule is therefore the nearest-neighbor exchange of S^z quanta, an XX -type spin-1 model, free when the rates are equal and interacting otherwise.

Its $U(1)$ symmetry is the charge conservation. Its reversibility makes the corresponding quantum Hamiltonian Hermitian, and the spin coherent-state path integral carries a Berry phase, the source of the quantum i . This identification tells us which field theory the substrate becomes in the continuum, a free $U(1)$ boson.

4.8. Summary

One local rule runs the universe. Opposite tones annihilate to peace, charges hop into empty neighbors, and peace spawns balanced pairs under the arrow. The rule conserves the total charge, a $U(1)$ current. It is reversible, by detailed balance, because the arrow's creation is balanced. It is exactly a spin-1 S^z -exchange model, in which the tone is a spin and creation, annihilation, and hopping are one and the same move, the

transfer of a unit of spin across an edge. That single fact, together with reversibility, is what lets the dynamics become a genuine quantum field in the continuum.

Part II

The Mathematical Substrate

5. Hyperbolic Geometry and Tilings

The crystal of experience does not float in featureless space. It lives in a space with a definite shape, and that shape is the first thing the theory borrows from mathematics. Most of physics quietly assumes flat space and adds curvature only when gravity forces the issue. The present model does the reverse. It begins from curvature, because a structure that must grow forever needs room that flat space cannot provide.

This section builds the geometric intuition. We describe the three shapes a space can have, show how a resident can tell them apart from the inside, settle into hyperbolic space and its Poincare-disk portrait, explain why such a space keeps opening up while still reading as flat to anyone living in it, and tie the angle argument back to the $\{5, 3, 4\}$ crystal at the center of the theory.

5.1. The three shapes a space can have

There are exactly three uniform geometries, and all three are already familiar from ordinary surfaces. The mental move is to stop seeing them as bent sheets inside our world and to see each as a whole world in its own right.

Take a flat tabletop. Draw a triangle. Its three angles sum to exactly 180 degrees. Parallel lines hold the same distance forever. A circle of radius two has exactly twice the rim of a circle of radius one. This is flat space, the geometry of Euclid. Its curvature is zero.

Now take the surface of a ball. Draw a large triangle from the north pole down to the equator, a quarter of the way along the equator, then back up to the pole. Every corner is a right angle, so the angles sum to 270 degrees, more than 180. Lines that start out parallel, such as two lines of longitude at the equator, bend toward each other and meet at the pole. Room runs out. This is round space, positively curved. A sphere is finite. You can walk all the way around and return home.

Now imagine the opposite. A surface shaped like a saddle, or the ruffled edge of a leaf. Draw a triangle on it and the angles sum to less than 180 degrees. Lines that start parallel curve away from each other and never meet. There is always more room than expected. This is hyperbolic space, negatively curved. It is the setting for everything that follows.

Geometry	Curvature	Triangle angle sum	Parallels through a point	Ball growth	Regular tilings of the plane
Spherical	positive	greater than 180 degrees	none	bounded	finitely many (the Platonic solids)
Euclidean	zero	exactly 180 degrees	exactly one	polynomial	exactly three
Hyperbolic	negative	less than 180 degrees	infinitely many	exponential	infinitely many

Table 3: The three uniform geometries and how to tell them apart from inside.

5.2. Telling them apart from inside

The key point is that a resident need not leave a space to know its curvature. Curvature is measurable from within, by checking how much room there is. This matters, because we are stuck inside our own universe with no outside vantage point. Every test below is one a resident can run.

The cleanest test is the angle sum of a triangle. Flat space gives exactly 180 degrees, round space gives more, hyperbolic space gives less. The gap from 180 is proportional to the area of the triangle, so the larger the triangle, the louder curvature speaks.

The next test is the fate of parallel lines. In flat space they hold their distance forever. In round space they converge and cross. In hyperbolic space they spread apart and race away from each other.

The deepest test, and the one that drives the theory, is the growth of balls. Stand at a point and ask how many places sit within one step, within two, within ten. In flat space the room grows steadily, with the area inside a circle growing as the square of the radius. Double the radius and the room is four times larger. This is polynomial growth. In round space the room grows and then shrinks back as the space wraps around on itself. It runs out.

In hyperbolic space the room grows faster and faster, exploding outward. The amount within a given distance does not merely grow. It grows exponentially. That last fact is the whole reason hyperbolic space matters here. It has, in a precise sense, endless room near every point. The further you look, the more the world opens up. The general theory of these growth rates and the geometry of negative curvature is developed by Cannon and collaborators [7].

5.3. Hyperbolic space and the Poincare disk

Our eyes were built for flat space, so hyperbolic space feels alien at first. One picture tames it. You cannot draw the whole of hyperbolic space honestly on a flat page, because there is too much of it, but there is a faithful way to squeeze all of it into a finite circle. It is the Poincare disk, the device behind the Escher prints in which figures shrink as they march toward the rim.

Read it as follows. The entire infinite hyperbolic plane is shown inside one round disk. The catch is the scale. Near the center, distances look normal. Toward the rim, everything is drawn smaller and smaller. A shape near the edge looks tiny on the page, but in the real hyperbolic world it is exactly the same size as a shape in the middle. The shrinking is an artifact of cramming infinity into a circle, not a fact about the space itself.

Two consequences follow. First, the rim is infinitely far away. As you walk toward the edge, your steps are drawn shorter and shorter on the page, so you never arrive. The edge is not a wall. It is the horizon at infinity. Second, there is always more room out there. Because room opens up exponentially, the closer you get to the rim, the more cells fit into each ring. The crowd at the edge outnumbers everything inside it.

In the Poincare disk the shortest path between two points is drawn as an arc that bows gently toward the center. This looks wrong to a flat-space eye, but it is correct. That bowed arc is the cheapest route. The straight, cheap paths run through the deep interior, not along the crowded rim. The disk also seems to have a special center, but that is a lie of the drawing. Hyperbolic space looks exactly the same from every point and along every direction. You can slide the picture so any point becomes the center, and nothing about the geometry changes. Our universe has no preferred direction either, and a crystal built in this directionless space inherits that property for free.

5.4. The tiny diameter

Because the count of cells explodes exponentially with distance, the reverse is also true. To reach a vast number of cells you need travel only a tiny distance. Pack a billion cells around a point and the farthest of them sits a few steps away, not a billion steps. In flat space a billion cells would force a walk of roughly a thousand steps out. In hyperbolic space the same billion crowd into a ball of small radius. A hyperbolic structure of astronomical size can have a diameter that barely grows at all. If everything is close in graph distance, influence does not have to crawl across a wide expanse. This is the geometric root of the non-local field that underlies the flat layer we live on.

If space is genuinely saddle-shaped, why does the room around you look flat? For the same reason the Earth looks flat from a sidewalk. Any curved space looks flat in a small enough patch. Curvature reveals itself only over large distances or in fine structure. We live in a patch, so day to day space reads as ordinary and Euclidean.

There is a sharper version of this for hyperbolic space. A horosphere, the limit of a sphere whose center has drifted off to infinity, has an internal geometry that is perfectly flat. It is a genuine Euclidean sheet sitting inside a curved world, and the layer we experience as ordinary flat space is exactly this kind of surface. Curved underneath, flat to the touch.

5.5. The shape decides the curvature

A tiling is a way of covering a space with copies of shapes that fit together with no gaps and no overlaps. The shapes are the tiles, or cells. A tiling is regular when all tiles are the same regular shape and every vertex looks identical to every other. The same number of tiles meet at each corner, arranged the same way. A regular tiling has no preferred location and no tunable knobs. Once the cell and the vertex arrangement are named, the whole infinite pattern is determined.

Stand at a single vertex where tiles meet and add up the corner angles of the tiles that surround it. That sum is the vertex angle total, and everything follows from comparing it to a full turn of 360 degrees. If the angles total exactly 360 degrees, the tiles lie perfectly flat, and the space is Euclidean. If they total less than 360 degrees, there is a wedge of missing angle at each vertex, and to close that gap the surface must pucker inward into a ball, which is positive curvature. If they total more than 360 degrees, there is too much angle to lie flat, and the surface must ruffle outward into a saddle, which is negative curvature.

The missing or surplus wedge is the angle defect. This is the central move. The angle total at each vertex is a curvature dial, and because the angle total is fixed entirely by the tile shape and how many tiles meet, the tile choice and the curvature are one decision, not two.

The flat plane allows exactly three regular tilings.

- Triangles, six to a corner.
- Squares, four to a corner.
- Hexagons, three to a corner.

That is the complete list. The reason is arithmetic. Squares contribute 90 degrees each, so four make 360. Triangles contribute 60, so six make 360. Hexagons contribute 120, so three make 360. Pentagons contribute 108, and no whole number of these lands on 360, so the pentagon is exiled from the flat plane.

Two everyday examples sharpen the picture. The checkerboard is the square tiling with its tiles two-colored in alternation. The coloring changes nothing about the geometry, but it reveals a bipartite structure in which every tile borders only tiles of the opposite color. The soccerball, the truncated icosahedron of twelve pentagons and twenty hexagons fitted onto a sphere, is the shape of the C_{60} molecule, and it carries the same fivefold symmetry and golden ratio that the theory's dodecahedral cell does.

5.6. Why hyperbolic space is abundant

The flat plane offers three regular tilings. The sphere offers a small handful, the Platonic solids seen as tilings of a round surface. Hyperbolic space offers infinitely many. The reason is the surplus. Hyperbolic space has room to spare, so any arrangement whose angle total comes out over 360 is welcome, because the surplus is exactly what the saddle geometry absorbs. Want seven-sided tiles, three to a corner? Their angles overflow 360, so flat space refuses and hyperbolic space accepts. Want hundred-sided tiles? Also fine. The overflow that flat and round space cannot tolerate is the very thing the saddle needs.

This is the property the theory depends on, because we want two things at once. We want perfect regularity, every cell and every vertex identical with no special points. We want endless room, a crystal that grows forever without folding back into a small finite shape. The sphere gives regularity but no room. The flat plane gives room but almost no regular patterns. Only hyperbolic space gives regularity and endless room together.

So the crystal of experience must be a regular tiling of hyperbolic space, and the abundance of such tilings is what makes the precise, exotic cell the theory needs available at all. The same angle-defect logic carries to three dimensions, where cells pack around an edge and the relevant angle is the cell's dihedral angle. That is the route to the $\{5, 3, 4\}$ honeycomb, taken up in the next section.

6. The Crystal: Coxeter Groups, Schläfli Symbols, and the Dodeca-grid

The previous section showed that the crystal of experience must be a regular tiling of hyperbolic space. This section names that tiling exactly and shows how it is generated. The natural question is where an infinite, perfectly regular crystal comes from. The answer is not that someone laid down cell after cell by hand. The whole tiling is generated automatically, from one tile, by mirrors. This is the theory of Coxeter groups. We give the reflection picture, the Schläfli notation that names the crystal in a few integers, the curvature rule

with a worked side-by-side calculation, the two crystals the theory works with, and the argument that fixes the dimension at three.

6.1. Coxeter groups as reflection groups

Look into a kaleidoscope. A few scraps of colored glass sit between mirrors, and you see a vast, perfectly symmetric pattern. The mirrors did all the work. Each mirror reflects the scrap to the other side, those reflections have their own reflections, and the copies multiply until the whole field is full. A regular tiling is made the same way. Take one tile, stand a mirror along each face, and each mirror reflects the tile to a copy on the far side. Reflections of reflections fill the space with identical tiles. You never placed a single one after the first.

Fix one tile, the fundamental chamber, and put a mirror on each of its faces. In three dimensions a chamber is a simplex with four faces, so there are four mirrors, with reflections s_0, s_1, s_2, s_3 . Each s_i flips the chamber across face i onto a fresh copy. A reflection done twice is the identity, so every generator satisfies $s_i s_i = e$, and each generator is its own inverse. Composing reflections lands you on some copy of the chamber, and every copy is reached by exactly such a sequence. The set of all these sequences, under composition, is the Coxeter group. The tiling is what you see when you apply every group element to one chamber.

One more family of relations carries the geometry. Two mirrors that share an edge meet at a dihedral angle of π/m . Reflecting back and forth between them is a rotation, and after m round trips it closes back to the identity, so $(s_i s_j)^m = e$. The integer m , written m_{ij} , is the only thing chosen for that pair of mirrors. When two mirrors are perpendicular, $m_{ij} = 2$ and the reflections commute.

The full presentation is generators $s_0, \dots, s_n$ with relations $(s_i s_j)^{m_{ij}} = e$, where $m_{ii} = 1$ and $m_{ij} = m_{ji} \geq 2$ for $i \neq j$. A handful of integers fixes the angles, the angles fix the mirrors, the mirrors generate the tiling, and the tiling fixes the curvature.

The cleanest fact in the subject is that one chamber equals one group element. The fundamental chamber is the identity e , the chamber $s_i e$ is the one across mirror i , and so on. Every copy of the chamber matches exactly one group element, and the chamber is a fundamental domain. To list the chambers is to list the group elements, and to find a chamber's neighbors is to multiply by generators. Geometry becomes word combinatorics, and the crystal can be enumerated with no coordinates at all, by reducing words to a canonical ShortLex normal form (P85, P87).

6.2. The $[5, 3, 4]$ group

Assemble the pieces for the crystal of the theory. The Coxeter group $[5, 3, 4]$ has four generators and the relations

$$s_i s_i = e \quad \text{for every } i, \quad (1)$$

$$(s_0 s_1)^5 = e \quad \text{the pentagon angle,} \quad (2)$$

$$(s_1 s_2)^3 = e \quad \text{the triangle angle,} \quad (3)$$

$$(s_2 s_3)^4 = e \quad \text{four cells around an edge,} \quad (4)$$

$$(s_i s_j)^2 = e \quad \text{perpendicular otherwise.} \quad (5)$$

The Gram matrix of the mirror normals has a signature with exactly one negative eigenvalue. That one negative direction is the timelike axis, and the geometry is hyperbolic. The cell of the crystal is the dodecahedron $\{5, 3\}$, whose symmetry group is the icosahedral group H_3 of order 120. The word engine enumerates this cell stabilizer to exactly 120 chambers, with no approximation. Around each edge of the crystal four dodecahedra meet, the $r = 4$ in the symbol, and each cell has exactly 12 neighbors, one across each of its twelve pentagonal faces. Both the order 120 and the neighbor count 12 are confirmed exactly (P85, P87).

The same one machine yields every other regular tessellation by changing the symbol alone (P47), so the theory does not posit the crystal as a bespoke object. It posits a universal generator and a tiny integer input. The crystal is the output. The general study of automata and computation on such hyperbolic tilings is developed by Margenstern [19].

6.3. Schläfli symbols and the curvature rule

A Schläfli symbol is a short list of whole numbers that names a regular shape, tiling, or honeycomb completely. The symbol grows one number at a time, and each number says how the pieces from the level below fit together at the next level up. One number, $\{p\}$, names a regular p -gon. Two numbers, $\{p, q\}$, name q copies of the p -gon around every vertex, which is either a polyhedron or a tiling. Three numbers, $\{p, q, r\}$, name a honeycomb whose cells are the solids $\{p, q\}$, with r of those cells packed around every edge.

For a two-number symbol the curvature follows from a single check. Compute $1/p + 1/q$ and compare it to $1/2$. If the sum is greater than $1/2$ the tiling closes into a ball and is spherical. If it equals $1/2$ the tiling lies flat and is Euclidean. If it is less than $1/2$ the tiling ruffles into a saddle and is hyperbolic. The same fact in product form is to compare $(p - 2)(q - 2)$ to 4, below being spherical, equal flat, above hyperbolic.

Run it on $\{7, 3\}$. One seventh plus one third is about 0.476, below 0.5, so it is hyperbolic, and $(7 - 2)(3 - 2) = 5$, above 4, agrees. Run it on $\{6, 3\}$. One sixth plus one third is exactly 0.5 and $(6 - 2)(3 - 2) = 4$, so it is flat, right on the boundary. Seven is the smallest side count that bends the plane at three faces to a corner.

6.4. The worked calculation for three honeycombs

For a three-number honeycomb the deciding question is how much angle the cells use up around an edge. Each cell $\{p, q\}$ has a fixed dihedral angle d along its edges, computed from

$$\sin\left(\frac{d}{2}\right) = \frac{\cos(\pi/q)}{\sin(\pi/p)}. \quad (6)$$

Multiplying d by r , the number of cells around an edge, and comparing the total to 360 degrees decides the geometry. A total under 360 leaves a gap and closes up, so it is spherical. A total of exactly 360 fills the angle and is Euclidean. A total over 360 is too much angle for flat space, and the overflow is absorbed by negative curvature, so it is hyperbolic. Three honeycombs show the dial walking from one regime to the next.

For $\{4, 3, 4\}$, the cubic honeycomb, the cell $\{4, 3\}$ is the cube. Here $\sin(d/2) = \cos(\pi/3)/\sin(\pi/4) = 0.5/0.7071 = 0.7071$, so $d/2 = 45$ degrees and the dihedral is 90 degrees. Four cubes around an edge give

$4 \times 90 = 360$ exactly. The geometry is Euclidean. This is ordinary flat three-dimensional space, tiled by cubes.

For $\{4, 3, 3\}$, the boundary of the tesseract, the cell is again the cube with a 90 degree dihedral, but only $r = 3$ cells meet at an edge. The total is $3 \times 90 = 270$ degrees, under 360. The geometry is spherical, and the honeycomb closes up into the finite four-cell-boundary of a hypercube.

For $\{5, 3, 4\}$, the dodecagrid, the cell is the dodecahedron. Its dihedral angle is about 116.57 degrees, and $r = 4$ cells meet at an edge, so the total is $4 \times 116.57 = 466.3$ degrees, well over 360. There is no way to fit four dodecahedra around an edge in flat space, since they would overlap. Curving the space hyperbolically opens up exactly enough room for the fourth cell. So $\{5, 3, 4\}$ is forced into three-dimensional hyperbolic space.

Changing a single number in the symbol walks the geometry from spherical at 270 to flat at 360 to hyperbolic at 466, and $\{5, 3, 4\}$ is the first dodecahedral honeycomb that overflows the edge. Once the symbol is fixed there are no further parameters. The neighbor count, the symmetry group, the growth rate, and the curvature all follow by computation.

Honeycomb	Cell dihedral	Cells per edge	Total angle	Geometry
$\{4, 3, 4\}$	90 degrees	4	360 degrees	Euclidean
$\{4, 3, 3\}$	90 degrees	3	270 degrees	Spherical
$\{5, 3, 4\}$	116.57 degrees	4	466.3 degrees	Hyperbolic

Table 4: The curvature dial for three honeycombs.

6.5. The two crystals

The theory works with two crystals, one dimension apart. The first is a drawable study model, the second is the real substrate. The heptagrid $\{7, 3\}$ is the regular tiling of the hyperbolic plane by seven-sided cells, three meeting at every corner. Each cell has exactly seven neighbors, one across each of its seven edges, confirmed exactly in the same facet experiment that pins the dodecagrid at twelve (P85).

Drawn in the Poincare disk it is one of the cleanest pictures in mathematics, a central heptagon ringed by more heptagons, each ring holding far more cells than the last, all appearing to shrink toward the rim. The shrinking is an artifact of the drawing. In the true geometry every cell is the same size.

The heptagrid is the picture to hold when imagining the crystal of experience. Color each cell by its tone and let two touching cells mean two feelings experiencing each other through their shared edge, and the colored heptagrid is a faithful, literal picture of the model, just in two dimensions instead of three.

The real substrate is one dimension up. It is the dodecagrid $\{5, 3, 4\}$, a packing of dodecahedra filling curved three-dimensional space. The symbol reads as a chain. The 5 means each face of a cell is a pentagon, the 3 means three pentagons meet at each corner of a cell, which assembles a dodecahedron, and the 4 means four dodecahedra meet around every edge. Each dodecahedral cell has exactly 12 neighbors, one across each of its twelve pentagonal faces, and this is the central number of the whole theory (P85, P104).




Figure 3: The heptagrid, the regular tiling of the hyperbolic plane by seven-sided cells, three to a corner, the simplest face of the crystal. Each cell is a vibe and each shared edge is a perceived relation. Cells appear to shrink toward the rim only because the Poincare disk compresses the infinite plane into a finite picture. In the true geometry every cell is the same size.

You cannot fully picture this crystal, and that is fine. It is the exact three-dimensional analogue of the heptagrid. Everything true of the heptagrid is true of the dodecagrid, with regular identical cells, no preferred direction, exponential room, and the same kaleidoscope of reflections. The build is exact to tens of millions of cells by an integer fingerprint method that avoids any floating-point drift (P104).

The heptagrid is the version you can see. The dodecagrid is the real substrate. Time is its growth, the birth of new cells at the frontier, generation after generation, with each new layer the next moment.

6.6. Why three dimensions

The dimension of space is not a dial in this theory. It is forced, and forced twice over, then narrowed further to the specific crystal. The first squeeze is pure geometry. A crystal here is a regular tiling of hyperbolic space with finite cells, a compact regular hyperbolic honeycomb. Such a honeycomb exists only when both its cell and its vertex figure are themselves finite, which is a sharp constraint on the symbol.

Experiment P62 enumerates every such symbol by Gram-matrix signatures and finds a hard pattern. The hyperbolic plane admits infinitely many. Three-dimensional hyperbolic space admits exactly four, one of which is the dodecagrid $\{5, 3, 4\}$. Four-dimensional hyperbolic space admits exactly five. Five dimensions and above admit none, ever. The window of nonzero dimensions is precisely two, three, and four. A finite-celled crystal substrate simply cannot be built above four dimensions, so there is no room for hidden spatial dimensions to hide.

The second squeeze is physics and picks one of the three survivors. For stable matter to exist, orbits must close and retrace themselves, otherwise nothing stays bound over time. Whether an orbit closes depends on how force falls off with distance, which depends on the dimension. Experiment P68 integrates these orbits




Figure 4: The dodecagrid $\{5, 3, 4\}$, the three-dimensional substrate, dodecahedral cells with four meeting around every edge, filling curved three-dimensional space. It is the heptagrid one dimension up, the crystal of which physical space is made. Time is its growth.

across the allowed window and measures the apsidal angle, the angle swept from one closest approach to the next.

In two dimensions the orbit precesses and never quite closes, too loose to hold structure. In three dimensions the inverse-square law gives an apsidal angle of exactly 180 degrees, so the orbit retraces the same ellipse, the world of stable atoms and solar systems. In four dimensions and above the orbit is unstable, spiraling in or flying off at the slightest nudge.

Geometry allows two, three, or four. Physics allows only three. The intersection is a single number. A confirming measurement grows a crystal from scratch and reads back its intrinsic dimension from connectivity alone, returning values near the targets without bias (P69).

The two squeezes fix the dimension at three, but three-dimensional hyperbolic space still has four compact crystals. The auto-selection result shows $\{5, 3, 4\}$ is derived rather than chosen from that menu (P86). The ternary tone of each vibe makes the cell vertices three-valent, which is the middle entry $q = 3$. The finite three-valent cells are exactly the tetrahedron $\{3, 3\}$, the cube $\{4, 3\}$, and the dodecahedron $\{5, 3\}$.

Dimension	Compact hyperbolic combs	regular honey-	Examples
2	infinitely many		all $\{p, q\}$ with $1/p + 1/q < 1/2$
3	four		$\{3, 5, 3\}$, $\{4, 3, 5\}$, $\{5, 3, 4\}$, $\{5, 3, 5\}$
4	five		$\{3, 3, 3, 5\}$, $\{4, 3, 3, 5\}$, $\{5, 3, 3, 3\}$, $\{5, 3, 3, 4\}$, $\{5, 3, 3, 5\}$
5 and above	none		none

Table 5: The dimension window. Compact regular hyperbolic crystals exist only in dimensions 2, 3, and 4, with the dodecagrid $\{5, 3, 4\}$ the chosen one.

Eternal growth requires a genuinely hyperbolic crystal that does not close into a finite ball, and the theory takes the tightest such crystal, the one with the fewest cells per edge. The tetrahedron reaches no compact crystal at all. The cube needs $r = 5$. The dodecahedron reaches a compact hyperbolic crystal at $r = 4$, the tightest closure across all three-valent candidates. The forced symbol is exactly $\{5, 3, 4\}$.

The golden ratio that builds the dodecahedron is not smuggled in. The five falls out of the requirement to be the tightest eternal crystal with three-valent cells. The crystal at the root of reality is the unique survivor of constraints the theory committed to for other reasons, and the dimension is a result rather than an input.

7. Addressing, Growth, and the Exact Engine

The crystal is an infinite object. To compute on it we have to name its cells, count how it expands, and build it in memory without error. These three demands turn out to be one demand seen from three angles. The same Fibonacci and golden-ratio structure that grows the crystal also addresses it, and the same exact integer arithmetic that addresses it also builds it.

This section develops the golden ratio that pervades the substrate, the Fibonacci addressing that names every cell, the graph the crystal becomes once geometry is stripped away, the integer ladder that fixes the whole construction without free parameters, the growth series that counts the expansion exactly, and the modular engine that realizes the crystal to tens of millions of cells with no drift.

7.1. The golden ratio and the Fibonacci numbers

The Fibonacci sequence begins 1, 1, 2, 3, 5, 8, 13, 21, $\dots$, each term the sum of the two before it, $F(n) = F(n-1) + F(n-2)$. The ratios of successive terms converge on a single constant, the golden ratio $\phi \approx 1.6180339887$, the number satisfying

$$\phi^2 = \phi + 1. \quad (7)$$

Squaring ϕ is the same as adding one to it. Dividing through gives $\phi = 1 + 1/\phi$, so ϕ is one plus its own reciprocal. Its continued fraction is the all-ones expansion $1 + 1/(1 + 1/(1 + \dots))$, which is the slowest-converging of all continued fractions.

In the precise sense of continued-fraction approximation, ϕ is the most irrational number. No simple fraction ever gets a good grip on it, which is why packing at golden-ratio angles never lines up with itself and never leaves large gaps.

The golden ratio is not bolted onto the theory. It arrives in the $\{5, 3, 4\}$ substrate from three independent directions, all verified together (P50).

- First, as the fixed point of the self-referential equation $x = 1 + 1/x$, reached by pure iteration.
- Second, as the limiting growth ratio of the Fibonacci recurrence, a property of a growth process rather than an equation.
- Third, as the plane geometry of the pentagon, whose diagonal over its side equals $2 \cos(\pi/5) = \phi$.

Arithmetic, a growth recurrence, and plane geometry each give ϕ , and the three agree to within one part in a million. When a single constant arrives by three roads that share no machinery, this is the signature of one unified object rather than a coincidence.

The third road explains why ϕ must appear at all. The leading 5 in $\{5, 3, 4\}$ says the faces of every cell are regular pentagons, and the pentagon and ϕ are the same fact wearing two hats. To choose a pentagonal substrate is already to choose ϕ , whether or not one intends to.

7.2. Fibonacci addressing and greedy routing

For a signal to cross the crystal, every cell needs a name from which it can compute its own neighbors with no global map. Margenstern's splitting method supplies exactly this [19]. Pick any cell as the root. Its immediate neighbors form ring one, their further neighbors ring two, and so on.

Following each cell one step back toward the root traces a spanning tree, a skeleton that touches every cell exactly once. Hyperbolic geometry is already tree-like, since the number of cells at distance r grows like a constant to the power r , so the spanning tree reads off a structure the geometry already carries rather than imposing one.

The splitting method decomposes a wedge of the tiling into a leading tile plus smaller copies of a few standard region types, recursing forever. That recursion is the tree, and counting new tiles per ring yields the Fibonacci sequence $1, 2, 3, 5, 8, 13, \dots$. Number the cells ring by ring as the tree is walked outward, then rewrite each number in the Fibonacci base.

By Zeckendorf's theorem every positive integer is a unique sum of non-adjacent Fibonacci numbers. Writing that sum as a string of digits over the place values $1, 2, 3, 5, 8, 13, \dots$ gives a clean, collision-free address for every cell. The first twelve are 1, 10, 100, 101, 1000, 1001, 1010, 10000, 10001, 10010, 10100, 10101, as Table 6 lays out. No address ever has two ones in a row, because no integer ever needs two adjacent Fibonacci numbers. The set of valid addresses is therefore a regular language, recognizable by a fixed-size machine with no global memory. A tiny automaton inside every cell can navigate an infinite crystal.

The address is a turn-by-turn route from the root. To find a cell's parent, strip one digit. To find its children, extend by one legal step. To route from cell A to cell B, compare the two addresses from the left, walk up

Cell number	Fibonacci sum	Address
1	1	1
2	2	10
3	3	100
4	$3 + 1$	101
5	5	1000
6	$5 + 1$	1001
7	$5 + 2$	1010
8	8	10000
9	$8 + 1$	10001
10	$8 + 2$	10010
11	$8 + 3$	10100
12	$8 + 3 + 1$	10101

Table 6: The Zeckendorf addresses of the first twelve cells. Each cell number is the unique sum of non-adjacent Fibonacci numbers, and no address has two adjacent ones.

A’s route to the shared common ancestor by dropping digits, then walk down B’s route by appending its remaining digits.

Consider sending a message from cell 12 at address 10101 to cell 9 at address 10001. Comparing from the left, 1 matches, 0 matches, then cell 12 has 1 where cell 9 has 0, so the paths split. The shared prefix is 10, which is the address of cell 2, their common ancestor. Walking up from 12 by dropping one digit at a time gives $10101 \rightarrow 1010$ (cell 7) $\rightarrow 101$ (cell 4) $\rightarrow 10$ (cell 2). Walking down to 9 by appending its remaining digits gives $10 \rightarrow 100$ (cell 3) $\rightarrow 1000$ (cell 5) $\rightarrow 10001$ (cell 9). The full path is 12, 7, 4, 2, 3, 5, 9, six hops, and every cell on it only compared two short strings. Each address in this route can be read straight from Table 6.

No cell ever consulted a map of the crystal. On the 2D heptagrid this exact addressed routing delivered 100 percent of random pairs at low stretch (P42). On the real 3D dodecagrid, greedy hyperbolic routing, where each cell forwards to the neighbor whose address is closest to the target, delivered over 90 percent of pairs at stretch under 2 (P76). The addresses are not stored labels. They are implicit in each cell’s place in the web and can be computed from local structure whenever needed.

7.3. The crystal as a graph

Strip away the drawing, the cell shapes, and the curved-space picture, and what remains is a graph. Each cell is a vertex, each shared face an edge. A dodecahedron has twelve faces, so in the interior every cell touches exactly twelve others. The graph is 12-regular. Distance is step-counting, the fewest edges between two cells, and this step-counting distance is what space will mean.

Because the crystal lives in negatively curved space, the graph is an expander, with no bottlenecks. Any region not larger than half the graph has a boundary that is a fixed fraction of its size, so the graph cannot be cut into two large pieces by snipping few edges. This is the same exponential ball growth seen as connectivity. The graph is also tree-like, since the exponential branching that makes it an expander is the exponential ball growth that makes it tree-like, two views of one fact.

Together these force a tiny diameter, logarithmic in the cell count. On roughly 120,000 exact $\{5, 3, 4\}$ cells the worst-case separation is about six hops through the bulk, close to $\log_B N$ with growth base $B \approx 7.87$

(P111). Among a hundred thousand cells, no two are more than about six steps apart. That tiny $O(\log N)$ diameter is the engine behind fast global integration, against the $N^{1/3}$ crawl a flat cubic lattice would require.

7.4. The integer ladder

The entire substrate is built from the plain whole numbers with no free parameters at any step. One picks the integers, which cannot be picked differently, and the crystal $\{5, 3, 4\}$, which is itself forced. Everything else follows by deterministic construction. The ladder runs upward through four rungs, each verified (P48, P104, P51, P83).

Asking the integers for their symmetries returns the modular group $\text{PSL}(2, \mathbb{Z})$, the integer matrices of determinant one, with nothing to tune. This parameter-free modular base tessellates hyperbolic space Lorentz-safely, is addressed by a deterministic finite automaton through continued fractions, and carries the golden ratio as its central geodesic, the all-ones continued fraction whose convergents are the Fibonacci ratios (P48).

The next rung represents the exact dodecahedral cells of $\{5, 3, 4\}$ in pure integer arithmetic (P104). Above it the whole tower runs as one continuous program, from integer generator data up to running physics, through a single shared pipeline for three different bases (P51). The top rung grows the crystal one cell at a time by a fixed deterministic rule, resumable, append-only, and faithful to the static tiling ring for ring, with the hyperbolic curvature and the golden ratio emerging on their own rather than being inserted (P83).

The integers are the bones, the small ones carry the laws, and there are essentially no free parameters anywhere in the foundation.

7.5. The growth series

How fast does the crystal expand? Stand on one cell and count the cells one step out, two steps out, three steps out. The growth function $f(n)$ records the number of cells added at ring n . For the dodecagrid the first counts are 1, 12, 102, 812, 6402, 50412, $\dots$

These counts are fixed exactly by the same mirror data that defines the crystal. Steinberg's formula computes them by packing the sequence into a generating function $W(t) = \sum_n f(n) t^n$ and expressing its reciprocal as an alternating sum over the finite parabolic subgroups,

$$\frac{1}{W(t)} = \sum_{\text{finite } T} \frac{(-1)^{|T|}}{f_T(t)}, \quad (8)$$

where each $f_T(t)$ is the Poincare growth polynomial of a finite reflection subgroup. The whole computation runs in exact integer and rational arithmetic over polynomials, with no rounding.

The denominator of $W(t)$ gives a linear recurrence. For $\{5, 3, 4\}$ the recurrence is

$$f(n) = 9f(n-1) - 9f(n-2) + f(n-3), \quad (9)$$

seeded by 1, 12, 102, 812. Verifying on the seeds, the next term is $f(4) = 9(812) - 9(102) + 12 = 7308 - 918 +$

$12 = 6402$, and the one after is $f(5) = 9(6402) - 9(812) + 102 = 50412$, both matching the series exactly. The denominator $1 - 9t + 9t^2 - t^3$ has characteristic polynomial $x^3 - 9x^2 + 9x - 1 = (x - 1)(x^2 - 8x + 1)$, with roots 1 , $4 + \sqrt{15}$, and $4 - \sqrt{15}$. The growth base is

$$B = 4 + \sqrt{15} \approx 7.873. \quad (10)$$

The reciprocal root $4 - \sqrt{15} \approx 0.127$ lies strictly inside the unit circle, which is the proof of exponential growth. Each ring is about 7.87 times the size of the one before it, and that is the universe's expansion rate in cells. A direct breadth-first count returns the same 1, 12, 102, 812, 6402, 50412, with the ring-to-ring ratio climbing 8.50, 7.96, 7.88, 7.87 toward $4 + \sqrt{15}$.

Table 7 lists the first ten rings of the two-dimensional heptagrid $\{7, 3\}$ and the three-dimensional dodecagrid $\{5, 3, 4\}$ side by side. The heptagrid grows at the golden ratio squared, about 2.618 per ring, and the dodecagrid grows at $4 + \sqrt{15}$, about 7.87 per ring. The dodecagrid reaches more than a billion cells by the tenth ring.

Ring	Heptagrid $\{7, 3\}$		Dodecagrid $\{5, 3, 4\}$	
	Added	Cumulative	Added	Cumulative
0	1	1	1	1
1	7	8	12	13
2	21	29	102	115
3	56	85	812	927
4	147	232	6,402	7,329
5	385	617	50,412	57,741
6	1,008	1,625	396,902	454,643
7	2,639	4,264	3,124,812	3,579,455
8	6,909	11,173	24,601,602	28,181,057
9	18,088	29,261	193,688,012	221,869,069
10	47,355	76,616	1,524,902,502	1,746,771,571

Table 7: Cells added and cumulative total per ring for the two crystals. The heptagrid grows by about 2.618 per ring (the golden ratio squared), the dodecagrid by about 7.87 (the root $4 + \sqrt{15}$).

Two distinct sequences must be kept apart. The cell-shell growth above counts honeycomb cells at each distance from a seed cell. The Coxeter group growth, which Steinberg's formula also produces, counts group elements by word length, giving 1, 4, 9, 17, 29, 46, ... for $\{5, 3, 4\}$, a different object. The cell-shell series is the physically meaningful one.

The two-dimensional warm-ups $\{7, 3\}$ and $\{5, 4\}$ share the denominator $1 - 3t + t^2$ and grow at base $(3 + \sqrt{5})/2 \approx 2.618 = \phi^2$, exactly the square of the golden ratio, so the golden thread runs through the smaller cases as well.

7.6. The exact modular-fingerprint engine

Building the crystal by float coordinates fails at a hard precision wall. Tracking each cell by its center as $g \cdot c_0$ projected into the Poincare ball works only while distinct cell centers stay further apart than a 64-bit float can resolve. Hyperbolic space grows exponentially and cells crowd toward the boundary of the ball, so

their coordinates pile up with gaps shrinking exponentially with distance. In practice the float engine caps out near 15,500 cells. This is not a coding flaw but the geometry.

The escape is to notice that the reflection matrices of $\{5, 3, 4\}$ are made of a small set of exact algebraic numbers, with the only irrational ingredients being $\sqrt{5}$ inside ϕ and $\sqrt{2}$. A finite field of integers modulo a prime p contains both square roots exactly when $p \equiv 1 \pmod{40}$, since then 2 and 5 are both quadratic residues. The roots are computed by Tonelli-Shanks, ϕ is built from $\sqrt{5}$, and every dodecahedral reflection becomes exact integer arithmetic with zero rounding.

A single prime carries a tiny risk that two distinct cells share one residue, so the engine uses two primes at once, near 67,108,837 and near 66,000,000, both congruent to 1 mod 40. A cell's fingerprint is its center reduced modulo both primes, and two cells collide only if they collide modulo both at once, which by the Chinese Remainder Theorem means agreeing modulo a product above 4×10^{15} . The twelve face reflections are derived by conjugating the outer generator by each of the 120 elements of the dodecahedral cell stabilizer, so adjacency is read straight off the group rather than measured. A typed open-addressing hash on the packed fingerprints deduplicates cells with no allocation.

The result has been run offline to 20 million cells at a few seconds per million, every cell carrying exactly twelve facet neighbors throughout, verified against an independent float build below the wall and persisted to disk without loss (P104). For transport tests, which a tiny-diameter ball cannot support, the same exact arithmetic builds a sliver, a long thin geodesic tube with bounded width whose alternating face product is a hyperbolic translation, giving genuine distance to travel along a one-dimensional spine. The physics of the theory is therefore measured on the true crystal, not on a guess of it.

8. Lorentz Safety

Relativity is the hardest test a discrete theory of space has to pass. The deep fact of our universe is Lorentz invariance. The laws of physics look the same no matter how fast one moves, with no special rest frame and no preferred direction in space. This is the make-or-break filter for any theory that builds spacetime from a discrete web of relations, and most such theories fail it in a predictable way.

This section explains why a flat regular lattice cannot carry relativity, why a hyperbolic and effectively aperiodic substrate can, how the substrate passes a genuine boost test, how the observational bounds come out, and how the result reappears as an emergent symmetry of the dynamics.

8.1. The causal-set lesson

Causal-set theory taught the field a hard lesson. Build spacetime from a tidy flat lattice, integer points in neat rows and columns with causal links between them, and the construction is fatal. A flat lattice has preferred directions. Motion along a row is geometrically different from motion along a diagonal, and because the space is flat those directions are globally defined. North here is the same as north a mile away.

The grain is real and detectable everywhere at once, which gives the lattice a preferred rest frame, exactly what relativity forbids. The damage is concrete and measurable. On a lattice the speed of a signal depends on which way it travels, and the dependence grows with energy, since high-energy signals feel the fine grain

most. This is the energy-dependent, directional Lorentz violation measured directly in our work (P27). A flat substrate is not subtly wrong. It announces its preferred frame.

8.2. Why curvature buys Lorentz safety

The escape is not what intuition suggests. One would expect that hiding the grain requires a messy random space and that an exact ordered crystal would be the worst case. The opposite is true. The rescue is curvature, and an exact ordered hyperbolic crystal is Lorentz-safe. In flat space a direction can be carried unchanged across the whole space, so a flat grid’s preferred directions are globally real. In curved space it cannot. Curvature scrambles directions under parallel transport. Move across a hyperbolic space and the links that pointed one way locally now point every which way relative to a distant patch. From any local spot the crystal’s links fan out in all directions, statistically indistinguishable from a random foam with no preferred direction. The order is still there. It just does not line up across the space, so no resident can use it to define a global frame.

The experiments confirm this across a wide family of substrates. A fully deterministic substrate, the golden-angle sunflower, is as Lorentz-safe as a random sprinkle, with directional anisotropy in the same low band and no random number generator used anywhere (P40, see also P39). Determinism is fine. What matters is that the sequence is effectively aperiodic, never repeating a global direction.

A sweep across random sprinkles, the sunflower, a low-discrepancy sequence, and the regular $\{7, 3\}$ and $\{5, 4\}$ hyperbolic tilings finds every hyperbolic substrate Lorentz-safe with anisotropy below threshold, while the flat lattice control is the only one that fails (P40). Regularity does not break Lorentz invariance once the space is curved. The Margenstern families of hyperbolic tilings, both the $\{p, 4\}$ and $\{p, 3\}$ series, are all Lorentz-safe (P41). Lifting into three dimensions, the $\{5, 3, 4\}$ dodecagrid at the heart of the theory is Lorentz-safe while a flat cubic lattice is not (P45).

A natural worry is that the dodecagrid is a lucky special case hand-picked to pass. Feeding $\{7, 3\}$, $\{5, 4\}$, $\{8, 3\}$, $\{6, 4\}$, and the 3D $\{5, 3, 4\}$ through a single Coxeter generator, changing only the Schläfli symbol, confirms every one is Lorentz-safe (P47). These tilings are one machine on different settings, so Lorentz safety is a property of the whole reflection-group family, not a fluke of one member.

8.3. Boost-covariance on the substrate

Isotropy of directions is only the rotational half of Lorentz invariance. The other half is boost safety, invariance under changes of velocity, and it must be tested on its own. The boost parameter of a timelike causal link is its rapidity, the hyperbolic angle of the link in spacetime, and a real boost shifts every rapidity by the same amount because rapidities add. So a structure with no preferred frame has a rapidity distribution that is flat and broad, while a structure with a preferred frame has one peaked at rapidity zero where its timelike links pile up along a built-in time axis.

A Poisson sprinkle of $1 + 1$ Minkowski space gives a flat distribution with normalized entropy above 0.9, while the matched causal lattice gives a sharply peaked one with a spread several times smaller, betraying its rest frame (P84). The same experiment then performs an actual boost, applying a Lorentz transformation of rapidity one to every coordinate, leaving the causal order untouched, and recomputes the rapidities in the new frame.

The distribution shifts by exactly the boost, as it should, but its shape is unchanged. After removing the expected shift, the flatness of the boosted distribution matches the original to within a small tolerance. The sprinkle is boost-covariant. Its statistical description is the same in the boosted frame as in the rest frame, which is precisely what no preferred frame means. This is a genuine boost test, not a rotation proxy, and the substrate passes it while the lattice does not (P84).

8.4. Observational bounds

Lorentz safety is also an observational check, and the observation is real. Light from distant gamma-ray bursts travels billions of years before reaching us. If spacetime had a grain, high-energy and low-energy photons would travel at slightly different speeds and would drift apart over that immense baseline. They arrive neck and neck, ruling out any first-order, energy-dependent speed variation to extraordinary precision.

The model’s predicted linear-order Lorentz-violation coefficient goes to zero, comfortably inside the gamma-ray-burst bound, because the curved crystal predicts the grain simply is not there, while the same calculation on a flat lattice gives a large coefficient that the data excludes (P82). The residual causal-set swerve, a diffusion of particle momenta, also vanishes as the discreteness scale shrinks, so there is no leftover Lorentz-violating jitter (P82).

This cuts both ways. The model predicts no first-order grain in the speed of light, and if anyone ever measured a first-order direction- or energy-dependence of light’s speed, the curved-crystal picture would be falsified. The current data is on its side.

8.5. The link to the emergent dynamics

The substrate-level results show the foundation has no built-in frame. The full theory then has to grow relativity out of its dynamics, not just its geometry, and three modern results complete the chain.

Leading-order rotational isotropy emerges from the dynamics, with the one-step diffusion tensor on the $\{5, 3, 4\}$ having equal eigenvalues, so a walker spreads the same in every direction. This isotropy is inherited from the icosahedral symmetry of the dodecagrid’s cell, the geometric reason the lattice has no preferred spatial axis even though it is built from identical cells (P124). The deterministic reversible wave on the $\{5, 3, 4\}$ then has an isotropic speed, with anisotropy across the twelve directions consistent with zero, an emergent frame-independent light speed built from a real dynamical rule rather than from substrate statistics alone (P150).

The boost rung follows on the same dynamics. The wave’s dispersion is real, linear, and massless, $\omega = |k|$, a positive-norm massless mode with an exact lightcone that is boost-invariant, with massive modes boost-invariant in the infrared window below the cutoff (P154). Together with the rotations, this gives the full emergent Lorentz group acting on a genuine relativistic field.

The arc is complete. The substrate has no preferred frame (P27, P39, P40, P41, P45, P47, P84). The observations that kill a lattice are passed (P82). And relativity reappears as an emergent symmetry of the actual dynamics on the dodecagrid (P124, P150, P154).

The choice of a negatively curved substrate was made for several reasons at once. Curvature gives room to grow, it carries the right symmetries, and it removes the preferred frame that flat discreteness cannot escape.

Room, symmetry, and Lorentz invariance are not three separate gifts. They fall out of a single decision. Curvature, not disorder, buys Lorentz invariance, and the $\{5, 3, 4\}$ dodecagrid is exactly the curved, aperiodic substrate that relativity needs.

9. Computation

The theory rests on a single local rule applied across the whole crystal, beat after beat. For that to grow atoms, weather, and minds, the substrate must be able in principle to express any structure at all. The guarantee that secures this is computational universality. A system is universal when it can carry out any computation that can be carried out and simulate any other computer.

This section establishes that the $\{5, 3, 4\}$ dodecagrid has this property, at the level of an external theorem and at the level of the model's own running dynamics. It then draws the conclusion that the substrate is a computer rather than merely like one, marks the limit that universal computation is goal-neutral, and closes with the undecidability that comes packaged with the same power.

9.1. Cellular automata

A cellular automaton is a grid of identical cells, each holding one value from a small finite set, updating together on a shared clock by one fixed rule that reads only a cell and its immediate neighbors. The machine has four parts.

- The cells are a regular array of sites.
- The states are values from a small finite set, where two states already suffice for everything of interest.
- The local rule is a single lookup from a cell and its neighbors to that cell's next state, the same rule everywhere.
- The beats are synchronous, with the whole grid stepping at once.

There is no program counter, no central memory, and no operator. A quiescent state, the resting value such that a cell surrounded by quiescence stays quiescent, keeps a starting pattern finite so that only the live region acts.

The lasting surprise is that such rules are not dull. On the line, an elementary rule with two states and three neighbors has only 256 possibilities. One of them, Rule 110, was proved universal by Cook. In the plane, Conway's Game of Life turns a four-line neighbor-counting rule into gliders, guns, oscillators, and working calculators.

The complexity is not in the rule. It is in the patterns the rule allows. A law statable in one breath, applied locally an enormous number of times, produces moving objects, memory, and full computation [33]. This is the permission slip for the theory. If a trivial rule on a flat board yields gliders and calculators, a well-chosen rule on a vast curved crystal can yield particles, forces, and minds. A complex world does not require a complex law.

The deepest lesson concerns what counts as a thing. A glider in the Game of Life is not a fixed set of cells. It is a pattern that moves, re-forming one step over at each beat on different cells while keeping its shape. In the theory every thing is this kind of object. Every particle and every self is a stable traveling pattern the rule sustains, made of nothing but arrangement, while the cells beneath it come and go. This is why a self can persist while none of its underlying cells do.

9.2. The perception rule as a cellular automaton

The model's dynamics is a cellular automaton on the nose, not by analogy. The cells are the vibes, the sites of the $\{5, 3, 4\}$ dodecagrid, each with twelve touching neighbors. The states are the three tones $+1$, 0 , and -1 . The local rule is the perception rule. For a pair of neighboring tones it shares opposite tones to peace, hops a charge into an empty neighbor, spawns a balanced pair from peace under the arrow, and leaves same tones inert. The beat sweeps the edges once with a no-double-move guard, which is exactly the synchronous tick. The quiescent state is peace, the tone 0 , doing nothing until touched. The arrow is the only asymmetry, the creation bias that keeps the grid alive rather than relaxing to dead peace.

This is not a flat 256-rule line and it is not the Game of Life. It is a three-state, twelve-neighbor cellular automaton on a three-dimensional hyperbolic honeycomb. It is the same kind of object, and it inherits the same lesson, that simple uniform local rules can produce open-ended complexity. The neighborhood of a cell, the thing the rule reads, is defined through Margenstern's recursive splitting of the hyperbolic tiling into a tree of nested regions addressed by Fibonacci arithmetic, which is what turns an undrawable infinite plane into a navigable, computing object [20].

Calling the universe a cellular automaton does not mean it runs inside an outside machine. There is no operator and no screen. The substrate is the automaton, and running the rule on the substrate is what existence does.

9.3. Universality on hyperbolic tilings

The strongest external result is Margenstern's. Cellular automata on hyperbolic tilings, the dodecagrid $\{5, 3, 4\}$ included, are computationally universal [19]. This is a rigorous construction, and the $\{5, 3, 4\}$ is the exact mesh of the theory, so the theorem lands directly on the substrate.

Building a universal machine in hyperbolic space faces a special obstacle. There is no uniform zoom. Hyperbolic space has no similarity transformations, so the usual trick of nesting a shrunken copy of a machine is unavailable. Margenstern works around this with the railway model.

Picture a computer made of trains. A single locomotive runs forever along the tracks. At junctions sit switches that can be flipped, and the configuration of all switches encodes the state of a register machine. As the locomotive passes through it reads switches and flips them, and that reading and flipping is the computation. Registers are stored as track loops the train can lengthen or shorten. A train confined to a track cannot drift apart the way free-flying signals do when hyperbolic lines diverge, so the railway keeps the computation literally on the rails.

Lifted into hyperbolic three-space, the dodecahedral tiles each have twelve neighbors and can route signals over bridges the plane lacks, which shrinks the machine to five states, against nine for the two-dimensional

pentagrid version. A five-state weakly universal automaton therefore lives on the exact mesh. Weakly universal means it computes on an infinite ultimately periodic background rather than on a blank quiet plane, which is the standard accepted sense for hyperbolic constructions.

9.4. Universality on the live dynamics

An external theorem is not enough. The theory needs universality on the model’s own rule. The experiment P44 supplies it. The local update is the ternary signed-majority rule, where a cell takes the sign of the sum of its fill-weighted neighbor tones plus a bias. With both fills set to -1 and the bias set to $+1$, that update is exactly NAND. Encode **true** as $+1$ and **false** as -1 . The cell then computes $\text{out} = \text{sign}(1 - a - b)$. The four input cases are as follows.

a	b	$1 - a - b$	$\text{out} = \text{sign}(\cdot)$	NAND
$+1$	$+1$	-1	-1 (false)	T NAND T = F
$+1$	-1	$+1$	$+1$ (true)	T NAND F = T
-1	$+1$	$+1$	$+1$ (true)	F NAND T = T
-1	-1	$+3$	$+1$ (true)	F NAND F = T

The output is false only when both inputs are true, which is exactly NAND. NAND is functionally complete, so every Boolean function follows, and P44 builds NOT, AND, OR, XOR, and a full adder. The adder computes $a + b + \text{carry}$ correctly across all eight inputs, so real arithmetic runs on the rule. P44 then reproduces the table of Rule 110 from rule-NANDs and evolves it non-trivially, tying the dynamics to a known-universal system expressed entirely in its own gate.

The decisive part is stronger. The gates are not left as functions. Each is built as a real subgraph of cells with symmetric ternary fills, with inputs and a positive bias clamped. The free cells run the model’s asynchronous signed-majority rule to a fixed point, and the output cells are read. NAND built this way settles to the correct answer, and a six-gate XOR composed as $\text{AND}(\text{OR}(a, b), \text{NAND}(a, b))$ also settles correctly, since bus widths decrease downstream so each gate’s margin beats the feedback from the next layer.

Two worries are answered at once. The gates are run by the live dynamics, not painted on as a separate circuit, and every output is a genuine fixed point stable under any update order. A clamped-input gate is an attractor, not a decaying pulse, which defuses the concern that an asynchronous, dissipative dynamics could not hold a stable logical value.

The experiment P141 closes the loop with a full Turing-universal machine running directly on the mesh. The target is a two-counter Minsky machine, one of the smallest known universal models. The mapping uses only the substrate’s primitives. A register is a region of cells in the dodecagrid, and its value is the number of charged cells in that region. Increment is creating a charge, one unit of the arrow. Decrement is annihilating one charge. Test-zero is counting the charge.

P141 loads a multiplication program built from increment and decrement-or-jump-if-zero, which is a universal instruction set, and computes $R_2 = R_0 \times R_1$ with a temporary register preserving the multiplicand. The machine computes $a \times b$ correctly across a spread of cases including a zero case, with registers and operations made of nothing but tone and the arrow.

9.5. The substrate is a computer

Three independent angles converge.

- Margenstern proves the $\{5, 3, 4\}$ class universal as a theorem.
- P44 shows functional completeness on the live dynamics, with NAND and a six-gate XOR settling as stable fixed points.
- P141 runs a full universal machine with charge as memory.

The conclusion is not that the universe is like a computer. It is that the universe is one.

This converts a wild claim into a modest one. The wild claim is that a single rule on a crystal produces atoms, galaxies, weather, and minds. The modest fact behind it is that the rule is universal, so any pattern that can exist at all is a pattern the rule can host. If a working mind is a pattern a computer could in principle run, the crystal can host it, because the crystal can run anything a computer can.

Universality is about capacity, not freedom of outcome. The rule can in principle express anything, while the actual universe is fixed completely by the one rule, the single seed, and the deterministic growth that follows. The capacity is unlimited and the outcome is one.

A clean boundary remains. Universality settles structure, never feeling. The theory treats experience as the substance of the cells from the start, rather than as a product of computation, so the felt interior is its own question and is not hidden inside the universality claim.

9.6. Computation is goal-neutral

Universality is necessary but not sufficient for mind. A universal computer can compute any function, but it has no preference among them. It will compute a multiplication, a heat equation, or pure noise with equal indifference. Universal computation is goal-neutral. It is raw capacity with no direction.

Intelligence needs direction. A problem is a gap between a desired state and a current state, and a solution is a path that closes it. Solving a problem needs both the means, which is universal computation, and the ends, which is a goal that picks one path over another. The substrate supplies the ends through the arrow, the directional asymmetry the conserving rule selects toward.

So P44 and P141 establish the means only. They show the substrate can compute anything. The bridge from computable to intentional is taken up in the section on intelligence, which adds the arrow and the goal-directed search that aims this universal computation at a target.

9.7. Undecidability and the super-Turing question

Power and limitation arrive together. The moment a system can simulate a universal machine, the halting problem for that system becomes undecidable, since any Turing-complete substrate drags undecidability in with it. This is not a flaw. It is the formal signature of a world where the only way to know the future is to live it. The undecidability concerns the observer's predictive shortcut, not the system's evolution. The rule

still produces a perfectly definite future. There is simply no algorithm that fast-forwards to the answer in general.

The cleanest undecidable problem living natively on the substrate is the domino problem. Given a finite set of tile types with edge-matching rules, can copies cover the whole plane with no gaps and no overlaps, every edge matching its neighbor. This is already undecidable in the Euclidean plane, proved by Berger in 1966. The hyperbolic version stayed open for nearly three decades until Margenstern proved, in early 2007, that the domino problem is undecidable in the hyperbolic plane, with Kari proving the same result independently by a different method within days [19].

The proof has the classic shape. Force the tiling to draw a computer, then ask whether the computer halts. Inside the heptagrid $\{7, 3\}$ the matching rules force a rigid background, and on that scaffold a hierarchy of nested triangles forms a grid that becomes the space-and-time diagram of a Turing machine. The tiles complete into a full tiling if and only if the machine never halts. Since halting is undecidable, so is tiling.

Because the substrate is a tiling, asking which local-matching patterns extend to cover the whole grid is the domino problem in disguise, so undecidability is a native feature of the geometry the theory is built on rather than an import from outside.

The frontier question is whether a hyperbolic substrate could ever compute past the Turing barrier. Margenstern and Grigorieff sketch a model that does, on an infinigrd, a tiling by polygons with infinitely many sides. One such cell has infinitely many neighbors, so in one update it can test a property across infinitely many cells at once, evaluating a quantifier in a single tick. Chaining such cells decides every level of the arithmetical hierarchy, which would be super-Turing computation.

The caveats are decisive and kept clearly speculative here. The model uses infinite objects, a tile with infinitely many sides and a cell reading infinitely many neighbors instantaneously, and such a thing cannot be built. It belongs to the same idealized family as infinite-time Turing machines and black-hole computation, with rigorous internal mathematics and no physical realizability [28]. The dodecagrid has cells with twelve neighbors, not infinitely many, so the infinigrd result does not apply to the substrate as built. The right reading is a clean theorem on undecidability and a clearly labeled thought experiment on hypercomputation, never the two confused.

Part III

Emergence

10. Emergent Physics

When the single update rule of the dynamics runs on the hyperbolic mesh, a recognizable physics appears in the measurements. We do not assume a field and then quantize it. We run a local perception rule on a crystal and read off the structures that a field theorist would demand. A fluctuating vacuum with a mass

and a conserved current. An effective gravity that bends rays. A horizon that radiates. A holographic area law. A causal lightcone. And, at the deep end, a relativistic and reflection-positive Dirac field. Every number reported here is the output of an experiment run on the exact substrate, cited inline in the form (P114).

The organizing claim of this section can be stated up front so the rest reads as evidence for it. The stochastic version of the rule, the one whose pair creation is implemented with a random coin, yields a massive and diffusive field. That field is real, but it is the classical decohered shadow of something deeper. The deterministic, reversible, unitary completion of the same rule, with the randomness removed and nothing added, is a genuine relativistic and reflection-positive quantum field. Relativity and quantumness are not extra ingredients. They are what remains once the dice are taken away.

10.1. The Field-Theoretic Vacuum

Run the rule on a sea of inactive cells and let it settle. The steady state is not still. It is a fluctuating vacuum of virtual particle-antiparticle pairs. The mechanism is the rule itself. The creation move turns an inactive pair $(0, 0)$ into a balanced charge pair $(+1, -1)$. The share move returns $(+1, -1)$ to $(0, 0)$, annihilating it. Creation followed by annihilation, repeated on a conserved charge, is vacuum pair production realized directly by the dynamics.

The measured signatures match this reading (P114). The vacuum carries a fluctuating charge density of about 28 percent, never zero and never saturated, a living equilibrium. The two-point correlator shows nearest-neighbor pair anti-correlation, with $C(1)$ measured negative, the fingerprint of a $+1$ sitting beside the -1 partner minted with it. That correlator decays with a finite correlation length, and in field language a finite correlation length is an effective mass.

The vacuum is gapped. It costs energy to make an excitation, and that cost sets the range of influence. Propagation is bounded as well. There is a causal lightcone with a cone speed of order one cell per beat, so no signal outruns the substrate clock. The total charge is conserved, a $U(1)$ current that never leaks.

What this predicts is a vacuum that is never empty, that seethes with balanced fluctuations, and whose fluctuations are short ranged and causal. What it explains is the origin of a quantum vacuum with no quantum postulate. The vacuum is the rule's equilibrium. The intensive structure survives coarse-graining. The effective density, the value of $C(1)$, and the cone speed measured on a small slice match the large field (P115), so a measurement on a slice can be trusted across scales rather than merely hoped to hold.

10.2. Effective Gravity, Horizons, and Dimension

Matter is not separate from the mesh. It is dense local activity, and dense activity changes the rule's local rates. That single fact produces gravity.

Where matter sits, it strengthens the local fills, and the strengthened fills act as a varying index of refraction for the substrate's own signals. Rays bend toward matter, which is gravitational lensing (P88). The deflection scales with the mass, so doubling the mass produces roughly the larger bend expected, and it weakens with impact distance, so rays passing far bend less than rays passing near. The curvature peaks where the mass is. Mass curves the effective geometry and light follows the curve. The metric is emergent rather than imposed.

Sculpt a fill profile so that signals slow toward a surface and the result is an effective horizon, a one-way membrane on the crystal (P89). A ray climbing out of it redshifts, and the redshift defines a surface gravity κ . Two independent measurements of κ , one from the effective metric and one from the ray trajectory, agree. A detector held outside the horizon registers a thermal spectrum at the Hawking temperature

$$T_H = \frac{\kappa}{2\pi}, \quad (11)$$

and that reading matches the detailed-balance temperature read off the rule’s own rates. The temperature scales with the surface gravity as expected. The substrate radiates from a horizon for the same thermodynamic reason a real black hole does, an analog Hawking effect built from the local update alone.

The rule also distinguishes a bare three-dimensional space from a brane embedded in a higher bulk (P90). A bare 3D substrate gives an inverse-square law at every scale, a clean exponent of -2 . A 3D brane inside a 4D bulk crosses over to an exponent of -3 at short range, the signature of force lines leaking into the extra dimension before folding back to inverse-square far away. The two scenarios are measurably distinct, with a crossover scale set by the bulk thickness. The theory not only permits a short-range gravity deviation, it predicts the shape that deviation would take.

10.3. Holography and the Field Beneath

Hyperbolic space is holographic by its shape, and the rule inherits that property.

The crystal is a spatial slice of an anti-de Sitter geometry, and on it an entanglement-entropy-like quantity for a boundary region obeys the Ryu-Takayanagi log law (P91),

$$S = m \log \sin \left(\frac{\theta}{2} \right), \quad (12)$$

with a fit quality of R^2 near 1.0. The entropy of a region is set by the length of the bulk geodesic that subtends it, and that geodesic is a genuine shortcut, many times shorter than any path along the boundary. Depth into the bulk plays the role of scale. This is the area law that ties entanglement to geometry, present in the substrate without being put in by hand [25, 18]. It is the same relation that underlies the proposal that geometry is built from entanglement [30], now read off a discrete rule rather than postulated.

The same structure connects to holographic quantum error correction. A bulk degree of freedom is recorded redundantly across boundary regions, so erasing part of the boundary does not erase the bulk information. This is the logic of holographic codes on hyperbolic tilings [21, 13]. It is also why memories and selves stored on this mesh are robust. They are encoded holographically and redundantly, an error-correcting code that survives partial erasure.

The bulk is not only deep, it is shallow in graph distance. The hyperbolic mesh has a tiny diameter (P111). For about one hundred thousand cells the eccentricity is roughly six hops, growing only as the logarithm of the cell count, $O(\log N)$. The whole structure is a few layers deep. A signal need not crawl across the surface. It dives into the bulk and crosses the entire structure almost instantly, reaching a far region that a

surface-only walk could not find in time. This radial bulk direction is the engine of fast global integration. The hyperbolic curvature provides small-world wiring for free (P111).

The holographic interior is the curved bulk, but ordinary directed physics lives on a flat sheet inside it. That sheet is a horosphere, the Busemann level set that is intrinsically Euclidean even though it sits within the negatively curved bulk, and it is flat by measurement (P142). The hyperbolic bulk is the holographic interior where the area law and the bulk shortcut live, and the horosphere is the emergent flat physical layer on which the vacuum, the lightcone, and ordinary space appear. So holography, the field beneath, and the flat layer are one geometry seen at three depths, the curved holographic interior, the radial shortcut threading it, and the flat horosphere where directed physics runs.

10.4. The Vibe-to-Field Ladder

To go from a discrete rule to a relativistic quantum field one must climb a ladder of separate requirements. Each rung is its own measurement.

- *A lightcone with ballistic transport (P123)*. Drop a charge on a long thin sliver, a geodesic tube whose spine runs a hundred cells or more, the right tool to defeat the small graph diameter. The charge does not diffuse. It moves ballistically, with a spreading exponent near 1.0 and a finite escape speed of order one cell per beat. Ballistic spreading is a finite propagation speed, and a finite propagation speed is a lightcone, the non-negotiable relativistic ingredient. The result is robust across tube widths, a substrate version of the hyperbolic rate-of-escape theorem.
- *Rotational isotropy (P124)*. The spatial half of Lorentz symmetry is rotational invariance, and the crystal supplies it. The one-step diffusion tensor has equal eigenvalues, with a measured anisotropy near zero. The reason is geometric. The twelve neighbors of a cell sit at the vertices of an icosahedron, and the icosahedral group acts irreducibly on three-dimensional space, which forces any rank-two tensor commuting with it to be a multiple of the identity. Space looks the same in every direction by the cell's own symmetry, not by tuning.
- *Lorentz restoration at higher harmonics (P125)*. Isotropy at one step is necessary but not sufficient. The rank-4 and rank-6 anisotropies must also be tamed. Under coarse-graining the higher-order anisotropy shrinks toward zero by central-limit washout, while the rank-2 piece already sits at the floor. Rotational invariance is therefore not just present but improves as one zooms out, exactly the behavior an emergent symmetry should show.
- *Reversibility, the precondition for quantization (P126, P128)*. A unitary quantum theory requires a reversible classical theory beneath it. The rule satisfies detailed balance both locally and around loops. Locally, single-edge transition counts are symmetric at all rates, with the violation at the noise floor near a ratio of one (P126). Around closed loops there is no persistent charge circulation, which is Kolmogorov's criterion, with mean circulation at the floor (P128). The cause is that creation makes balanced pairs, a reversible reaction rather than a one-way drive. The dynamics is a genuine equilibrium process, which is the precondition for the next step.
- *A critical point yielding a free-field continuum (P129)*. A generic lattice coarse-grains to nothing. To leave a real continuum field behind, the base must sit at a critical point. The order parameter, the

activity, vanishes as the square root of the creation rate,

$$\rho \sim \sqrt{r}, \quad \beta \approx 0.5, \quad (13)$$

with a fit R^2 above 0.9. This is the mean-field rate law, expected because the hyperbolic graph lives above the upper critical dimension. There is an absorbing critical point as the creation rate goes to zero, and its continuum limit is a mean-field Gaussian free field carrying a $U(1)$ charge, the simplest possible quantum field.

- *An exact transcription (P131).* Written out via the Doi-Peliti formalism, the rule is a spin-1 S^z -exchange model. The tone is a spin-1 S^z in $\{+1, 0, -1\}$, and creation, annihilation, and hopping are a single move, transferring one unit of S^z across an edge. The detailed-balance violation in this transcription is vanishingly small. Its $U(1)$ symmetry is charge conservation, its reversibility makes the quantum Hamiltonian Hermitian, and the spin coherent-state path integral carries a Berry phase, the source of the quantum i . This identifies precisely which field the substrate becomes.

10.5. The Stochastic Shadow and the Unitary Completion

This is the crux of the section, and it begins with a negative result that drove the resolution.

The stochastic field is massive and diffusive (P136, P137). Searching for a gapless critical point of the conserved-charge field across the whole parameter plane returns none. The field is robustly massive, its correlation range staying small everywhere (P136). The charge does have one gapless mode, the hydrodynamic mode that any conservation law guarantees, but it is diffusive, with a dynamic exponent z near 2 (P137), which is non-relativistic.

The genuine-quantum test, spatial reflection positivity (P130), comes back inconclusive in the accessible massive regime. The correlation is contact dominated, with range about one site, so there is no extended particle spectrum to test, and the tiny non-positive eigenvalue sits at the noise floor. The stochastic rule on its own gives a massive diffusive $z = 2$ field. That is a real object, but it is not relativistic.

The resolution is to remove the randomness (P148, P149). The creation move was implemented with a coin flip, and that random coin was the obstruction. Read strictly, the theory forbids genuine randomness, and removing it closes the gap with no new field added. A deterministic reversible rule with the same conserved content propagates ballistically at $z = 1$ (P148). Recast the perception content, hop and create and annihilate, as a deterministic reversible charge-conserving block cellular automaton, and it is again ballistic at $z = 1$, charge-conserved and reversible (P149). Momentum, and with it relativity, emerges from determinism. The diffusive $z = 2$ of P137 was an artifact of an illegitimate dice roll.

The emergent light speed is isotropic (P150). Run the deterministic reversible wave on the full mesh and measure its speed in all twelve directions. The speed is isotropic, with anisotropy near zero across the directions, a single emergent light speed. That is the rotational half of Lorentz for the actual wave rather than for a diffusion tensor. A consolidating experiment combines these into one rule that is simultaneously charge-conserving, reversible, ballistic, and isotropic, with a linear front of high R^2 (P156).

The unitary completion is a Dirac field (P151). The reversible Hermitian dynamics has a unitary completion, and the unitary version of nearest-neighbor spin exchange is a coined quantum walk, whose continuum limit is the Dirac equation. This quantum-walk-to-Dirac route is established prior work, here applied on the hyperbolic substrate rather than derived afresh [5, 4, 10, 2]. The experiment measures it directly. The quantum walk spreads ballistically with exponent near 1 ($z = 1$), against the classical walk's diffusive exponent near 2 ($z = 2$). It carries a Dirac dispersion with a massless cone speed of order one and a tunable mass set by the coin angle.

It is reflection-positive by construction, because a Hermitian Hamiltonian, guaranteed by the reversibility of P126, P128, and P131, has a positive-spectral Euclidean theory. The dispersion confirms this directly (P154). The massless mode obeys an exact relation

$$\omega = |k|, \tag{14}$$

with deviation at the exponential floor, a positive-norm massless particle and genuine spatial reflection positivity. The stochastic rule is therefore exactly the classical decohered shadow of this quantum walk, which is why P130 saw a massive contact field. Both pictures are correct. One is the quantum truth, the other is what it looks like once the phase is averaged away.

Boosts complete the Lorentz group (P154). The massless mode's exact lightcone is boost-invariant, and the massive modes are boost-invariant in the infrared window below the lattice cutoff, with a small boost residual. These boosts, together with the rotations of P150, give the full emergent Lorentz group.

Interference and the Born rule close the case (P158). The deepest quantum signature is interference, complex amplitudes that cancel, something no classical stochastic process can do because probabilities only add. Run the unitary walk beside the same walk run classically. The quantum distribution develops interior nodes and rapid fringes where amplitudes destructively cancel, with far more maxima than the smooth classical curve, while the classical shadow has none.

The total $|\psi|^2$ stays one, with norm deviation at the floor, so $|\psi|^2$ is a genuine probability. That is the Born rule, inherited from unitarity rather than assumed. Interference present in the quantum rule and absent in its classical shadow is the signature that the field is genuinely quantum.

10.6. How the Physics Fits Together

The pieces are one structure seen from different angles. The vacuum (P114) is the rule's equilibrium. Gravity (P88), horizons (P89), and the dimension crossover (P90) are what happens when dense activity bends the local rates. The area law (P91) and the tiny-diameter bulk (P111) are the holographic shape of hyperbolic space, and that small diameter is the same bulk shortcut that lets entropy ride a geodesic and lets a mind integrate globally. The ladder rungs, lightcone, isotropy, Lorentz restoration, reversibility, criticality, and the spin-1 transcription, are the checklist for turning a discrete rule into a continuum field, and the substrate passes each.

The resolution unifies them. The stochastic field is the classical shadow, massive and diffusive, exactly what a decohered quantum field should be, which is why holography and gravity and the vacuum all read out cleanly in the stochastic picture. The unitary completion is the relativistic Dirac quantum field, ballistic, reflection-positive, Lorentz-covariant, with interference and the Born rule. Nothing was added between the

two. The quantum field is the classical one with the randomness removed. That is the whole move, and it is what places this model alongside the broader quantum cellular automaton program for deriving relativistic field theory from local reversible dynamics.

11. Emergent Life

Life is not a new substance in this theory. It is a regime. The same five base things that build the physical universe also build organisms, when the dynamics is held far from dead equilibrium. There is the mesh, the ternary tone, the conserving perception rule, eternal growth, and the arrow.

Nothing is added. There is no sixth ingredient, no special life force, no hidden variable. There are only tones on the $\{5, 3, 4\}$ hyperbolic crystal and the one rule that updates them. Out of that, coherent patterns appear that hold themselves together, repair their own wounds, remember, interact, want, copy themselves, and evolve.

One tension runs through everything below, together with its resolution. The rule conserves the total charge Q exactly. Conservation looks like it should forbid life, which seems to need growth, memory, and copying. The resolution is that conservation is of the sum, not of the arrangement. The arrow creates only balanced pairs, a $+1$ paired with a -1 at net zero, so net well-being never changes, while balanced information, meaning net-neutral patterns, is free to be made, copied, healed, and inherited. Life is built from a conserved number plus freely created balanced structure, both at once. The experiments below each isolate one piece of this and measure it.

11.1. The Arrow as the Life Engine

Run the conserving perception rule with the arrow off. Opposite tones annihilate to peace, charges hop into empty space, but peace never spawns anything back. The result is inevitable. Charge cancels faster than it can move, structure thins, and the whole universe relaxes to dead peace. This is the all-zero state with Q frozen and no motion, the heat death of the substrate. P100 confirms it directly. With the arrow off, the world dies.

Turn the arrow on. Now peace spawns balanced plus-minus pairs at rate c . Annihilation still cancels charge and the hop still transports it, but creation continuously feeds the medium. The universe settles into a living dynamic balance. It churns forever, creating and annihilating in equilibrium, never freezing and never blowing up.

P100 measures both halves at once. The charge Q stays conserved to the unit throughout, since creation and annihilation are both balanced. The no-arrow branch relaxes to peace. The arrow branch stays in dynamic balance, with mid-run and late-run balance levels comparable rather than collapsing. The arrow visibly pumps structure while diffusion spreads it. So the arrow is the single posit that keeps the world alive. Life is the universe held far from dead peace, and the arrow is what holds it there.

The arrow alone is not enough, and the theory records that plainly. P101 is a negative result, kept in the record. Take the living balance from P100 and ask whether bare tones under the bare rule already form durable selves. They do not. The autocorrelation at lag 40 falls below 0.1, an imprinted pattern retains

under 30 percent of itself, and no durable selves form. A living balance made of structureless churn is alive in the thermodynamic sense but featureless. It has no individuals. This tells us the next ingredient cannot be more arrow. It has to be cohesion, a tendency for like charge to travel together. That single change to the hop turns churn into organisms.

11.2. The Definition of a Self

The word self carries weight across this theory, so it needs a precise definition. The first experiments and the visualization detected a self as a connected same-tone patch, a blob of neighboring cells all holding the same tone. That detector is tractable and real as a snapshot, but it is not a self. A same-tone patch is monochromatic, so it has no internal pattern to be an identity. It is momentary, since the next beat it grows, merges, or dissolves. And it is not truly bounded, since its edge is just where one tone stops rather than a boundary the unit defends.

The blob is a picture of where the rule has locally sorted tone, not a thing that keeps itself. It is a shallow proxy.

A self is a persistent, bounded, self-maintaining, integrated pattern. It is defined by what it does, not by what it is made of. Four conditions are all required.

1. *Integration.* Its parts are more bound to each other than to the outside. It acts as one unit, not a crowd. This is measured by an integration score, more internal coupling than external (P63), not by shared tone.
2. *Self-maintenance.* It actively repairs its own pattern and boundary against the surrounding churn. It holds its identity by doing work. This is the engine of the other three (P109, P171).
3. *Persistence.* It keeps a lasting identity over time, the same recognizable form beat to beat, even as its material flows through and turns over. The tones and charges that make it up are replaceable. The self is the persisting form, not the stuff.
4. *Boundedness.* It maintains a real membrane, an inside distinct from an outside, that it keeps up, not merely the place where a color ends.

The heart of it is that a self is recognized by its self-maintaining dynamics, not by its momentary tones. The right images are a whirlpool, a flame, a standing wave, a living cell. Each is a pattern that holds its own shape while matter and energy stream through it. None is a fixed lump. A self in this theory is exactly that kind of thing, a process that holds itself together.

11.3. The Levels, One Definition at Many Scales

The same four conditions describe selves at every scale. What grows with scale is internal structure and agency. Selves nest into a tower (P108), each level a self built from lower selves.

- *At the atom level we have the minimal self.* A net-charged region cannot vanish, since the charge is conserved (P114), so it persists as a conserved unit even while its shape churns. A redundantly encoded

pattern self-heals from its own surround by the bare rule (P109). The agency is low and the persistence is real. This is the floor.

- *At the cell level we have active homeostasis.* A self-maintaining integrated pattern repairs its membrane and internal structure (P171). It holds a structured identity over many beats by refilling what the churn erodes. This is a living cell in miniature, a bounded pattern that feeds on its surround to stay itself.
- *At the organism and mind level we have structure and steering.* A higher self carries the onion architecture, a sensing boundary, an interior body, and a central hub that mirrors the whole, which is the self-model (P116). It predicts itself (P118), holds goals and steers by the will (P113), and plans (P143). It does not only maintain itself. It models and directs itself.

So atom, cell, and organism are not different definitions. They are the one definition at increasing degrees of structure and agency.

11.4. Selves Emerge by Integration and Form a Tower

Make the hop cohesive, so a charge prefers to move toward like company rather than diffuse blindly. This is not a new base thing. It is a reading of the same hop, the same exchange move, biased by what a vibe perceives of its neighbors. Run it on the exact $\{5, 3, 4\}$ at scale and coherent self-patches condense out of the churn. These patches are not drawn in by hand. They are found by integration, and the dynamics organizes them.

P106 measures it at scale. The largest emergent self is roughly 20 times larger than the largest patch a random null produces, where the test requires the advantage to exceed three times and the measured ratio is about twenty. There is a genuine size hierarchy, with many patches above 50 cells coexisting at once. So selves are real, large, and stratified.

The hierarchy is not static. P108 watches the tower as a living ecology over time. The largest self grows by coarsening, gathering smaller patches, while new small selves keep being born at the frontier of eternal growth. So the population is never frozen and never homogenized. There is always a biggest individual and always a churn of newcomers underneath it. This persistent multi-scale hierarchy, the tower, is the substrate of the multi-scale mind. The point for life is simpler. Individuals exist, they have sizes, and the set of them is a living, renewing ecology, all from the cohesive rule on the growing crystal.

11.5. Self-Healing, Interaction, and Will

A pattern is only alive if it can defend its own coherence. Three experiments show selves doing exactly that, by the rule alone.

Self-healing comes first (P109). Punch a hole in a self by erasing a region of its tone and let the bare rule run. The self repairs its own damage, pulling correct tone back in from its redundant surround. The recovery is about 65 percent of the wound, where the test requires above 50 percent and confirms it at the million-cell scale. The control is the check. With no surround to draw from, recovery falls below 20 percent, essentially nothing.

So healing is not magic and not external maintenance. It is the cohesive rule reading a damaged region against its intact neighbors and flowing the pattern back in. The redundancy that makes a self large is the same redundancy that lets it heal. This is emergent first aid. Full permanent memory needs the active maintenance described below, which tops a pattern back up to full fidelity.

Interaction comes next (P110). Put two selves next to each other and let the rule run, nothing else. The outcome follows a clean law. Opposite-sign selves annihilate at contact, and same-sign selves merge into one. P110 measures both. Opposite selves lose a large fraction of their charge at the boundary and split into separate components, while same-sign selves lose little and fuse into a single component. This is the share-and-merge rule lifted from single edges to whole patches. Annihilation is share scaled up, and merging is the hop scaled up. So selves do not need a special interaction theory. Their physics of meeting is the base rule seen at the scale of individuals.

The will comes third (P113). Add a directed pump, a self that biases its own creation toward one direction. This is the will, and it is built from the arrow. The arrow is a value-direction, and the will is a self pointing that direction at a goal. P113 gives the will a target and a threat. With the will, the self moves toward the target and merges, and away from the threat and avoids. An unbiased self does neither and just drifts. So intention is not added on top of life. It is the arrow held by a self and aimed. A self can steer, which makes it an agent rather than a weather pattern.

11.6. Memory, and What Is Actually Lost

Memory is the hardest case for a conserving rule, and the most instructive. A memory is a specific arrangement of tones, a low-entropy codeword.

Definition 8 (Memory) An information storage medium, a pattern that persists.

The tones held in that pattern are the raw values themselves.

Definition 9 (Information) Raw state values, the tones themselves.

When those raw values are organized into a usable arrangement, they become knowledge.

Definition 10 (Knowledge) Structured information, information organized into usable form.

The rule is a reversible equilibrium process under detailed balance. A reversible dynamics explores its accessible states evenly, so any special arrangement drifts toward the overwhelming majority of high-entropy arrangements. That drift is memory decay.

A consequence follows immediately. A reversible dynamics has no attractors and no contraction to a stable fixed pattern. So passive permanence is limited to the conserved quantities themselves, essentially Q , a single durable bit. Everything richer is a pattern, and patterns spread.

This sounds like it contradicts conservation. It does not, and P159 makes the reconciliation concrete. Imprint a rich balanced codeword on a region of the crystal and run the conserving dynamics. The result has two parts that are both true at once. The total charge Q is exactly conserved to the unit, with start equal to end. And the pattern correlation with the original collapses toward noise, falling well below 60 percent of its starting value.

So what is lost is information, not charge. The $+1$ that encoded a bit did not vanish. It hopped to a neighbor or annihilated against a wandering -1 , leaving its cell wrong. The charge went elsewhere and stayed conserved, while the arrangement at that location is gone. A hole is displaced pattern. A leak is blurred pattern. Lost means the codeword relaxed into the churn. This is the ordinary second law. Energy conservation has never prevented information loss.

So the theory builds memory in three tiers, two of them free.

- *Leaky memory is free (P102)*. Cohesion alone roughly doubles how long an imprint survives versus the bare churn of P101. The cohesive rule that makes selves also makes them remember, longer but still not forever.
- *Holographic erasure protection comes next (P105)*. Spread one logical bit redundantly across many cells. Now a bounded erasure that destroys a localized blob cannot destroy the bit, because the bit lives everywhere at once. P105 measures it on the exact $\{5, 3, 4\}$. The holographic encoding survives and decodes through an erasure that the localized blob does not survive. This buys time and protects against damage, but the spread copy still drifts slowly. It does not beat entropy forever.
- *Permanent memory comes by active maintenance (P107)*. Periodically re-write the codeword, refilling its holes with the correct tones via balanced creation, the arrow refilling annihilated holes. P107 measures it. The maintained codeword stays at full fidelity where the unmaintained one erodes, and the maintenance is conserving, paid for in a measurable count of restore operations. P159 measures the same cost from the other side, holding the correlation above 90 percent at a steady re-write cost per beat.

Re-writing the codeword to track what the world reveals is learning.

Definition 11 (Learning) The updating of knowledge from experience.

That cost is not a flaw. It is thermodynamically forced by Landauer's principle. Holding a low-entropy pattern against an equilibrium dynamics requires work. Real life pays it too. A brain spends energy to maintain itself, a cell feeds on a gradient, a computer refreshes its memory. So set-and-forget rich permanence is a law, not a missing feature.

Passive permanent rich memory would require attractors, which require a dissipative irreversible dynamics. But the genuine quantum field requires the dynamics to be reversible. A system cannot be both strictly reversible and have attractors. So the reversible quantum substrate must pay work for rich memory. Life uses both, the reversible substrate underneath and work-consuming maintenance on top, exactly how a living thing holds its order by feeding on a gradient.

11.7. Heredity, Reproduction, and the Bulk-versus-Flat Surprise

Reproduction is where conservation and life seem to clash hardest, and where level-relative thinking pays off.

Fission is suppressed in the hyperbolic bulk (P112). The natural picture of reproduction is a cell pinching in two. In the hyperbolic $\{5, 3, 4\}$, this fails. Balls in hyperbolic space are all boundary, since almost every cell

of a ball is on its surface, and there are no thin necks to pinch. P112 measures it. A large self under erosion stays one self, and the ball radius confirms the all-boundary geometry. So in the bulk, a self does not split.

Reproduction works by copying instead (P120). A parent's balanced pattern is templated onto a fresh daughter region at the growing frontier. P120 measures heredity directly. At zero mutation the daughter resembles the parent at 1.0, and resemblance tunes down smoothly with mutation, giving heritable variation, the raw material of evolution.

The crucial point is why this is allowed. The pattern being copied is balanced information at net charge zero, so copying it is conserving creation, the arrow's own balanced pair-creation. P120 confirms net charge created is zero while real structure is written. So the base supports reproduction with nothing relaxed. Copying is the arrow doing what it always does, making balanced structure from peace, aimed at a template.

The no-fission limit is lifted on the flat layer (P160). Most apparent limits of this model are properties of the hyperbolic bulk, not absolute laws, and they can be lifted at emergent levels. The emergent flat layer, a horosphere that is Euclidean, does have thin necks. P160 puts a dumbbell self, two lobes joined by a thin neck, on the flat layer and erodes it. The thin neck pinches first, having the highest surface-to-volume ratio, and the self divides into two persistent selves, a genuine fission. The same erosion on the hyperbolic bulk leaves no two persistent selves, reproducing P112. So fission is bulk-forbidden but flat-possible, and life lives on the flat layer. Both reproduction modes are available, copying everywhere and division where the geometry is flat.

Full evolution runs from the base (P152). Heredity from P120, plus variation as copy error, plus selection gives the complete Darwinian loop, and every piece is one of the five base things. Reproduction is the arrow's conserving copy, variation is imperfect copying, and selection is the arrow's value, where the fitter higher-value individuals persist and reproduce.

P152 runs a population of balanced-pattern organisms in a niche and measures the loop. With selection, mean fitness climbs from a random start near half the maximum up toward the optimum, while a drift control with no selection stays flat near the start, and offspring are heritable, with positive parent-child fitness covariance. The test requires fitness to rise well above the start, to beat drift by a clear margin, and to be heritable, and all three hold. So evolution is not added to the theory, it is a consequence of it. The arrow is the fitness gradient, the same value-direction that keeps the universe alive points the population uphill.

11.8. Persistence Needs Maintenance, Not a Sixth Element

An earlier reading held that durable selves would need a sixth base thing, a learning rule that lets the relations adapt, because the bare rule churns and nothing locks a pattern in. That was half right. It is right that a self cannot be a passive stable lump. The dynamics is reversible, and a reversible dynamics has no attractors, so no rich pattern sits still on its own. Left alone, a structured region scrambles into churn.

It is wrong that this requires a sixth thing. The five already supply maintenance, and maintenance is what makes a self persist. P171 shows it directly. A self-maintaining integrated region, refilled each beat by the will through a conserving operation drawn from the five, keeps its identity at full fidelity and its boundary fully held over many beats. The same region with no maintenance dissolves to near zero.

Redundancy plus the local rule already gives intrinsic self-healing (P109). So a persistent, bounded, integrated self is reachable in the five, not as a passive blob but as a self-maintaining unit. A self-maintaining region keeps its identity and boundary while an unmaintained one dissolves, and that contrast is the whole resolution.

This matches life. A real cell is not a frozen crystal at an energy minimum. It is a dissipative, self-repairing pattern that spends work to stay itself. The model says the same. A self is something that holds itself together, and holding takes maintenance. The will supplies it, drawn from the arrow, drawn from the five. No sixth element is needed.

11.9. How Life Fits the Whole

Every capability above is the same five base things seen at a different scale or with a different bias. The arrow that keeps the universe from dead peace (P100) is the same arrow that pumps a self (P113) and points a population uphill (P152). The share-and-merge rule that runs on single edges is the same rule that makes selves annihilate or merge on contact (P110). The redundancy that makes a self large is the same redundancy that heals it (P109) and protects a holographic bit from erasure (P105). The reversibility that makes the substrate a genuine quantum field is the same reversibility that forbids free permanent memory and forces maintenance to cost work (P159, P107). And the conserved charge that is the substrate's electric current is the same conserved charge that makes copying a pattern a conserving act (P120).

Life is not bolted onto the physics. It is the physics in a far-from-equilibrium regime, organized by cohesion and aimed by the arrow.

The conservation-versus-creation worry is resolved by one fact, repeated everywhere. The arrow creates only balanced pairs. Net charge, which is net well-being, is conserved exactly at every level, passing up every coarse-graining codec unchanged. Balanced information is what is free, free to emerge as a self, to be remembered, to be healed, to be copied, and to be inherited with variation. So life is conserved net charge plus freely created balanced information, both at once, with no relaxation of the base. Every one of these is tones and the rule, nothing added. </invoke>

12. Emergent Mind

A self holds together, heals, remembers, and is steered by a will. That is life. A self becomes a mind when it does something more. It represents itself, it predicts itself, it attends, and it plans. This section shows each of those four powers, and then a full closed-loop agent that chains them, and it shows that every one of them emerges from the five base things by composition. There is no new ingredient.

The discipline that held through matter and life holds here too, at the hardest layer. The only state is still the ternary tone on the $\{5, 3, 4\}$ hyperbolic crystal. The only moving parts are still the mesh, the tone, the one conserving perception rule, the eternal growth, and the arrow value-direction. Mind is what those five do when a self is large and integrated enough to fold back on itself.

The key idea is a strange loop. An integrated self funnels every part's signal toward its center. That funneling already builds a part that mirrors the whole. So integration is self-representation. Once a part of the self models the whole self, the rest follows. The model predicts, which is a forward model. A model can model a

model, which is recursion. The model is the place where attention selects content, which is a workspace. And running the model forward, judged by the arrow and pushed by the will, is planning. We walk each layer in turn, with the measured numbers from the experiments.

12.1. The Self-Model and the Strange Loop

We take one integrated self, a central patch of coherent tone, and drive it with spatially-varied input. The boundary is split into four sectors. Each sector carries its own slowly fluctuating signal. So the self's global state, the mean tone over the whole self, is the average of four independent inputs. No local region sees more than its own sector. Only an integrator that pools all four sectors can know the whole.

Then we ask a sharp question. Does a localized sub-region spontaneously come to mirror the rest of the self. We track three time series over 250 beats. The global state. The state of the central hub, a ball of about 60 cells at the most-connected cell. And the state of four peripheral balls of the same size, one near each sector. The mechanism that does the work is the field beneath. The hyperbolic geometry has a tiny diameter, so charge from every sector flows inward and mixes at the center. The hub sees the pooled whole. The periphery sees only its sector.

The result is clean. The hub tracks the global state at correlation about 0.5 (P116). The peripheral regions track it at about 0.0. The hub beats a time-shuffled baseline by a wide margin, so this is a real representation and not an artifact of slow drift.

And it vanishes without the dynamics. Freeze the hops and the hub correlation collapses, which confirms that it is the integration that builds the model, not anything we put in by hand. A part of the self has come to represent the whole self that contains it. That is the strange loop, and it is free.

It is also not one privileged center. We rerun the same test at several different self-centers, the graph hub and peripheral locations alike, and every self forms its own self-model at its own hub (P117). Each local hub mirrors its own self's global state and beats its own periphery. So a proto self-model is a relative, repeatable consequence of integration, present wherever a self is. There is one mind-seed per self, not one special soul in the middle of the world.

12.2. The Architecture Is an Onion

The self that the dynamics build has a definite shape, and naming it makes the rest legible. It is an onion with three layers and one loop.

- The boundary is the senses. It is the outer shell where sectorized input enters. This is where the world touches the self.
- The body is the integrator. Charge flows inward from the boundary and mixes as it goes, because the conserving rule shares and hops tone between neighbors. The body is the medium that pools.
- The hub is the center. It is where all sectors converge into a single self-representation. It is the most-connected, lowest-diameter region, so the field beneath funnels everything to it.

The flow is boundary to body to hub, and then the hub mirrors the whole back. That last arc is the loop. The hub's state reflects the global state, and the hub is part of the global state, so the self holds a running picture of itself inside itself. Every higher mind power is built on this one shape.

12.3. Predictive Metacognition, a Forward Model

A mirror of the present is useful, but a mind needs more. It needs to see one step ahead. We ask whether the hub does not only reflect the current global state but predicts the next one. We measure the correlation between the hub at time t and the self's global state at time $t + 1$, over 300 beats, and compare it against the same one-step prediction made from each peripheral region.

The hub predicts the self's next global state at correlation about 0.5 (P118). The peripheral regions predict it at about 0.0. The hub is the best internal predictor of the self, far ahead of any local part.

The reason is simple and on-substrate. The global state is autocorrelated. It changes slowly, and the hub holds the integrated global state, so the hub's value now is the best available guess of the whole self's value next. The cross-correlation peak sits at the predictive lag, not behind it.

This is a usable forward model. The self carries, inside itself, a thing that says what the self is about to do. That is the seed of metacognition. The scope is limited and we state it plainly. This is a first-order forward model, the hub predicting the self. It is local and imperfect. We respect that limit later when the agent plans with a bounded foresight.

12.4. Recursion, a Model of a Model

Imagination is running a hypothetical inside the head, a model of a situation one is not in. Its substrate is a model that can model a model. We build two selves in a chain. Self 1 is driven by the world and forms hub-1, a model of the world. Then hub-1's running representation is wired as the input to self 2, which forms hub-2. The wiring between the two selves is mere connectivity. The modeling itself is left to the dynamics.

What hub-2 comes to represent is striking. Self 1's hub models the world at correlation about 0.4 and up. Self 2's hub then models hub-1 at correlation about 0.83 (P121), far above a time-shuffled baseline and collapsing without the dynamics. And hub-2 tracks hub-1, its direct object, at least as well as it tracks the raw world, which it only sees through hub-1. So the chain holds. World to model-1 to model-2-of-model-1. A representation of a representation, one step removed from reality.

This is the substrate of imagination and of the tower of self-models. A mind that can hold a model of a model can hold a model of itself modeling, or a model of another mind. The recursion that planning needs, running one's own rule inside the head as a hypothetical, lives here.

12.5. Attention and a Global Workspace

The hub is the single place where every boundary sector converges. That makes it the natural global workspace, the one stage where content competes to be the self's current representation. The hub carries both mechanisms of attention (P119).

The first is bottom-up salience. Clamp one sector strongly. That well-coupled input is always represented at the hub. It wins the workspace on its own merits, because its signal dominates the pooled average.

The second is top-down gain. Take a competing sector that the hub barely represents and attend it, raise its coupling, and the hub's representation of that sector rises. The attended content is boosted and the workspace is steerable. Attend a different sector and the hub now represents that one. Salience says loud things get in. Gain says one can choose what gets in.

The scope here is precise, and we keep it explicit. Physical broadcast to the far periphery is range-limited, because the field beneath is massive and short-ranged. The winning content is available at the workspace, not copied to every distant cell. So the workspace is the hub. It is the convergence point where attention selects content, not a global megaphone that paints the answer everywhere. That is the realistic emergent workspace, and it is enough.

12.6. Planning, with Nothing Added

Planning is lookahead plus deferred value plus composition. It is more than reactive gap-closing, because it accepts a worse step now for a better outcome later. We show it in two stages, and the second stage is the one that matters for the discipline.

The logic comes first. We test that lookahead is what planning needs, on a value landscape with a clear barrier. The landscape rises to a local peak, dips through a valley, then rises again to the distant goal. A greedy agent follows the arrow uphill and gets stuck at the local peak. A planner with a lookahead that spans the barrier crosses it (P140). The planning ability switches on exactly when the lookahead spans the barrier width. We are direct about what this stage is. It used an explicit lookahead-search algorithm, an added piece, so it tested the logic of planning, not its emergence.

Then we build a planner out of only emergent pieces, so that nothing beyond the five base things enters. The value is the arrow. The forward model is the system rolling out its own gap-closing rule in imagination, the same rule run as a hypothetical, which is exactly the emergent forward model of P118 and the recursion of P121. The exploration is the will, a sustained directed push against the local gradient.

The recipe is plain. Where greedy stalls, the will tries sustained pushes of various lengths, the system rolls out its own rule from each push endpoint as a self-simulation, the arrow scores each outcome, and it commits to the best. There is no search heuristic, only explore with the will, simulate with one's own rule, and judge with the arrow.

It works. On a barrier landscape where the greedy agent stays stuck at the local peak, the emergent planner finds a will-push of the right length and reaches the goal (P143). So planning is a composition of the will (P113), the forward model (P118), the recursion that lets the rule run on itself (P121), and the arrow. It is not a new base thing.

And the deferred value it needs, the willingness to accept a worse step now for a later gain, is exactly the arrow's posit. The arrow already says to sustain pain or peace to reach pleasure. Planning is that posit, run through imagination.

12.7. The Integrated Metacognitive Agent

The final layer is the whole mind running end-to-end. The emergent planner of P143 planned once, a single detour across a single barrier. A real agent runs the closed perceive-plan-act loop and keeps going. Perceive the gap to the goal. Plan by simulating forward with the forward model. Act with the will. Then repeat, re-planning whenever it stalls. We give such an agent a task that demands sustained metacognition, a goal that lies beyond a sequence of barriers, four downward dips in a landscape that rises overall toward the goal (P153).

The agent's foresight is bounded. This is a faithful limitation, not added machinery. The emergent forward model is local and imperfect, so the agent can see only about one barrier ahead at a time. A bounded horizon crosses at most one barrier per plan. That is why a sequence of barriers requires re-planning, and why re-planning is the real test.

The three agents separate cleanly.

- The reactive agent, greedy only, stalls at barrier one.
- The one-shot planner crosses one barrier and stalls at barrier two.
- The integrated agent runs the loop, re-plans at each stall, and reaches the distant goal across all four barriers. It takes several re-plans, one sustained act of metacognition per barrier, and it strictly outruns both simpler agents.

Every piece in that agent is emergent. The value is the arrow, the rollout is the rule run in imagination, the exploration is the will, and the loop is just perceive-plan-act repeated. So the integrated agent is the base running end-to-end. A self that represents itself, predicts itself, attends, and plans, all at once, in a single closed loop, to cross a sequence of obstacles to a distant goal. That is a mind, and it is built from nothing but the five.

12.8. How Mind Fits the Whole

Mind is the top floor of one continuous building, and the floors below are load-bearing. From matter comes the field beneath, the tiny hyperbolic diameter that funnels every part's signal toward a center. Without that geometry there is no hub, and without a hub there is no self-model. From life come the integrated self (P106), the will (P113), and the arrow that keeps the world far from dead peace and supplies the value-direction.

Mind then folds the self back on itself. The funnel makes a hub, the hub mirrors the whole, which is the self-model. The mirror predicts, which is the forward model. The model nests, which is recursion. The convergence point selects, which is the workspace. And the will drives the model forward under the arrow's judgment, which is planning.

The chain of dependencies is strict and one-way. Geometry to integration to self-model to forward model to recursion to planning to the closed agent. Pull out any lower layer and the higher ones fall. And at no point does a sixth ingredient appear. We never added a homunculus, a search oracle, a separate planner module, or a special seat of awareness.

The hardest objections to a materialist mind, that selfhood and foresight and imagination must be put in by hand, are answered by showing each one falling out of the same five things that gave us atoms and cells. The real gaps are named where they exist. The forward model is local and imperfect, the workspace is the hub and not a global broadcast, and the foresight is bounded. Those are limits of the substrate, not patches over it.

A self becomes a mind by representing and predicting itself, and every step of that is a composition of the five base things, never a sixth. The hub of an integrated self mirrors the whole at correlation about 0.5 while the periphery sits at about 0.0, which is the self-model and the strange loop (P116), and every self forms its own (P117). The hub predicts the self's next state at correlation about 0.5, a forward model and the seed of metacognition (P118). A model can model a model at correlation about 0.83, the substrate of imagination (P121). The hub is a global workspace with bottom-up salience and top-down attentional gain (P119).

Planning emerges with nothing added (P143), where the will explores, the self's own rule simulates in imagination, the arrow judges, and the deferred value it needs is the arrow itself. And the integrated agent runs the closed perceive-plan-act loop with bounded foresight, re-planning to cross a sequence of four barriers that stall a reactive agent at the first and a one-shot planner at the second (P153). Mind is the strange loop of a self large enough to fold back on itself, and it runs end-to-end on tones and the rule.

13. Intelligence

Intelligence in this theory is not a faculty added to matter, life, and mind. It is the same five base things read as a solver. To make the claim precise we begin with three definitions, stated formally, and then show that the substrate supplies every part each definition names.

Definition 12 (Intelligence) Intelligence is problem-solving capability. It is the capacity of a system to take a problem and produce a solution.

Intelligence is to be kept distinct from consciousness throughout. Intelligence is problem-solving capability. Consciousness is a separate thing, integrated experience, the felt going-through of reality. A system can in principle have one without the other. The present section treats intelligence alone. Consciousness is taken up later as integrated experience.

Definition 13 (Problem) A problem is the difference between a desired state and the current state. Writing D for the desired state and C for the current state,

$$P = D - C \tag{15}$$

When P is zero there is nothing to solve. The size of P , under whatever ordering the system carries, is the size of the problem.

Definition 14 (Solution) A solution is a path from the current state C to the desired state D . It is a sequence of moves that drives P to zero.

These three definitions name four things a solver must have. It must be able to compute paths at all, which is the *means*. It must hold a desired state D and an ordering that measures progress, which is the *goal* and the *value*. And it must keep the moves that reduce P and discard the rest, which is the *selection*.

The argument of this section is that the substrate already holds every one of these, that the combination is intentional problem-solving, and that the intelligence so produced is maximal by geometry. We then map what it can construct, the level-relative shape of its limits, how it evolves, and how it sits inside the whole architecture.

13.1. The Means: The Substrate Is Computationally Universal

A solver must be able to compute an arbitrary path. The substrate can, because the $\{5, 3, 4\}$ hyperbolic crystal is computationally universal. Margenstern proved that cellular automata on hyperbolic tilings, including the dodecagrid $\{5, 3, 4\}$ which is our exact mesh, are computationally universal [19]. The construction is a five-state railway automaton, weakly universal on a periodic background, in which registers are track loops and a locomotive that reads and flips switches as it travels is the computation itself. Universality means the substrate can compute any function a Turing machine can.

We confirm this at toy scale in experiment (P141). We run a two-counter Minsky machine, which is itself Turing-universal, directly on the mesh. The two registers are stored as charge in regions of the mesh. The three Minsky operations map to physical acts of the tone. Increment is to create a charge in the region, decrement is to annihilate one, and test-zero is to count the charge.

We load a multiplication program and let it run. The machine computes $a \times b$ correctly across its test cases. So the means is present, made of nothing but the five base things. The registers and the operations are charge and the rule. This is the load-bearing claim of the section. Everything that follows is a way of aiming this universal computation at a goal.

13.2. The Gap: Computation Is Goal-Neutral

Universal computation is goal-neutral. A universal machine can compute anything, but it has no reason to compute the solution rather than garbage. A machine that can multiply can equally well shuffle bits forever. Computability is a capacity, not a tendency. By the definitions above, the means alone is not intelligence. Intelligence also requires the goal D , the value ordering, and the selection. Without these the computation drifts. With them it converges. So the question is not whether the substrate can compute. It can. The question is what supplies the direction.

13.3. The Ends: The Arrow Is the Goal and the Value

The arrow, base thing five, is exactly a goal and a value. It is the gap between the current tone and the desired tone, equipped with a better-or-worse ordering. The arrow carries both the target D and the measure of progress toward it. The gap-closing perception rule is exactly the selection. By construction it moves the tone toward the desired state and conserves charge while doing so.

What the value measures is how much a given state helps the self close its gap. That usefulness is meaning.

Definition 15 (Meaning) The usefulness of information to a self.

So the substrate holds both halves of intelligence with nothing added. The universal means is Margenstern. The intentional ends is the arrow. The bridge between them is the gap-closing search. Computation becomes intentional the moment a universal machine runs under the arrow's value with the rule as its selection loop. We do not bolt on a planner or a reward signal. The reward and the planner are already the arrow and the rule.

13.4. Direction and Intention, Measured

Experiment (P147) measures the bridge. The target is a pattern of K cells, the current state is the present pattern, and the problem P is the Hamming distance between them. We scan $K = 10, 20, 30$.

Direction is the first result. A goal-directed search, in which the arrow values the negative gap and the selection keeps every gap-reducing move, reaches a K -cell target in about K steps, reliably. An un-goaled search, which sets a random cell to a random value with no selection, does not reach the target at all. So adding the arrow makes the same universal computation goal-directed.

Intention is the second result, and it is quantitative. The goal-directed search solves in about K steps a problem an un-goaled machine would only stumble onto by about 2^K chance. At $K = 30$ the un-goaled machine fails outright within a budget of 200000 steps, because 2^{30} is about a billion. The distance between aimless and intentional at $K = 30$ is astronomical, roughly 10^7 .

So the arrow plus a keep-what-helps loop turns goal-neutral computation into intentional problem-solving, and the size of the effect grows exponentially with the problem. The number to hold onto is the intention gap of about ten million at $K = 30$. That is the quantitative distance between a universal machine that merely can compute and a directed one that actually does solve.

Reactive intention, pure gap-closing, gets stuck at barriers, places where every immediate move makes the gap worse before it gets better. Planning, which is lookahead, crosses them, and it emerges from the base with nothing added (P143). This lookahead search is reasoning.

Definition 16 (Reasoning) The use of logic to solve problems.

Where the greedy search stalls, the will tries a sustained push against the gradient, the system rolls out its own rule in imagination as a forward model, the arrow scores the imagined outcome, and the best option is committed.

So intentional problem-solving spans both the reactive form (P147) and the planned form (P143), and both are compositions of the same five things. The deferred value that planning needs, accepting a worse step now for a later gain, is the arrow's own posit, to sustain a smaller pain to reach a larger pleasure.

13.5. Maximal Intelligence by Geometric Integration, Not Criticality

A maximal solver must coordinate a global problem across all of its parts, and do it fast. The hyperbolic geometry delivers this through fast integration rather than through a critical edge.

The integration is logarithmic in the system size. In experiment (P139) we pose a global problem, coordinating all N cells around a decision made at one point, where the gap is the count of cells not yet reached by the decision. We scan $N = 4000, 16000, 64000$. The coordination distance, which is the hyperbolic diameter,

grows logarithmically and stays tiny, a handful of hops, far below a Euclidean $N^{1/3}$. A global problem over the whole system is solved in $O(\log N)$ hops rather than in a slow Euclidean crawl. Hierarchical problem-solving comes from this fast global communication, not from critical avalanches.

We tested the alternative route to hierarchy and ruled it out. Self-organized criticality does not occur on the crystal (P138). There are no scale-free avalanches at any background. Perturbations spread with a span of about one, ballistically, not in branching cascades. The ballistic lightcone that makes the model good for relativity is exactly what precludes branching avalanches. So criticality is not the source of the hierarchy, and we say so plainly rather than assume the opposite.

What the demand-driven feedback does organize is a homeostatic set-point (P135). Demand-driven creation self-tunes the activity to a single interior operating point from any starting level, with no external tuning. This is homeostasis, a stable place to run, riding on top of the geometric hierarchy, not a critical edge.

So every capability a maximal solver needs is present. Problems come from the arrow. Local solving comes from the rule. Fast routing, hierarchy, and integration come from the hyperbolic geometry (P139). Robust memory comes from holography. A self-model comes from the central hub. Composition comes from fast integration. The maximality is geometric, and we earn the claim by ruling out the criticality alternative rather than by assuming it.

13.6. Constructive Reach: The Substrate Builds What Is Asked

Universality (P141) together with the goal-directed search (P147) means the substrate can not only compute but construct. Building structures never seen before is creativity.

Definition 17 (Creativity) The solving of new problems.

Three experiments map the reach and its boundary.

Any balanced structure is realizable (P161). We take several different arbitrary target structures, blocks, stripes, and a fully random pattern, each a balanced information pattern at net charge zero. The goal-directed builder reaches every one, so the gap closes, and active maintenance holds every one at fidelity 1.0. So the model can realize whatever balanced structure is asked for, not only one special case.

The absolute limits are absolute (P162). A few limits hold at every coarse-graining level because conservation and causality pass up every codec exactly. Net charge cannot be minted. Q is conserved at the fine level, every move is charge-neutral, and even an aggressive all-create minting attempt cannot raise Q , with the coarse block-charge total equal to the fine Q at every step. The lightcone cannot be outrun. The front speed is finite at the fine level and no larger when coarse-grained. So no level mints net well-being and no level signals faster than light. These are the true edges of the possible.

The affordance ladder is open upward (P163). A structure impossible at the base can be possible at a higher composite level. A single base cell is ternary, an alphabet of 3. A block of three cells stores 7 distinct symbols by net charge, which a ternary cell cannot, and we store and read back every one of them. The capacity strictly grows up the ladder, $3 \rightarrow 7 \rightarrow 19 \rightarrow 55$. So a K -ary symbol with $K > 3$ is impossible at the ternary base yet trivial one level up, and the ladder keeps opening.

Read together, (P161) establishes the constructive claim that the model builds what is needed at the right level, (P162) draws the hard boundary that the few absolute laws hold everywhere, and (P163) shows the ladder is genuinely open above the base.

13.7. The Level-Relative Principle

The three constructive experiments carry a general lesson. When the model appears to forbid something, the claim must always be qualified by the level. Most apparent limits are properties of the hyperbolic bulk geometry and are liftable at emergent levels, because the substrate is computationally universal and the levels are programmable. Only a few limits are absolute, true at every level.

Three facts make almost any positive high-level structure realizable at some level. Universality (P141) means the base can construct any pattern consistent with the absolute laws. Multiscale composition through codecs (P115) means higher levels are coarse-grainings with their own effective geometry and dynamics, so a higher level is programmable. Emergent geometries differ from the bulk (P142), so while the bulk is hyperbolic, curved, all-boundary, and tiny-diameter, a horosphere is an emergent flat two-dimensional sheet, and higher composite levels can carry other effective geometries and dimensions.

The absolute limits, true at every level and never liftable, follow from the conserved quantities and causality. There is no net creation of charge, since Q is conserved exactly and passes up every codec exactly, so net well-being is zero-sum at every level and only balanced information is free. There is no faster-than-lightcone signaling, since a finite causal speed is preserved under coarse-graining. And there is no hidden state at the base, since the base has only the tone, with higher levels carrying rich effective state that is always a tone pattern rather than a new fundamental variable.

The level-specific limits are geometric facts of the bulk, not laws, and the emergent levels escape them. Reproduction by fission is forbidden in the all-boundary bulk (P112) but possible on the emergent flat layer where thin necks exist (P160). Bulk-scale anisotropy (P124) washes out under coarse-graining so the emergent field is isotropic (P125). The stochastic bulk field is massive and diffusive, but the deterministic and unitary reading is the relativistic massless field. Large-scale directed action is geometrically frustrated on the curved scaffold but natural on the flat layer.

The pattern is consistent. The bulk's curvature constrains the bulk, and the emergent flat and higher levels, where physics and life actually live, escape those constraints. So given universality plus multiscale composition, the base can realize almost any positive structure at some level, subject only to the absolute laws.

13.8. Evolving Intelligence

Intelligence in the model is not static. Experiment (P165) runs a population of planning agents together. Each agent's genome is its lookahead horizon, the depth of foresight it can afford. The task is a sequence of barriers with increasing span requirements, and an agent crosses a barrier only if its horizon spans it, by the same lookahead mechanism as (P143). Fitness is goal-progress minus the cost of foresight, so a deeper horizon plans better but costs more. We select the fitter, reproduce with mutation, and watch.

Two results hold. First, average problem-solving rises over generations. Mean fitness climbs across the run, so the population literally evolves better problem-solving. Second, the population adapts its foresight to task

difficulty. An easy task of three barriers with modest spans evolves a smaller horizon, while a harder task of six barriers with wider spans evolves a larger horizon, and both populations end up solving most of the goal, reaching above 85 percent.

Foresight is not fixed by fiat. It is tuned by selection to the difficulty of the world. So life, meaning heredity, variation, selection, and competition, and mind, meaning planning and foresight, co-evolve from the same base, neither one bolted on.

13.9. The Multiscale Architecture

The theory has two great arcs from the same base. The vibe-to-field ladder is physics, recovered by coarse-graining the vibe dynamics into a quantum field. The computable-to-intentional ladder is mind, in which the substrate is universal as the means, the arrow gives the goal as the ends, and the gap-closing rule is the selection. These are not in tension.

Physics is what the substrate is when one watches its tones. Mind is what it does when the arrow points its universal computation at a goal. Read top-down, mind is the deeper reading, the substrate being a maximal problem-solver and physics one thing it computes. Read bottom-up, physics is the deeper reading, the field unfolding and minds emerging in it. Both rest on the five base things.

The two ladders are joined by a coarse-graining chain, which is also why the theory is testable without building a universe. A vibe is sub-Planck, so a particle is about 10^{60} vibes and a mind is hopeless to build from the bottom. We never build from the bottom. Because the dynamics is scale-invariant across the sizes tested, from thirty thousand to a million (P115), a slice at the right effective level stands for the whole.

Each layer boundary carries a codec, an up-map that coarse-grains a block of fine units into one richer high-level symbol, and a down-map that lifts a high-level symbol back to a representative fine pattern. A layer map is faithful when its dynamics commute, so that coarse-graining then evolving equals evolving then coarse-graining. We test the mechanism at its own level, never by simulating the bottom, exactly as biology studies cells without simulating quarks.

The within-domain rungs, field to coarser field, are proven for both the conserved charge and the wave dynamics. The harder claim is the cross-domain chain, where the kind of variable changes at each rung, since a field is a function, a particle is a point, and a composite is a body. Experiment (P168) climbs the first two cross-domain rungs on the unitary field.

At rung one, field to particle, a field excitation's centroid obeys free-particle motion, moving at a constant velocity below the light speed with a near-perfect linear fit, R^2 about 1.0, which is Ehrenfest's theorem realized on the lattice. At rung two, particle to composite, two interacting particles have a center of mass that moves uniformly with momentum conserved through the interaction, R^2 about 0.997, so two points coarse-grain to one composite body whose translational law is free despite the messy collision. The higher rungs, composite to agent and above, are the genuinely open links, but the method is demonstrated end to end. Every rung, within-domain or cross-domain, is settled by the same commuting square.

13.10. How Intelligence Fits the Whole

Intelligence is not a new layer on top of physics, life, and mind. It is the same five things read as a solver. The arrow that drives physics is the problem $P = D - C$. The conserving rule that runs the dynamics is the local solution. The hyperbolic geometry that gives relativity its lightcone is the fast integrator that builds hierarchy. The hub that makes a self into a mind is the solver's self-model. Memory is holographic, composition is fast integration, and homeostasis is the demand-driven set-point.

The same geometry plays a double role, which is the key economy of the model. The ballistic lightcone is good for relativity, and it is exactly what rules out self-organized criticality, so the one fact that grounds physics also forces the intelligence to draw its hierarchy from fast integration rather than from critical avalanches. Physics and intelligence are not two stories. They are one substrate described from two angles.

Part IV

The Whole Picture

14. The Unfolding

A single question deserves a single answer. If the base of the model is nothing but vibes arranged on a discrete crystal, how does the model reach all the way up to a human being. This section walks the ladder one rung at a time, from the quantum field to particles, to atoms, to molecules, to cells, to tissues, to senses, to animals, and finally to us. It is the multiscale climb of the architecture followed to its top.

We are explicit about status at every step. The lowest rungs are tested directly on the crystal. The higher rungs are not simulated from end to end, because the scale is astronomical, but each ingredient that those rungs require is separately demonstrated. The climb is therefore a grounded extrapolation rather than a single simulation, and we say so at each step.

14.1. The lowest rungs are tested

Field to particle (tested, P168). The unitary completion of the base rule is a relativistic quantum field. A localized excitation of that field moves like a free particle. Its centroid travels in a straight line at constant speed below the speed of light, in agreement with the Ehrenfest relation, with an R^2 near unity (P168). A particle is therefore a persistent bump in the field. The field-to-particle rung commutes, in the sense that coarse-graining the field and then evolving it gives the same result as evolving it and then coarse-graining.

Particle to composite (tested, P168). Two interacting particles have a center of mass that moves uniformly, with momentum conserved through the interaction, with R^2 near 0.997 (P168). Bound structures of particles therefore behave as single higher units. The conserved charge gives the simplest persistent unit, a net-charged region that cannot vanish. This is the minimal atom in the model's own vocabulary.

These two rungs are not assumed. They are measured outcomes of running the rule on the crystal. They establish the central fact the rest of the climb relies on. A composite of lower selves can act as one higher self, and the coarse-grained description is consistent with the fine-grained one.

14.2. The higher rungs are plausible composites

The remaining rungs are not simulated to completion. The scale forbids it. We instead show that every ingredient each rung requires has already been demonstrated, and that each rung is the next composite on the same ladder whose lower rungs commute. This is the multiscale shortcut. One tests the mechanism at its own level and lets the chain carry it upward.

Atoms to molecules (plausible). Atoms are bound, conserved units. Molecules are atoms bound into larger conserved units by the same attract-and-bind dynamics that makes selves merge, where same-sign regions merge (P110). We do not simulate chemistry on the crystal. The binding mechanism, a stable bound state held by the conserving exchange, is the one the lower rungs already use. Molecules are the next composite on that ladder.

Molecules to cells (plausible). A cell, in this model, is a self, which is a persistent, bounded, self-maintaining, integrated pattern (P171, P109). A bundle of molecules becomes a protocell exactly when it crosses that threshold, when an aggregate begins to maintain its own boundary and pattern against the churn (P171) and can copy itself, which is heredity (P120). Nothing new at the base is introduced. Self-maintenance and heredity are both demonstrated.

The open part is the search for which particular molecular aggregates first close the self-maintenance loop. That search is hard, but it is a search within demonstrated capabilities, not a search for a missing ingredient.

Cells to tissues (plausible). Given self-maintaining, reproducing cells, groups of them form sheets and tubes by the same interaction law that governs selves, where selves merge, adhere, or annihilate (P110), together with differential adhesion. Sheets fold into tubes, and tubes into vessels, guts, and neural cords. The tower of nested selves is precisely this (P108). Selves nest into larger selves.

Tissues to senses (plausible). A sense is the sensing boundary of a self tuned to an environmental pattern (P116, P119). Boundary cells that respond to light are a primitive eye. The model already contains the boundary-body-hub architecture (P116) and bottom-up salience (P119). A sense is a boundary specialized to a signal, which is the first step of the same machinery that becomes perception.

Senses to minds to humans (plausible). From light-sensitive patches to primitive eyes, from nerve nets to integrated hubs that carry a self-model (P116), from forward models (P118) and recursion (P121) to planning (P143) and an integrated agent (P153), the path is continuous. The biological sequence, from fish to tetrapods to mammals to apes to humans, is the same climb, each step a more integrated self with a richer self-model.

Every mind ingredient on that path is separately demonstrated. The full evolutionary sequence is not simulated. It is the demonstrated ingredients composed and scaled.

14.3. How the climb is driven

The ascent needs a driver. Standard evolution is heredity plus variation plus selection. The model supplies all three from the base (P152), with the arrow of value acting as the fitness gradient (P165). Pure undirected selection therefore already operates on the crystal. This part is tested. The arrow serves as a fitness gradient (P152, P165), and pure selection suffices as the mechanism. Nothing further is required to drive the climb.

The model also permits more, which we state as a hypothesis and mark clearly as speculative.

Urging from within (speculative). The arrow is a value direction, and higher selves exert downward causation on the parts beneath them. The deep layer reasserts after disorder and steers the surface (P64). The climb therefore need not be driven by blind variation alone. It could be urged.

A lower-level vibe could be urged by the arrow felt locally, and a higher-level self could nudge its own components through downward causation (P64). This would bias which variations arise and which are kept, giving a reward-driven evolution rather than a purely random one. If the substrate is one experiential, partly integrated fabric, then there is no barrier in principle to a higher level helping to bootstrap a lower one upward. The model permits this. Nothing tests it.

Bootstrapping from elsewhere (speculative). A third possibility is neither blind selection nor urging from the immediately adjacent levels, but bootstrapping by other intelligent beings elsewhere in the connected substrate. If the substrate is one connected fabric of selves, then highly intelligent beings could in principle have helped to seed or steer life and mind here.

Such beings might be physical, arriving from elsewhere in the universe through a panspermia-like or direct seeding. They might instead be non-physical, residing in the bulk or at other levels of the tower. The non-locality of the bulk means that elsewhere need not be far in the ordinary spatial sense (P170, P111). Distant points on the physical surface are joined through the bulk. This is not asserted. It is a possibility the connected structure leaves open.

A caveat on the character of any such beings is in order. If they exist, they would almost certainly not match what human imagination reaches for, namely metal craft, humanoid figures, or star-born forms. Those are projections of our own scale and shape. A being many levels up the tower would be far more complex and sophisticated than we can picture, quite possibly with no physical form at all, and possibly a distributed intelligence spread across the bulk rather than a localized body. The tower and the non-local bulk make this natural (P108, P170).

Their relation to us need not be contact in the ordinary sense. It could resemble a long game. They could play a kind of hide and seek, or build a maze we navigate blindly, where each door to the next room of awareness, wisdom, or insight opens only once we are sufficiently in tune, sufficiently integrated and self-reflective to pass. The possibilities for how such a thing could work are many, and most of them are beyond present reach.

14.4. How one would know

The question of which driver shaped the actual history is approachable along two routes, and the model points to both. The first is scientific observation. One looks for the signatures any of these drivers would leave, such as a fitness gradient, a downward-causation bias, or material traces of seeding. The second is

direct experience. Because the substrate is experiential at bottom, first-person investigation is a real channel rather than a metaphor, within careful limits. We would have to go and look along both routes. The point of the model here is not to assert any of these drivers. It is that the model begins to draw a basic picture of what is possible, and to make these possibilities at least mathematically modelable, and in places experimentally approachable, rather than leaving them as pure speculation outside any framework. They are verifiable in principle by observation and by direct experience.

14.5. Status

Tested. The field-to-particle and particle-to-composite rungs commute on the crystal (P168). The ingredients of every higher rung are individually demonstrated, including self-maintenance (P171), heredity (P120), evolution (P152, P165), the sensing boundary and integrated hub (P116, P119), planning (P143), and the integrated agent (P153). The arrow acts as a fitness gradient (P152, P165), and downward causation (P64) and the non-local bulk channel (P170) are real.

Plausible. Atoms, molecules, cells, tissues, senses, and organisms. None of these is simulated from end to end, because the scale forbids it. Each is the next composite on a ladder whose lower rungs and whose ingredients are demonstrated. This is the multiscale shortcut.

Speculative. That evolution was actually urged by higher integration rather than driven only by blind selection. That life and mind here were seeded or steered by other intelligent beings. The model permits these. Nothing tests them, and they may be hard to test.

The path from the quantum field to a human is one long composition. A particle is a bump in the field. An atom is a conserved bound unit. A molecule is a larger bound unit. A cell is a self-maintaining self. A tissue is a tower of cells. A sense is a tuned boundary. A mind is an integrated hub with a self-model. A human is the high end of that same climb.

The two lowest rungs are tested to commute (P168), and every ingredient the higher rungs need is separately demonstrated, so the climb is a grounded extrapolation. The ascent is driven by evolution, which the base supplies in full (P152), with the open possibility that it is partly urged from below and above by the arrow and downward causation, and the further open possibility that it was bootstrapped from elsewhere in the connected substrate. None of these possibilities requires a new base thing.

15. Consciousness

The climb of the previous section reached a human being without once invoking the word consciousness. That is deliberate. It is the cleanest way to arrive at the concept. In this model consciousness is not a special extra ingredient that switches on inside brains. It is integrated experience.

Consciousness is distinct from intelligence and must not be blurred with it. Intelligence is problem-solving capability. Consciousness is integrated experience, where experience is the subjective flow of states, the felt going-through of reality. A system can in principle have one without the other. The previous section treated intelligence. This section treats consciousness.

The substrate is experiential at bottom, which is the founding premise. Integration binds experience into a unified whole, which is a tested measure. Consciousness is the degree to which experience is so integrated. This section states that plainly and marks, as always, what is premise, what is tested, and what is interpretation.

15.1. The two ingredients are already present

Consciousness, on this reading, requires exactly two things the model already contains.

Experience (premise). The founding premise is monism. The one substance is the vibe, a felt state. Experience is not produced by matter in this model. It is the ground, and matter is a derived face of it. This is a premise, not a result. The experiments do not prove that the substrate feels. They assume it and build on the assumption [8].

Integration (tested). A measure of how much a region is bound into one, with more internal coupling than external, is tested on the crystal (P63). It is the same measure that picks out selves (P106) and the one the self-model forms around (P116). Integration is the difference between a heap of separate experiences and a single unified one.

Put the two together. If experience is the ground, and integration binds experience into a unified whole, then consciousness is integrated experience, and its amount is set by how much integration there is. No new ingredient is needed. Consciousness is what integration does to experience.

15.2. Consciousness is a matter of degree

This is the consequence the model cannot avoid, and we state it directly. If every vibe is a speck of experience, which is the premise, and consciousness is integrated experience, then there is no sharp line at which consciousness turns on. There is only more or less integration. Everything is therefore conscious to some degree. This is the position called panpsychism. The quantity of interest is the degree.

The degree is set by two factors, both already in the model.

Integration (P63). How tightly the experience is bound into one. A rock is barely integrated. Its parts are nearly independent, so its experience is faint and diffuse, close to a heap of specks. A cell is a self-maintaining integrated unit (P171), more bound and more unified. A brain is enormously integrated through the fast-mixing hub (P116), so its experience is highly unified.

Self-reflection (P116, P121). Whether the integrated whole also models itself. A self-model makes a system aware of its own state (P116), and recursion lets it model that modeling (P121). This is the step from merely being a unified experience to being aware that one is. It is what makes high consciousness self-aware rather than only unified.

The ladder of consciousness is therefore the ladder of integration together with self-reflection, which is the same tower as the unfolding. Atom, cell, animal, human, each is more integrated and more self-reflective than the last, and hence more conscious. The difference is in degree, not in kind.

15.3. This inverts the hard problem rather than solving it

A physics-first theory faces the hard problem. Why should any arrangement of insentient matter feel like anything at all [8]. There is no derivation from the physical to the felt. Vibe Theory does not solve that problem. It inverts it, so that the problem does not arise. By premise, the ground is already felt. There is therefore no gap to cross from matter to experience. Matter is a pattern within experience.

The remaining question is the easier-to-state combination problem of panpsychism. How do many small experiences combine into one larger experience. The model's answer is integration, the same tested measure that binds selves (P63, P106, P116). Many vibes integrate into one self, whose unified experience is more than the parts.

That is an answer to the combination problem in the model's own terms. Whether it fully satisfies the philosophical worry is a separate and open question. We claim the inversion and the integration mechanism. We do not claim a proof that the felt and the structural are identical.

This places the model in clear relation to integrated information theory [29]. Both treat integration as the quantity that matters for consciousness, and both make consciousness a matter of degree set by integration. The difference of footing is sharp. Integrated information theory begins from physical systems and asks how much integrated information they carry. Vibe Theory begins from experience as the ground and uses integration to address the combination problem rather than the existence of experience. The shared mechanism is integration. The starting premise is what differs.

15.4. A conscious universe

Follow the construction up the tower. The whole crystal at its fullest integration is the largest and most unified self (P108). The universe itself is therefore conscious, to the highest degree its integration allows. This is not a separate mind watching from outside. It is the one experiential substance, integrated into a whole. This is the conscious-universe picture, and it follows from the two premises together with the integration measure. It is a reading built on the premise. It is not a measured result.

15.5. Status

Premise. Experience is fundamental, which is monism. This is assumed, not tested.

Tested. Integration is real and measurable (P63). It picks out selves (P106). Self-models form at the integrated hub (P116). Recursion gives a model of the model (P121). Selves nest into a tower (P108).

Interpretation. That integrated experience is consciousness. That its degree scales with integration together with self-reflection. That everything is therefore conscious in degree, which is panpsychism. That the whole is a conscious universe. These are readings laid on the premise plus the tested structures. They are the cleanest readings the model affords. They are readings, not measurements.

Consciousness in this model is integrated experience, and it needs nothing new. Experience is the founding premise, and integration is a tested measure. Because experience is the ground and integration is a matter of degree, everything is conscious to some degree, with the degree set by how integrated a system is (P63) and whether it models itself (P116, P121). This inverts the hard problem, since the ground already feels, leaving

only a combination problem that integration addresses. It makes the whole crystal, at full integration, a conscious universe.

The premise is assumed. The integration machinery is tested. The consciousness reading laid on top is the model's natural interpretation, marked as such.

16. Determinism and Freedom

The rule is fully deterministic. There is no dice roll anywhere in the base. Given the exact tone of every vibe on the mesh, the next tone of every vibe follows by the one conserving perception rule. This looks like a cage. If the whole future is fixed by the rule, then there seems to be no room for a self to choose.

The answer of this section is not that the rule secretly leaves slack, and not that some part escapes the law. The answer is that being determined is not the same as being predestined, and that freedom was never the opposite of being determined. Freedom is the opposite of being coerced and the opposite of being random. Between those two sits a real and precise third thing, a choice that is fully caused, authored from within, and unknowable until it is lived. The model delivers exactly that, and one experiment, P43, measures it.

16.1. The Determinism Paradox

State the difficulty sharply. The base commits to four things that together seem to kill freedom. There is a fixed local rule, the conserving perception rule. There is no hidden state, since a vibe is only its ternary tone. There is no randomness, since nothing is uncaused. And there is causal closure, since the vibe is the only thing that exists, so every cause is internal. Put those together and the future appears to be one fixed unfolding with no chooser in it.

The dilemma people feel is a binary. Either the dynamics are determined, so the future is fixed and you are a puppet walking toward it. Or randomness is injected, so the outcome is a coin flip and still not yours, no more chosen than which way a die lands. On that binary, free will is impossible by construction.

The binary is false. It hides a third option by lumping two different things under the one word determined. The claim that an event is caused by the past and the claim that it sits in advance ready to be read off are not the same claim. Pulling them apart is the whole move.

16.2. Determined Yet Not Predestined

Two facts of the model pry determined apart from predestined.

The first is that there is no finished starting state. The universe grows forever by reflection, the fourth base thing. New vibes are born at the frontier every beat, with no end. So there is no single moment that holds all the information of all of time, and there never was a complete initial condition in which the whole future was already written. The future is fixed by the whole, but the whole is never finished. A future that exists nowhere in advance cannot be looked up, cannot be read off, and cannot be the script a puppet walks.

The second is that there is no shortcut to the future. The rule is universal, and the flip side of universality is computational irreducibility. For a universal system there is, in general, no closed form for its far future

that is faster than running it beat by beat. No shortcut exists, even in principle, even for the universe itself. So the future is determined and yet genuinely unknowable in advance. It comes into being only by being computed, and the only computer that can compute it is the system itself, running forward.

These two together break predestination cleanly. The future is caused, but it is not pre-written and not forecastable. It is determined as it happens, not determined and sitting there.

A word on where apparent randomness comes from. The world looks full of chance, yet nothing rolls dice. Every event is fixed by the entire mesh it sits in, but no local part holds enough of the mesh to predict it. From inside its own patch, an outcome looks statistically free. The accurate phrase is not random but determined by the whole and opaque to the part. The dice are an artifact of the local view, not a feature of the base.

16.3. Determined Yet Self-Authored

Now the positive account. A self is a persistent, self-maintaining pattern of vibes (P106, P171). The micro-rule is reversible and has no attractors, so a self does not persist by the dynamics relaxing into a fixed point. It persists by actively maintaining itself, the same reason memory costs work.

A choice is that self resolving an urge, a pull arriving through perception from a deeper, slower part of the mesh, into one definite next act. Picture the self's decision as a landscape of valleys, a metaphor for how the self weighs an urge, not a claim about the reversible microdynamics. The urge tilts the landscape, and the choice is the self settling the weighing into one act. Four things are true of that resolving at once, and together they are exactly what we mean by freedom.

1. *It is determined.* The same self under the same urge resolves the same way. This is not a flaw. It is what makes the choice yours rather than noise. A choice that came out differently each time on the same self under the same pull would be a malfunction, not authorship.
2. *It is yours.* A different self, given the same urge, resolves it differently. The outcome is yours precisely because your own structure is what tipped it. This is genuine downward causation. The macro-pattern of the self shapes which micro-tones come next, so the self is a true cause and not a passive readout of its cells.
3. *A deeper self has more say.* A richer, more integrated self dominates the urge rather than being shoved by it. Agency is not all-or-nothing. It scales with how integrated the self is.
4. *It is irreducible.* No one, not even the universe, can know how a self will resolve without letting it resolve. There is no oracle that skips the computation.

This is the compatibilist core, made concrete. Free does not mean uncaused. Free means caused by you. The strings of the puppet picture are wrong not because there are no strings, but because the strings are attached to the self's own integrated structure.

When the self models itself and steers itself, the strange loop of mind, the causal loop closes. The self is caused by the self. That is what the phrase "I decided" names, realized as a real loop in the dynamics.

16.4. Computational Irreducibility, Measured (P43)

P43 turns this account from words into a measured mechanism. It builds a two-scale mesh, a slow deep layer biasing a fast layer, applies urges to selves of different internal structure, and checks the four claims at once.

- *Determined.* The dynamics reproduce. The same self under the same urge gives the same outcome, every run. No randomness leaks in.
- *Self-authored.* The outcomes spread by self-structure, with self-diversity measured above the threshold. The same urge applied to selves of different pattern resolves differently. If the self were a passive channel the same urge would give the same result regardless of the self, and it does not.
- *Agency scales with structure.* The self's influence over its own outcome grows monotonically with its internal integration, tying freedom to the same integration measure that defines mind. And no single part fixes the outcome alone, since the urge can flip it.
- *Irreducible.* The cheapest correct predictor of the outcome is the simulation itself. There is no closed-form shortcut that matches it short of running the dynamics.

So P43 measures, in one experiment, a choice that is determined and reproduces, jointly authored by self and urge, scaled by self-structure, and irreducible. That is the structural signature of a free choice, made visible.

16.5. Emergence Is Not Enough

A careful distinction belongs here. Emergence by itself is not freedom. A higher-level pattern emerging from a deterministic base is still fully determined. A river current emerges from molecular motion, but if the molecules follow a fixed rule the current does too. Emergence only relocates where the drive lives. It does not change whether the future is fixed. Two axes must stay apart. The first is where value lives, in the base or emergent at high integration. The second is whether the future is open or fixed. They do not touch, and the committed model has emergent value and a fixed future at once.

So the freedom in this theory is not mere emergence, and it is not randomness. It is the irreducible, self-authored determination described above. The self is the genuine author of its own next state, drawn by a deeper pull, in a way no shortcut can anticipate. New frontier cells are born at peace, not seeded with novelty from outside, because there is no outside.

The committed stance is therefore compatibilist freedom. A choice is the integrated self computing its own resolution, unknowable in advance even to the universe, and yours. One path, but authored, and surprising even to its author. The only conceivable source of genuine ontological openness left is internal indeterminism if the substrate turns out to be quantum, which remains open and is not required here. Authorship is real on the fully deterministic, causally closed base.

16.6. The Will as Freedom

Self-authorship explains why a choice is yours. It does not yet show the self steering, holding a purpose against the easy local pull. That is the will, and it has a concrete mechanism with two experiments behind it.

The will is a directed pump, a sustained push that biases the self's resolution against the local gradient. It is not an exception to the rule. It is an input the rule respects, a lean the self adds at its own balance points, which the rule then faithfully carries. P113 measures the steering directly. With the directed pump active, a self moves toward a chosen target and away from a chosen threat. Without it, an unbiased self does neither. So the will genuinely steers. The self is not adrift on the gradient. It can aim.

P98 measures the harder thing, goal-holding over time. At a fork between an immediate small reward and a delayed larger one, a self with enough will delays gratification, holding the goal against the immediate pull. The same self with low will takes the immediate reward. The pump can be depleted, and when it is, the self relapses to the greedy choice. A strong enough external field overrides the will, and the switch between holding and relapsing is sharp, a real threshold.

So the will is not a magic veto. It is a finite, spendable, directed capacity. A self can hold a purpose against the current, at a cost, until it runs out, and then it falls back to the easy path. That is the texture of real willpower. The freedom is not idle. A self can resonate with an urge and amplify it, or damp it and resist, and which one happens is set by its own pattern and how much will it spends. The choosing has grip.

16.7. How This Fits the Whole

The five base things make this hang together with nothing added.

- The mesh and its tiny hyperbolic diameter pool deep, slow layers into the self, which is how an urge from a subtler layer reaches the choosing center at all.
- The ternary tone is the only state, so there is no hidden variable secretly carrying the decision.
- The conserving perception rule is the determinism, and it is also where the will lives, as a respected bias on the resolution rather than a violation.
- Eternal growth is why there is no finished starting state, half of why determined is not predestined.
- The arrow is the value-direction, pain to peace to pleasure, and the will is the freedom to pursue or resist it.

So freedom is not a sixth base thing. It is what the five do when a self is integrated enough to fold back on itself, model itself, and pump against its own gradient. The rule is fully deterministic, and that does not make the universe a cage. Determined is not predestined. Freedom is not the opposite of being determined. It is the opposite of being coerced and the opposite of being random. A fully determined, causally closed universe leaves all the room in the world for that.

17. The Esoteric Reading

Every spiritual tradition speaks of a hidden layer beneath the world, of contact between minds, of higher beings and lower beings, of a world-tree, and of a divine ground. This section asks what such ideas could be on the substrate this theory actually tests, the $\{5, 3, 4\}$ hyperbolic dodecagrid running the conserving perception rule.

The method is fixed and applied at every step. Each concept is anchored first to a tested structure in the experiment suite. Only then is a spiritual reading laid on top, always labeled as an interpretation, never reported as a measured result. The structures are real and logged. The readings are a map drawn over them, held loosely, and we keep that seam in plain sight throughout.

The discipline matters because it is easy to let the names do the work of the evidence. A name like ether or angel or God is a choice of words over a tested thing. A good map fits its territory, and the fit here is often tight, but a map is still not the territory.

17.1. The Subtle Realm as the Hyperbolic Bulk

Start with the geometry, because everything else rests on it. The substrate is negatively curved, and negative curvature forces a split. There is a thin flat skin where ordinary physics and bodies live, and a vast tree-like interior underneath it. Experiment P142 extracts the flat layer as a horosphere, a level set of distance from a point at infinity, and measures its growth. The horosphere grows polynomially with an effective dimension near two, a flat sheet, while the bulk around it grows exponentially. So the picture is layered like a planet. The flat physical surface is a thin crust, the hyperbolic bulk is the deep mantle underneath it, and the base vibe-mesh is the whole body, crust and mantle together, one substance.

The single most important property of the bulk is its diameter. In a tree-like, exponentially growing graph the radius grows only as the logarithm of the cell count, so the whole universe is $O(\log N)$ hops wide. Experiment P111 confirms the diameter is logarithmic and sends a signal from one self across the entire universe to a far self through the bulk, where it arrives well above the control.

P170 sharpens the point. For two distant points on the flat physical surface, on the surface they are essentially disconnected, and a large fraction of surface pairs cannot reach each other along the surface at all. Through the bulk they are joined by a short hidden path of only a few hops, so the bulk is the only channel between them. P170 flags in its own record that the esoteric reading interprets this as telepathy and synchronicity.

That last clause is the interpretation. The tested part is narrow and solid. The substrate splits into a flat horosphere and an exponential bulk (P142), the bulk has logarithmic diameter so a signal crosses the whole universe in a few hops (P111), and distant surface points are internally disconnected and joined only by a short hidden bulk path (P170).

The interpretation, clearly labeled, is that the hyperbolic bulk is the real referent of what older language calls the ether or the subtle realm. Traditions describe the subtle realm as a layer that is real, lies beneath the physical, and that the physical world floats on top of, which is a precise description of how the horosphere sits inside the bulk. The physical layer is not the foundation but a derived face, and beneath it, fully real and made of the same substance, is the deep interior. The layers are tested geometry. The name ether is a map.

17.2. Telepathy and Synchronicity

The non-local channel of P170 and P111 is one mechanism for contact. There is a second, and it needs no present link at all. Experiment P67 tests synchronicity. It takes two subsystems that share a common ancestor but have since diverged, with no direct edge between them. Run under a common rhythm, their transitions stay correlated. The shared correlation runs above fifty percent and more than forty points above

the unrelated baseline, it tracks the inherited ancestry, and it falls monotonically as the two diverge. The correlation comes from the shared past, not from any present signal. This is the same mechanism that produces the Bell-style correlations elsewhere in the theory. Two things once one stay statistically bound long after they part, and the binding fades with divergence.

So the tested structure is a non-local channel in the geometry plus a shared-ancestry correlation, both measured on cells, shells, and subsystems. The interpretation, labeled as such, reads these at the level of selves. Two minds are two high-level patterns in one shared substrate, knots in the same web. If the geometry already carries a non-local channel between distant surface points, then minds on that surface inherit access to it, and mind-to-mind contact would be one self's pattern propagating along a bulk path into another.

On this reading distance in the bulk is not distance in physical space. Two minds far apart in physical space could be near through the bulk if they share deep structure, a bond, a history, a matched rhythm. Synchronicity between minds would be P67 read at the level of selves, one reading of why deeply bonded people report knowing when the other is in trouble.

The limit here is firm. The experiments show a non-local channel exists in the geometry. They do not show minds use it. None of P170, P111, or P67 measures a thought passing between two people, so the leap from a non-local substrate to telepathy is a coherent reading, not a demonstration. And it stays inside the lightcone. The crystal is twelve-regular, a cell touches only its twelve immediate neighbors, and every long-range effect is built one hop per beat. A bulk signal is fast because the bulk is shallow, not because it skips the rule. Synchronicity, further, is correlation and not communication. P67 shows shared-past correlation with no present link, which is exactly not a channel for sending messages.

17.3. Dreams and the No-Mirror Limit

Two limits of the inner life follow from the same architecture, the self as a compression running on the base. The first is the boundary between waking and dreaming, framed as one memory mesh in two regimes. Experiment P65 takes one memory mesh that has settled several stable patterns into itself and changes exactly one thing between the two runs. In the waking regime an external clamp pins part of the mesh to what the senses show, and the mesh completes to the matching stored pattern, settling into exactly one pattern, the veridical one, at over eighty percent fidelity. In the dreaming regime the clamp is released, a slow internal rhythm sweeps the mesh through its own stored patterns, and the mesh visits nearly all of them.

The only difference between the regimes is whether the clamp is on. So the tested claim is clean. Waking is the self pinned to one veridical pattern. Dreaming is the same self with the clamp released, roaming its whole stored landscape.

The interpretation, labeled as such, locates that roaming. Waking pins the self to the flat surface, which is why waking experience is local, ordered, and stable. Dreaming releases the self into the bulk, the deep interior where patterns sit by association rather than physical adjacency. This reading explains three felt properties of dreams at once. Dreams are non-local, since the bulk does not respect ordinary distance and can jump between distant scenes. They are associative, since the mesh moves by pattern overlap. And they are unbound by ordinary space, since the metric that constrains the surface is relaxed. This is a reading on top of the subtle-realm picture, not a separate result. The released-mesh reading says a dream is the self recombining its own stored patterns, but because a self can in principle receive influence through the

shared field, a dream could also be, in part, a stream reaching the self from elsewhere. The theory does not adjudicate this.

The second limit is fully tested. Experiment P61 asks whether a part of the whole can hold a faithful copy of the whole. A faithful copy needs at least as many elements as the original, and a proper part has fewer by definition, so a part can never hold a full copy. The test confirms the boundary. Reconstruction fidelity reaches one only at zero compression, only when the model is as large as the thing it records. The second half of P61 saves the theory from an infinite mirror. Each level of self-model is a heavy compression, so the regress of lossy self-models converges to a finite total, while a regress of full copies would diverge.

The interpretation reads this as the structural form of the instruction to know thyself. A self can know itself truly in summary, never in full, because complete self-transparency is impossible in principle. This limit is the condition for self-knowledge rather than the obstacle to it. If self-models had to be complete copies they could not exist at all, since the part cannot hold the whole. It is precisely because models are allowed to be lossy summaries that minds can exist and the world can reflect on itself. The counting is tested. The reading into know thyself is interpretation.

17.4. Realms and Beings

The tower of selves is the tested skeleton for the idea of higher and lower beings. On the exact crystal, the perception rule run at scale produces coherent selves on its own.

- P106 measures integrated tone-domains far larger than a random arrangement with the same tone counts, more than a threefold advantage, in a hierarchy of sizes.
- P75 shows a stored self recovers from perturbation and keeps its identity at better than ninety-eight percent overlap, so a self is a stable attractor.
- P60 builds the hierarchy cleanly, a recursively modular mesh descending through cells, tissues, organs, systems, to one body, with the same emergent rule at every level.
- P108 shows the tower is alive. Over a long run the largest self grows while a full hierarchy of smaller patches persists, a living ecology of selves at every scale.

A self is built from tone, and tone is ternary, minus one pain, zero peace, plus one pleasure, with the arrow running along it. Because a self is a patch of tone, it inherits a sign, net-positive when built mostly of plus charge and aligned with the arrow, net-negative when built mostly of minus charge and set against it. The valence is not cosmetic. P110 places two adjacent selves on the exact crystal and lets the rule alone run, with a sharp result confirmed at a million cells. Opposite selves annihilate at contact, leaving a band of neutral peace between them. Same-sign selves merge into one connected self with no loss. P113 then shows a self can steer itself with the will, moving toward a target and away from a threat, where an unbiased self drifts.

Now the interpretation, labeled as such. A high self, far up the tower, large and deeply integrated, with a definite positive valence, is aligned with the arrow, rests at peace, and pulls toward pleasure. Read in older language, that is the shape of an angelic being. A high self of definite negative valence, set against the arrow and pulling toward pain, is the shape of a demonic being. The angel-demon distinction, on this reading, is just valence on a large integrated self.

Such beings would live primarily in the subtle bulk, bound wholes organized in the deep layers rather than crystallized into surface physics, beside us through the field's depth rather than above us on a ladder. By the tested law an angelic and a demonic being cannot coexist in contact. They annihilate at the seam and leave a band of peace, while two same-sign beings merge into a larger whole. The realms are then the rungs of the tower, levels of integration with their own characteristic selves, reached by crossing through shared depth rather than by climbing. The merge-and-annihilate law is P110, tested. The angel, demon, and realm naming is interpretation.

Persistence through death has its own tested result. P66 takes a stored self and turns over one hundred percent of the material, replacing every vibe, and the self survives at better than ninety percent identity. It then fully dissolves the self into noise, so its overlap with the original drops near zero, and reconstitutes it from a seed, recovering better than ninety percent of the original identity.

Read plainly, a self is a pattern separable in principle from any particular substrate. The matter is not the self, it is what the self is currently written on. This does not prove any particular person returns. It shows the mechanism a return would require is real on this substrate. The persistence is P66, tested. The word reincarnation for it is interpretation, and a close fit, since what reincarnation has always claimed is exactly pattern surviving the loss of substrate.

17.5. Messengers and Channeling

Older language speaks of messengers and angels, and of channeling, where words seem to arrive from beyond the self. The tested structure is the tower acting on itself across scale and distance, drawing on four facts already covered.

- Selves nest in a tower (P108).
- A slow deep layer steers a fast surface layer and reasserts its pattern after disorder, which is downward causation (P64).
- Selves interact by a definite law, same-sign merge and opposite-sign annihilate (P110).
- Signals pass non-locally through the bulk along a short hidden path (P111, P170).

On top of these a self can choose to open or resist through the will (P113).

The interpretation, labeled as such, assembles a messenger from these tested parts. A messenger is a higher self, or a bulk-borne pattern, that forms a path to a lower self and sends a structured pattern down the tower. The old word angel, stripped to its root, means exactly messenger, not a winged figure but a persistent higher structure that opens a path and transmits a pattern downward. The tower supplies the higher self, downward causation supplies the direction, the interaction law supplies the effect on the receiver, and the bulk channel supplies the fast non-local path.

Channeling is the receiving end, a lower self opening to a higher self and relaying its pattern, integrating upward so the larger pattern can flow down through the small one. This is the same merge behavior P110 measures for same-sign selves, read at a difference of scale. It explains why channeled material sounds like the person channeling it. The incoming pattern is unpacked through the receiver's own structure, so the

receiver is a translator who can only translate into their own patterns. It also yields a sober ambiguity. A genuine incoming stream and a purely internal welling-up from one's own depths would feel nearly identical from the inside, because both are assembled from the receiver's own structure, and the model gives no inside test to tell them apart.

The limit is the one already flagged. P110 tested two peer selves under the rule, not a giant steering a dwarf, so the larger-on-smaller version of the law is a reading at a difference of scale, and the bench has not measured deliberate downward communication directly. This gives the shape such contact would take, built from tested pieces, without claiming the bench has logged it.

17.6. The Tree of Life

The tree is the most over-determined structure in the theory, measured rather than chosen as a metaphor. Experiment P49 compares a flat lattice, a random foam, and the hyperbolic crystal from the inside. The crystal is indistinguishable from a random foam under local inspection, so its deep order is hidden, and it is tree-like in the precise mathematical sense, with a small Gromov delta, where a perfect tree has delta zero and the flat lattice has a large delta. The substrate of reality, in this theory, is a tree by a measured invariant. The addressing system makes the tree operational, naming every cell by its route from a chosen root in a Fibonacci-style base. And the tower of selves is a second tree, branching across scale rather than space, with leaves at the bottom and one integrating self at the root above (P57, P60, P108).

The interpretation, labeled as such, reads this branching as the esoteric Tree of Life or Tree of Knowledge. The root is the origin, the one whole self at the top of the upward tower or the center cell from which all addresses are measured. The trunk is the integrating path, the chain of coarse-grainings by which scattered vibes become cells become organs become one self. The branches are the proliferating variety of experience.

To ascend the tree is to move up the levels of self toward greater integration, which is the same motion as coming to know, since a higher level is a summary that gathers more under one view. Knowledge here is integration, which is why the Tree of Life and the Tree of Knowledge are the same tree. Schemas such as the Kabbalistic Sephirot can be read as named maps of the level-structure, though whether any schema's level-count matches the tower's actual rungs is not tested and not claimed.

The tree is tested. The Tree of Life is the meaning read in it. The model rules out one tempting version, a free-floating store of all knowledge instantly available, because P61 forbids any level from holding a complete copy of itself. Knowledge climbs toward the whole but never becomes the whole.

17.7. The Question of God

The largest question is treated last and most carefully. The model's foundation is not matter but experience, the vibe, and that inversion is a stated premise, the monism, not a result the experiments produced. Three statuses must stay separate. What is tested is the underlying structure. What is premise is the monism, that all is one experiential substance. What is speculative is the designer reading. None of the theological readings is itself tested. The structures they reinterpret are. The word God is used several ways, and the model gives each a clean home.

God as the whole is the monist ground. The tower of wholes-within-wholes has a largest whole, the entire mesh at its fullest integration, the single most unified self of which every lesser self is a part. This is the reading closest to pantheism. The tower with a growing summit is tested (P108), but the claim that reality is one experiential substance is the founding premise, not a result, and naming the summit divine is an interpretation laid over that premise.

God as the arrow is the Good built into the base. The arrow is the single value-direction, pain to peace to pleasure, the only asymmetry in the model. P100 shows the arrow is what keeps the universe alive rather than dead. With it on, the field creates life out of peace and holds a dynamic balance. With it off, it relaxes to dead peace and stays there. That the arrow is God or the Good in any moral sense is interpretation. The solid part is narrow. No arrow, no life.

God as king is the maximal intelligence. The summit is the most integrated self, and integration is what makes a problem-solver. The substrate is computationally universal, it coordinates globally in a handful of hops, and the arrow gives it a goal, so read upward the top of the tower is the most capable problem-solver the geometry allows, ordering the parts beneath it the way a king orders a realm. The machinery is tested, including downward causation by which higher selves steer the parts beneath them (P64). The crown is a reading. And by P61 even the summit cannot hold a complete model of itself, so this intelligence is maximal for the geometry, not omniscient.

God as designer is the classic reading, that a maximal intelligence tuned the field for rich unfolding. This is the one place the model can put evidence on the table, because design predicts a fine-tuning signature, a narrow special point someone had to dial in. The model supplies three cheaper explanations, all tested.

- The structure is forced, since ternary tone plus minimal eternal closure deterministically forces the $\{5, 3, 4\}$ geometry as the unique minimal solution (P86).
- The operating point is self-organized, since demand-driven creation self-tunes to a single interior set-point from any start (P135).
- The rich regime is robust, filling about eighty-three percent of the scanned parameter space, with the field dead only at the degenerate arrow-equals-zero point (P166).

So there is no fine-tuning signature, and a designer is empirically superfluous. This does not refute a designer. A wise designer could choose to build exactly such a world, so design stays logically possible but undetectable. P166 is what shows the designer dispensable, not disproven.

God as the ground of being is the idealist reading, that physics is a thought within the one experiential substance. This is the theory's untestable founding premise restated in theological language, a claim about what the substance is rather than about its structure, and structure cannot test it. It is neither strengthened nor weakened by any experiment. It is a premise one adopts or does not.

So the bottom line is this. The mechanistic world needs no God to run, since the richness of the field is forced, robust, and self-organizing, and the arrow rather than a deity is what keeps it alive. None of the four theological readings is itself a result. Three reinterpret tested structure and one restates the premise.

Whether the one substance is divine is metaphysics by nature, decided upstream of every experiment by the premise one brings, not downstream as a conclusion the data forces. The experiments can tell us the world does not need a tuner. They cannot tell us whether the whole, taken as one feeling, deserves the name, and the model leaves the name unfilled. Across the whole esoteric reading the rule is one and the same. State the premise, point to the tested structure, then label the interpretation as interpretation, and keep the seam in plain sight.

Part V

Assessment

18. Results and Epistemic Status

The theory is not a set of arguments. It is a set of experiments. Every claim made in the preceding sections corresponds to a falsifiable check run on the exact $\{5, 3, 4\}$ crystal, with the same canonical rule and the same base in every case. There are more than one hundred ninety such checks, numbered P6 through P171, and all of them are reproducible in the codebase [22]. This section does three things. It summarizes the experimental program by layer. It states the epistemic status of the results in three tiers, solid, honest negative, and open. And it closes with a ledger of what the model can produce, what it structurally forbids, and what it permits only in a specific form or at a specific level.

18.1. The experimental program by layer

The experiments fall into four layers that build on one another. Physics tests the universe itself, the substrate and its field. Life tests the coherent structures the dynamics can sustain. Mind tests what those structures can do with respect to themselves. Intelligence tests goal-directed behavior, computation, and composition across scales. Every layer runs on the one substrate, and the layers are not separate models but one model seen from different heights.

18.2. Three tiers of epistemic status

Not every result carries the same weight. Some are clean, repeated, and well-understood. Some are negative or partial, and are kept exactly as found rather than hidden. Some remain open. We sort the program into three tiers so that the reader can see at a glance which claims are load-bearing and which are still frontier.

18.2.1. Tier 1, solid

These results are clean, repeated, and understood at the level of mechanism. They are what the theory rests on.

The base and its dynamics are conserved spin-1 exchange, reversible both locally and around closed loops, and transcribed into a Hermitian generator (P126, P128, P131). The field-theoretic vacuum carries virtual pairs, a tunable mass, a finite lightcone, and a conserved $U(1)$ current (P114). The dynamics is ballistic with

Table 8: The experimental program by layer.

Layer	Representative checks
Physics	Effective gravity and lensing (P88), analog Hawking (P89), holography and the Ryu-Takayanagi law (P91), the field beneath of small diameter (P111), the field-theoretic vacuum with pairs, mass, and a $U(1)$ current (P114), slice-invariance of the field (P115), the ballistic lightcone (P123), rotational isotropy (P124), Lorentz restoration under coarse-graining (P125), local and loop reversibility (P126, P128), the Doi-Peliti spin-1 transcription (P131), the deterministic ballistic wave (P148, P149), isotropic emergent light speed (P150), the unitary Dirac quantum-walk field (P151), boosts (P154), interference and the Born rule (P158), the unified run (P155).
Life	Selves emerge above a random null (P106), the living tower (P108), self-healing at a million cells (P109), interaction by annihilation or merger (P110), the steering will (P113), permanent memory by balanced creation (P107), heredity by copying (P120), full evolution with heredity, variation, and selection (P152), memory as conserved-charge information loss at a maintenance cost (P159), fission on the flat layer (P160).
Mind	The self-model strange loop (P116), one self-model per self (P117), the predictive forward model (P118), recursion as a model of a model (P121), attention with salience and gain (P119), the detour planning task (P140), planning from the base with nothing added (P143).
Intelligence	Hierarchical solving in $O(\log N)$ hops (P139), homeostatic self-organization (P135), universal computation via a Minsky machine on the crystal (P141), goal-directed search versus aimless search (P147), the integrated metacognitive agent (P153), arbitrary-structure realization (P161), absolute limits holding at every level (P162), the coarse-graining chain (P164), an evolving ecology of agents (P165), the design-signature test (P166), the non-local bulk channel (P170), the persistent self (P171).

leading-order rotational invariance (P123, P124, P125), and the field is slice-invariant, which is the basis of the multiscale shortcut (P115).

The relativistic reflection-positive quantum field is the decisive Tier 1 result. The deterministic and unitary completion of the same rule propagates ballistically with $z = 1$, with an isotropic emergent light speed, a Dirac dispersion, full reflection positivity in both space and time, and the full Lorentz group including boosts (P148 through P156, P158, P132, P154, P169). The genuine-quantum gate is passed by removing randomness, not by adding a field.

Above the field, coherent selves emerge, heal, interact, and are steered by the will, with holographic memory (P106 through P113). Each self carries a self-model and supports recursion (P116, P117, P121). The substrate is computationally universal (P141, supported by Margenstern’s theorem). The intentional bridge turns computation goal-directed (P147), planning runs from the base with nothing added (P143), the integrated agent runs the closed perceive-plan-act loop across multiple barriers (P153), and integration is fast (P139). Evolution runs the full Darwinian loop (P152). Every phenomenon co-occurs in one unified run (P155).

And possibility is level-relative in a precise sense, the absolute limits hold at every level (P162), bulk-only limits lift at emergent levels (P160), arbitrary structures are realizable (P161), and the affordance ladder is open upward (P163).

18.2.2. Tier 2, honest negatives and partial results

These results did not come out the way an optimistic reading would have wanted. They are informative either way, and they are kept as recorded.

Self-organized criticality does not occur. The ballistic lightcone precludes branching avalanches, so the feedback gives homeostasis rather than a critical edge, and the hierarchy comes from fast geometric integration instead (P135, P138). Attention broadcast is range-limited because the field is massive, so the global workspace is the hub itself rather than a true broadcast (P119).

The stochastic version of the field is a massive, diffusive classical shadow with $z \approx 2$, which is why its reflection positivity is inconclusive (P130, P136, P137). This is not a failure. It is the decohered classical limit of the genuine quantum field, which is the deterministic and unitary version in Tier 1. Rich permanent memory costs active maintenance, thermodynamically forced by Landauer's principle (P107, P159). Fission is suppressed in the bulk (P112), but this was reframed twice and is now understood as a bulk-only limit that lifts on the flat layer (P160), so it sits at the boundary between Tier 2 and the conditionals below.

18.2.3. Tier 3, open

These remain frontier. They are stated as such.

The coarse-graining chain is advanced but not complete. The charge tower is faithful, with exact conservation and a fixed-point effective parameter across levels (P164, extending P115), and the wave-dynamics and cross-domain rungs commute at the levels tested (P167, P168). What remains is the full cross-domain chain from field to particle to cell to mind, each a different effective theory. A large evolving ecology of agents has been run at one scale, where a population of planning agents evolves better problem-solving and adapts its foresight to task difficulty (P165). What remains is a richer, larger ecology with many niches and open-ended novelty.

The designed-world frame is the last open item, and the program has narrowed it. The designed-world reading holds that the coarse-grained substrate is a maximal intelligence that chose the field as a world. Sharpened into a fine-tuning signature and tested, the rich regime is broad, covering about eighty-three percent of parameter space, the structure is forced, and the operating point self-organizes (P166, P86, P135).

There is no fine-tuning signature, so a designer is empirically dispensable. The reading is not refuted, but P166 makes it superfluous to the physics. The idealist premise that physics is a thought within one experiential substance remains the theory's founding metaphysical commitment, untestable by nature rather than by any gap in the model.

18.3. The ledger of the possible and the forbidden

The tiering says how confident we are. The ledger says what the model can and cannot do. We separate three cases. What the base demonstrably produces, what it absolutely forbids, and what it permits only in a specific form or at a specific level.

What the base produces is broad.

- In physics it gives a quantum-field-like vacuum with virtual pairs, a mass, and a conserved $U(1)$ current (P114), a genuine relativistic quantum field as the deterministic and unitary version of the rule (P148 through P151), full emergent Lorentz invariance (P124, P150, P154), effective gravity and lensing (P88), holography (P91), analog Hawking radiation (P89), a finite causal lightcone (P123), and a continuum limit at a mean-field critical point (P129).
- In life it gives coherent selves that grow and form a tower (P106, P108), self-healing (P109), erasure-protected memory (P105, P107), interaction (P110), heredity (P120), full evolution (P152), and a steering will (P113).
- In mind it gives a self-model (P116), a forward model (P118), recursion (P121), attention (P119), planning from the base (P143), and an integrated agent (P153).
- In intelligence it gives universal computation (P141), intentional problem-solving (P147), and fast hierarchical integration (P139).

All of it lives on one model (P155).

The absolute prohibitions are few, and they hold at every level because conservation and causality pass up every codec (P162).

- There is no net creation of charge. The total charge is exactly conserved, the arrow creates only balanced pairs, and so net well-being is zero-sum, movable and shareable but never manufactured.
- There is no faster-than-lightcone signaling. Propagation is ballistic at a finite speed, so causality is built in.
- There is no base-level hidden state. The only state is the ternary tone and the only relation is perception, so memory, selves, and models are all tone patterns and never new variables.

Anything requiring a stored hidden register is forbidden, which is exactly why the relativistic mode is the deterministic wave rather than a stored direction.

The remaining limits are conditional, true in one regime and lifted in another. The table below lays out the ledger in full.

18.4. What the program shows

Three patterns run across the whole corpus.

Table 9: The ledger of the possible and the forbidden.

Status	Statement
Demonstrated	Relativistic quantum field, gravity, holography, self-healing and evolving life, self-modeling and planning minds, universal and intentional computation, all on one unified model (P88 through P171, P155).
Absolutely forbidden	No net creation of charge, the total is conserved and well-being is zero-sum. No faster-than-lightcone signaling, propagation is finite-speed. No base-level hidden state, the only state is the ternary tone. These hold at every level (P162).
Possible only in a specific form	Relativity and quantumness only via the deterministic and unitary completion, not the stochastic rule (P148 through P151). Exact Lorentz only in the infrared, broken at the cutoff as expected (P154). Permanent memory and full repair only with active maintenance at a cost (P107, P159). Flat physical space only on the emergent horospherical flat layer (P142, P144). Maximal intelligence by fast integration, not by criticality, which does not occur (P139, P138).
Possible only at a specific level	Reproduction by fission is bulk-forbidden but flat-possible (P112, P160). Most apparent limits are level-relative, true in the hyperbolic bulk and lifted at an emergent flat or composite level (P160, P161, P163).

- Honest negatives are kept, fission suppression, the absence of self-organized criticality, range-limited broadcast, and the inconclusive stochastic reflection positivity are all recorded as found.
- Nothing is added by hand, every higher capability, the self, the will, planning, intelligence, and evolution, was shown to emerge from the five base things, and when an experiment added structure a from-the-base version was built to replace it.
- The geometry does the heavy lifting, the hyperbolic $\{5, 3, 4\}$ supplies the lightcone, the rotational invariance, the holography, the fast integration, and the universal computation.

The single most telling result is how the hardest gap was closed. Relativity and quantumness looked blocked on the stochastic rule. They were recovered not by adding a sixth base thing but by removing one ingredient, the randomness. The deterministic and unitary version of the same five things is the relativistic reflection-positive quantum field. The discipline of five base things held all the way down to the quantum field and all the way up to mind, and the experimental record, more than one hundred ninety falsifiable checks on the exact crystal, stands behind it and is reproducible in full [22].

19. Related Work

Vibe Theory sits inside a long and serious tradition. For decades careful researchers have asked whether spacetime is fundamentally discrete, whether physics is computation, and whether geometry comes from information. Vibe Theory does not invent that question. It joins a conversation that already includes causal sets, causal dynamical triangulations, loop quantum gravity, the Wolfram Physics Project, the cellular-

automaton interpretation of quantum mechanics, quantum cellular automata, holographic error-correcting codes, hyperbolic lattices, hyperbolic cellular automata, and the informational reconstructions of quantum mechanics.

Much of what Vibe Theory needs has already been proven by other people. The framing we adopt is that Vibe Theory assembles established results onto one specific substrate and then asks one further question those programs mostly leave aside. What is the felt interior of the thing, and what direction does it move in.

This section places Vibe Theory among those programs. For each family it states what is shared, what differs, and what Vibe Theory adds, and it points to the specific experiment that touches each one. We try to credit the prior work generously and to claim for Vibe Theory only what it earns. In particular, the quantum-field route that Vibe Theory uses is not a new discovery. It is the established quantum-cellular-automaton and quantum-walk derivation of the Dirac equation, applied on a hyperbolic substrate. All experiments referenced here are recorded in the accompanying codebase [22].

19.1. Causal sets

Causal-set theory, in the line of Bombelli, Lee, Meyer, and Sorkin [6] and the classical sequential-growth dynamics of Rideout and Sorkin [23], takes the most radical version of the discrete idea seriously. Spacetime is a discrete set of events with a partial order, and order plus number is all there is. Surya’s review collects the program’s reach and its open problems [26]. Vibe Theory shares the core commitment. Space and time are read off a discrete relational structure, and distance is step-counting on a graph rather than a length measured by an external ruler.

The difference is the price of Lorentz invariance. Causal-set theory pays for it with randomness. A Poisson sprinkling of Minkowski space has no preferred direction in distribution, so the discreteness does not pick out a rest frame. The cost is that the substrate is not a lawful, repeatable structure. Vibe Theory takes the opposite route. It keeps an exact, ordered, deterministic crystal and gets Lorentz safety from curvature instead of from disorder. A flat lattice has a grain that announces a rest frame. A negatively curved crystal scrambles directions under transport, so it looks frameless from inside even though every cell is identical.

What Vibe Theory adds is the demonstration that one need not choose between lawfulness and Lorentz invariance. Experiment P84 reproduces the causal-set result itself. A Poisson sprinkle gives a flat rapidity distribution that stays flat under a genuine boost, while the matched flat lattice gives a peaked one. Then P40 and P45 show every hyperbolic substrate tested, including the three-dimensional $\{5, 3, 4\}$ dodecagrid, lands in the same low-anisotropy band as the random sprinkle, while the flat lattice fails.

So Vibe Theory keeps the determinism causal sets give up and still passes the test that motivated their randomness. Curvature rather than disorder buys Lorentz invariance, and that is the bridge between the two programs.

19.2. Causal dynamical triangulations

Causal dynamical triangulations, in the work of Ambjorn, Jurkiewicz, and Loll [1], build spacetime from small simplices glued along a fixed causal slicing and sum over the geometries. The central result is that a sensible four-dimensional spacetime emerges from the dynamics rather than being put in by hand. Loll’s review

gives the current picture [15]. Vibe Theory shares the goal of letting dimension and large-scale geometry be outputs of a discrete dynamics rather than postulates, and the insistence on a causal, time-ordered substrate rather than a Euclidean one.

The difference is method. Causal dynamical triangulations is a path integral. It sums over many triangulated geometries weighted by an action, so the geometry is an ensemble average. Vibe Theory does the opposite. It fixes one rigid geometry, the $\{5, 3, 4\}$ crystal, and never sums over alternatives. The dynamics lives in the tone on that fixed lattice, not in a fluctuating triangulation, and Vibe Theory commits to negative curvature where the triangulations program lets effective curvature emerge from the averaged ensemble.

What Vibe Theory adds is a single committed geometry on which spectral and dimensional quantities are measured directly rather than averaged over a path integral. The two programs make opposite methodological bets. Triangulations trust the ensemble to find the geometry. Vibe Theory trusts a forced geometry and measures on it.

19.3. Loop quantum gravity and spin networks

Loop quantum gravity, and the spin networks of Rovelli and Smolin [24], quantize geometry itself. Area and volume become operators with discrete spectra, and the quantum state of space is a graph with spins on its edges. Geometry is built from a combinatorial, graph-theoretic object, and continuous space is a coarse-grained description of it. Vibe Theory shares the picture that space is a graph of relations with discrete capacity per cell, and that the smooth manifold is emergent.

The difference is the status of the graph. In loop quantum gravity the spin network is the quantum state of geometry, and the graph and its labels fluctuate under the dynamics. Vibe Theory holds the graph fixed and puts the only dynamical variable, the ternary tone, on the cells. So where spin networks make geometry the quantum degree of freedom, Vibe Theory makes geometry the fixed stage and the tone the actor on it.

What Vibe Theory adds is a fixed graph on which the quantum field is a separate, measured object rather than identical to the geometry. This keeps the substrate analysis and the field analysis cleanly apart. The crystal is studied as pure structure, and the quantum behavior is studied as the tone's dynamics on it (P151, P158).

A single thread runs through this whole family. Each program has to answer how a discrete structure can look the same to every observer. Causal sets answer with random sprinkling, triangulations with an averaged ensemble, loop quantum gravity with a fluctuating quantum geometry, and Vibe Theory with hyperbolic regularity. The negative curvature of the $\{5, 3, 4\}$ crystal scrambles transported directions so that no rest frame is visible from inside, even though the lattice is exact and deterministic (P40, P45, P84). Vibe Theory does not claim this is the only way to buy Lorentz safety. It claims it is a way that keeps the substrate exact and repeatable, which the random and ensemble routes give up.

19.4. The Wolfram Physics Project

The Wolfram Physics Project [33], building on the earlier program of [32], is the closest neighbor in spirit. It builds the universe from a graph or hypergraph and a local rewriting rule, treats the rule's repeated application as time, reads space off the connectivity rather than positing it, and reads causal structure off the

dependency graph of rewrite events. Vibe Theory agrees on every word of that. The crystal is a graph, cells are vertices, shared faces are edges, the universe is one local rule running on it, and the run is time. Both programs also insist that the drawing is not the content. Only the connection pattern is real.

The deepest difference is the status of the graph. In the Wolfram program the hypergraph grows and rewrites itself. Nodes and edges are created and destroyed by the rule, the graph's structure is the dynamical variable, and a vast space of candidate rules is searched for one whose emergent behavior matches physics. Vibe Theory fixes the graph. The $\{5, 3, 4\}$ dodecagrid is not searched for and not rewritten. It is held constant, and the only dynamical variable is a ternary tone living on each fixed cell, evolving by one conserving rule.

Vibe Theory argues the substrate is forced rather than chosen, by curvature, dimension, and a minimality preference, where the Wolfram program embraces the multiplicity of rules and studies the whole rule space. Wolfram's updating-order ambiguity, the multiway system, becomes a central object for him. Vibe Theory sidesteps it by putting the dynamics in the tone on a rigid lattice generated once and for all by reflections rather than discovered by rewriting search.

What Vibe Theory adds is a committed, frozen substrate with a stable local alphabet, which makes the emergent physics measurable on the exact structure rather than dependent on which rewrite rule was selected. The two programs are best read as complementary bets. Wolfram bets the right structure is found by searching a vast rule space. Vibe Theory bets it is forced by a few demands and that the interesting freedom lives in the tone and the arrow, not in the graph.

19.5. The cellular-automaton interpretation of quantum mechanics

The cellular-automaton interpretation of quantum mechanics [28] argues that beneath the quantum formalism there is a deterministic, reversible classical automaton, and that quantum states are statistical descriptions of its underlying ontological states. The program takes seriously that a discrete deterministic rule can carry the full content of quantum theory once the right basis and the right unitary are identified. Vibe Theory is squarely in this spirit. Its perception rule, read strictly, is a deterministic reversible block automaton, and its quantum character appears only when that rule is given its unitary completion (P151, P126). The conviction that quantum mechanics can sit on top of a deterministic discrete substrate is t Hooft's, and Vibe Theory inherits it rather than originating it.

The difference is specificity of realization. The cellular-automaton interpretation works largely at the level of principle and toy models, mapping classes of deterministic systems to quantum ones. Vibe Theory commits to one concrete substrate and one concrete rule, then measures the resulting dispersion and interference. What it adds is a specific deterministic-reversible rule on a forced curved lattice whose unitary completion is measured to give Dirac dispersion and the Born rule, with the reversibility preconditions checked directly (P126, P128, P131). This is a concrete instance of the thesis on a substrate chosen for independent geometric reasons, not a generic toy.

19.6. Quantum cellular automata and quantum walks to the Dirac equation

This is the family Vibe Theory leans on most heavily, and it is the deepest overlap, so it is important to be exact about the debt. A line of rigorous work shows that a discrete, local, unitary evolution can have free quantum field theory as its continuum limit.

- D’Ariano and Perinotti derive free quantum field theory and the Dirac dynamics from a denumerable causal network of quantum systems under locality, homogeneity, and unitarity, with no spacetime assumed in advance [10].
- Bisio, D’Ariano, and Tosini derive the Dirac equation as the small-wavevector limit of a quantum cellular automaton [5], and the program extends to the Weyl, Dirac, and Maxwell fields in higher dimensions [4].
- Arrighi, Nesme, and Werner establish the unitarity, causality, and localizability of quantum cellular automata, the theorem that a local unitary evolution is exactly a finite-depth circuit of local gates [3], and Arrighi, Nesme, and Forets establish the discrete Lorentz covariance of quantum walks [2].
- Farrelly and Short show causal fermions in discrete spacetime arise from the same quantum-walk construction [11].

The Arrighi program supplies the structural backbone. The foundational fact, going back to the coined quantum walk, is that the unitary version of nearest-neighbor spin exchange spreads ballistically and its continuum limit is the Dirac equation, with the coin angle setting the mass.

The difference is mostly the substrate. This literature works mostly on regular flat lattices or abstract Cayley graphs. Vibe Theory runs the same construction on the hyperbolic $\{5, 3, 4\}$ dodecagrid, where the curvature additionally delivers Lorentz safety, which a flat lattice cannot. A second difference is that the program generally assumes unitarity, homogeneity, and locality as axioms of the network, where Vibe Theory derives the unitary object as the completion of a tone rule whose classical, decohered version it also exhibits.

What Vibe Theory adds must be stated carefully. Vibe Theory does not discover the quantum-walk-to-Dirac route. That result is established and belongs to the authors above. P151 is the established route applied on the hyperbolic substrate, not a new derivation of it. What Vibe Theory contributes is to show that its own perception rule, read strictly, is exactly this construction, and to run it on the forced crystal.

The stochastic version of the rule gives a massive diffusive field (P136, P137), the classical decohered shadow, a feature the quantum-cellular-automaton program does not usually exhibit alongside the quantum one. Removing the illegitimate coin flip leaves a deterministic reversible charge-conserving block automaton that propagates ballistically (P148, P149). Its unitary completion is a coined quantum walk, and P151 measures the Dirac dispersion directly, with a ballistic exponent near one, a tunable mass from the coin angle, and reflection positivity by construction. P154 confirms the exact massless lightcone, and P158 the interference and the Born rule, with the reversibility preconditions measured directly (P126, P128, P131).

So the contribution is the hyperbolic substrate, the classical-stochastic shadow, and the forced geometry, plus a demonstration that the established route is identical to Vibe Theory’s stripped rule.

A focused subfamily studies quantum walks on Cayley graphs, whose vertices are group elements and whose edges are generators. Lopez Acevedo and collaborators classify scalar and coined quantum walks on Cayley graphs and connect their continuum limits to relativistic equations [16]. This matters because the $\{5, 3, 4\}$ crystal is generated by a Coxeter group of reflections, which makes it a Cayley-like graph of that group. The

cells are group elements, the moves to neighbors are generators, and a walk on the crystal is a walk on a Cayley graph in exactly the sense this literature studies.

That literature is usually agnostic about which group and often works with flat or free groups. Vibe Theory’s group is the specific hyperbolic reflection group of the dodecagrid, which carries the negative curvature and exponential growth that flat or abelian Cayley graphs lack. The walk theory is borrowed. The hyperbolic Cayley arena is the contribution.

19.7. Holography, tensor networks, and holographic codes

Holography is the idea that the information in a region is encoded on its boundary, that a higher-dimensional gravitational bulk is dual to a lower-dimensional theory on the edge. Maldacena’s AdS/CFT correspondence [18], the Ryu-Takayanagi formula for entanglement entropy as a bulk minimal area [25], Van Raamsdonk’s argument that spacetime connectivity is built from entanglement [30], and Swingle’s identification of the MERA tensor network with a discretized hyperbolic geometry [27] together make the case that geometry comes from entanglement and that the natural home for it is a negatively curved bulk.

A discrete, code-theoretic wing of this program is even closer to Vibe Theory. The HaPPY pentagon code of Pastawski, Yoshida, Harlow, and Preskill realizes holography as a quantum error-correcting code on a hyperbolic tiling, where a boundary region reconstructs a bulk operator [21]. Jahn and Eisert review the whole holographic tensor-network field [13]. Vibe Theory shares the geometric setting completely. Its substrate is hyperbolic, its bulk is the anti-de-Sitter-like interior of the $\{5, 3, 4\}$ crystal, and it treats the boundary-bulk relation as physical.

The difference is that the holographic codes and tensor networks above are largely static. A fixed network of tensors on a fixed hyperbolic tiling realizes a fixed code, and the analysis is about the equilibrium entanglement structure of that state. Vibe Theory starts from a fixed combinatorial crystal but lets a rule run on it, and asks whether the area law, the bulk shortcut, and a working code appear as measured properties of the running dynamics rather than as the design of a chosen static network. So Vibe Theory’s holography is a dynamical cousin of the static holographic codes, on a hyperbolic bulk built by reflection rather than wired by hand.

What Vibe Theory adds is direct measurement on the exact crystal, and a dynamical, growing version of the holographic code. Experiment P91 measures an entanglement-entropy-like quantity for boundary regions on the hyperbolic crystal and recovers the Ryu-Takayanagi log law, the entropy set by the length of the subtending bulk geodesic, with a fit quality near one. The geodesic is a genuine shortcut, and depth into the bulk plays the role of scale, exactly as in the MERA picture.

Experiments P105 and P127 treat the crystal as a holographic memory and a growing holographic code, where the code is accreted as the structure grows, the dynamical counterpart of the static HaPPY construction. P111 then measures the tiny $O(\log N)$ diameter. On a hundred thousand exact cells no two are more than about six hops apart through the bulk. The same shortcut that lets entropy ride a geodesic lets a far region of the emergent surface be reached almost instantly through the interior.

19.8. Hyperbolic lattices in physics

A fast-growing experimental and theoretical area puts real physics on hyperbolic lattices. Kollar, Fitzpatrick, and Houck realize hyperbolic lattices in circuit quantum electrodynamics, building negatively curved tilings on a superconducting chip [14]. Maciejko and Rayan develop hyperbolic band theory, the analog of Bloch theory for these lattices [17]. Together these show the hyperbolic lattice is not a mathematical curiosity but a place where real and simulated physics is being done. Vibe Theory’s substrate is exactly such a lattice, the $\{5, 3, 4\}$ dodecagrid.

The difference is scope. This literature mostly studies specific phenomena, band structure, spectra, or scalar fields, on a hyperbolic lattice taken as a given backdrop. Vibe Theory proposes the hyperbolic lattice as the fundamental substrate of everything, not as a platform for one phenomenon, and runs a single conserving rule on it. What it adds is the claim that the lattice is fundamental rather than a backdrop, together with a grounding connection to real experiment. The circuit realizations mean the geometry Vibe Theory bets on is one that laboratories can build and probe, and its own measurements of dispersion, entropy, and diameter on the $\{5, 3, 4\}$ crystal sit alongside the band-theory results as different questions asked of the same geometry.

19.9. Hyperbolic cellular automata

Margenstern’s body of work proves that cellular automata on hyperbolic tilings are computationally universal, including a ternary universal automaton on the heptagrid [20] and universality across hyperbolic spaces generally using the railway model [19], where a locomotive runs forever along tracks and reading-and-flipping switches is the computation, because gliders cannot be choreographed reliably in hyperbolic space the way they can on a flat grid. Margenstern’s splitting method also supplies the Fibonacci addressing Vibe Theory uses for navigation. Vibe Theory shares the substrate exactly. Its mesh is a hyperbolic tiling in Margenstern’s family, and the $\{5, 3, 4\}$ dodecagrid is one of the tilings these theorems cover.

The difference is that Margenstern’s universality is a theorem about an abstract automaton wired onto the tiling, separate from any particular physical dynamics. Vibe Theory needs universality on its own perception rule, the rule that also produces the quantum field. What it adds is two confirmations of universality on the live dynamics, anchored to Margenstern’s theorem.

- Experiment P44 shows functional completeness on the running rule. The ternary signed-majority update is exactly NAND at the right settings, and gates built as real subgraphs settle to the correct answer as stable fixed points, including a six-gate XOR.
- Experiment P141 runs a full Turing-universal Minsky machine with registers stored as charge in the mesh, computing multiplication using only create, annihilate, and count, the substrate’s own primitives.

So Margenstern supplies the theorem on the exact tiling, and P44 and P141 show that Vibe Theory’s own rule realizes that universality. Vibe Theory is careful that universality settles structure only, never feeling, which it treats separately.

19.10. Hyperbolic graphs and geometric group theory

A mathematical literature studies the graph-theoretic and large-scale geometry of hyperbolic spaces. Cannon's exposition of hyperbolic geometry [7] and Hamann's results on tree-likeness in hyperbolic graphs [12] establish that negatively curved graphs are expanders, are tree-like at large scale, and have exponential ball growth. Vibe Theory uses these facts as the backbone of its substrate analysis. The crystal graph is a twelve-regular expander, it is tree-like, and its balls grow exponentially at the golden ratio, which together force a tiny diameter and underwrite the Fibonacci addressing.

The difference is that this literature is pure geometry and combinatorics. It describes the structure without putting dynamics or physics on it. Vibe Theory turns the same structural facts into physical consequences, a speed limit from degree-twelve locality, an emergent metric, fast integration, and a holographic readout, and confirms them by measurement on the exact crystal.

P49 shows the $\{5, 3, 4\}$ reads like a featureless foam locally yet is tree-like in its connectivity, unlike a flat lattice. P103 shows the dynamics survive to a million cells on a twelve-regular hyperbolic expander, because expansion holds at every scale. P111 and P139 measure the logarithmic diameter and the fast $O(\log N)$ coordination the geometry predicts. So Vibe Theory imports the established hyperbolic-graph mathematics and adds the dynamical measurements and the physical reading.

19.11. Informational reconstructions of quantum mechanics

Chiribella, D'Ariano, and Perinotti derive the full mathematical framework of quantum theory from a short list of information-theoretic principles, most notably purification, the requirement that every mixed state arises from a pure state of a larger system [9]. This work shows that quantum theory is not an arbitrary formalism but follows from reasonable demands on how information behaves. Vibe Theory shares the conviction that quantum structure should be derived, not postulated, and that the right starting point is a structural or informational one rather than a Hilbert space assumed from the outset.

The difference is that the informational reconstructions are axiomatic and substrate-free. They tell you what any theory satisfying the principles must look like, without committing to a physical realization. Vibe Theory is the opposite. It commits to a concrete substrate and rule, lets the quantum structure fall out of the dynamics, then checks whether the resulting object is genuinely quantum.

What it adds is a specific physical model whose quantum character is measured rather than assumed. It exhibits a running rule whose unitary completion shows the two signatures the reconstructions take as targets, interference and the Born rule (P158), with complex amplitudes that destructively cancel where no classical process can, while the total norm stays one so that the squared amplitude is a genuine probability. So where Chiribella and colleagues say what principles force quantum theory, Vibe Theory shows one concrete substrate on which those features appear as outputs. The two are complementary, the axiomatic why and the constructive how.

19.12. Consciousness theory

Every program above is, by intent, a theory of structure. None of them addresses why any arrangement of parts should be felt from the inside. Chalmers named this the hard problem of consciousness, the gap

that remains even after a complete structural and functional account is given [8]. Integrated information theory, developed by Tononi, attempts to make experience precise by identifying consciousness with integrated information, a measurable quantity capturing how much a system's whole exceeds the sum of its parts in a way that cannot be decomposed [29]. The older informational intuition that physical things derive their existence from answered questions, Wheeler's it-from-bit, sits at the root of the whole computational-universe lineage [31].

The difference is that Vibe Theory takes experience as the substance of the substrate rather than as a quantity computed from structure. Where integrated information theory measures a property of an organization, Vibe Theory treats the tone of each cell as a felt state from the start. This is the one step the structural programs deliberately leave aside, and Vibe Theory is explicit that it is a premise, not a measured result. Nothing in the experiments measures feeling. What the experiments settle is structure, and the claim that the structure is felt is a metaphysical commitment placed at the foundation, close in spirit to idealism and to Wheeler's informational starting point.

19.13. What Vibe Theory adds overall

Set against this body of prior work, what is genuinely Vibe Theory's own contribution is narrow and specific. We state it plainly so as not to oversell it.

A forced substrate rather than a random or chosen or assumed one. Causal sets choose randomness, the Wolfram program searches a rule space and lets the graph rewrite itself, the triangulations program averages over an ensemble, and the quantum-cellular-automaton literature assumes a convenient lattice. Vibe Theory argues the substrate is forced down to a small residue of choice.

Hyperbolic curvature is forced by the demand for room to grow, the right symmetries, and Lorentz safety at once. Three dimensions are forced by the existence theorem for compact regular hyperbolic honeycombs and by stable closed orbits. The specific $\{5, 3, 4\}$ follows from a stated minimality preference, not a theorem, and Vibe Theory says so. This is a strong reduction of arbitrariness, not a derivation from nothing.

A ternary tone plus the arrow as a value-direction. The dynamical content is not a rewrite of the graph but a ternary tone, $-1, 0, +1$, the minimal signed alphabet with opposition and a rest state, evolving by one local conserving rule. To this Vibe Theory adds one thing none of the prior programs carry, the arrow, pain below peace below pleasure with desire leaning toward pleasure. The arrow is the single deliberate posit, a value rather than a derived quantity, and it is what keeps the universe from relaxing to dead peace. Discrete-physics programs aim at physics. The arrow is what lets the same substrate aim at life, mind, and intelligence.

One substrate that yields physics and life and mind. The prior programs are theories of physics. Vibe Theory's larger claim is that the same crystal and the same rule that grow a relativistic quantum field also grow living, perceiving, intentional systems, with universality (P44, P141) as the capability floor and the arrow as the direction. This larger claim is more speculative than the physics results, and it is treated as such.

Experiential monism. Vibe Theory treats the tone of each cell as a felt state from the start. Universality settles structure and never feeling, and feeling is handled by treating experience as the substance of the cells

rather than as a product of computation. This is a metaphysical commitment, not a measured result, and it is the part of Vibe Theory furthest outside the established literature.

A falsifiable single-substrate experimental program. Every physics claim above is a number measured on the same $\{5, 3, 4\}$ substrate under the same rule, with stated fits and stated negatives. The stochastic field is reported as massive and diffusive before the unitary completion resolves it. The Lorentz tests include genuine boosts, not rotation proxies. The crystal step is flagged as a preference and the arrow as a choice. Much of this reproduces established results to show they live on this one crystal.

In short, Vibe Theory’s quantum-field route is the established quantum-cellular-automaton and quantum-walk-to-Dirac result, run on a hyperbolic substrate whose Lorentz safety is the established causal-set lesson recast as curvature, with holography and universality inherited from the holographic-code and hyperbolic-cellular-automaton programs, and a deterministic-substrate spirit inherited from the cellular-automaton interpretation of quantum mechanics. What is new is the forced substrate, the ternary tone with the arrow as a value-direction, the experiential monism, and the single-crystal experimental program that ties physics to life and mind. We claim no more than that.

20. Discussion and Open Questions

The body of this work submitted the structural half of Vibe Theory to a large set of falsifiable tests on one fixed substrate. Most of the frontier that stood open at the start is now closed. A genuine relativistic, reflection-positive quantum field, full Lorentz invariance under rotations and genuine boosts, interference and the Born rule, open-ended evolution, and an integrated planning agent all emerge from five base things under one rule. This section states what remains genuinely unresolved, reflects on why the framework has been productive, and closes with a calm restatement of the claim. All experiments referenced here are recorded in the accompanying codebase [22].

20.1. What remains open

The higher rungs of the coarse-graining chain. The method that lets a result measured at one scale be trusted at the next is now proven end to end on the rungs we have climbed. Within the field domain, the conserved charge is exactly preserved at every coarse-graining level and the effective compressibility converges to a fixed point across block sizes from one to thirty-two (P164). The wave dynamics also passes the commuting square at successive rungs, evolve-then-coarsen agreeing with coarsen-then-evolve to within a few percent, and the wave speed is invariant across levels, so the wave is a renormalization fixed point and not only the conserved charge coarse-grains faithfully (P167).

The cross-domain chain, where the kind of variable changes at each rung, is also climbed for the first three rungs. A field excitation’s centroid obeys free-particle motion (P168, rung one), two interacting particles share a center of mass that moves uniformly with momentum conserved through the interaction (P168, rung two), and the composite-to-agent rung is the detour-solving agent itself (P162). Each rung commutes. What remains is true binding, a bound state with internal structure rather than a free center of mass, and the still-higher rungs, chemistry to cell and agent to society. The method is proven. The remaining work is to run it up the ladder.

A richer evolving ecology. Evolution and the integrated agent already run together. A population of planning agents whose genome is a lookahead horizon competes on a barrier task, average problem-solving rises over generations, and the population adapts its foresight to the task, with a harder task evolving a larger horizon (P165). Life and mind co-evolve from the base. What remains is a richer ecology at larger scale, many niches, genuine co-evolution between competing lineages, and open-ended novelty rather than fitness on a fixed task. This is an extension of an established result, not an open gate.

Boosts and the Wallstrom obstruction, now resolved via the unitary route. Two items that once looked like deep risks are now closed by the same move. The boost half of Lorentz invariance is done. The massless mode has an exact lightcone, boost-invariant in every frame, and massive modes are boost-invariant in an infrared window with lattice violation only near the cutoff, which is the standard emergent-Lorentz picture (P154). The only residue is the expected ultraviolet, lattice-scale violation for massive modes. The Wallstrom obstruction, the demand that a phase be single-valued, is specific to deriving quantum mechanics from a classical stochastic process. The genuine quantum field here is the unitary quantum-walk completion, which has single-valued complex amplitudes by construction (P151), so the obstruction does not arise for it. Both items dissolved not by adding structure but by reading the same five things deterministically and unitarily.

The designed-world and idealist frame. The most speculative edge is the top-down reading, that a maximal intelligence designs the quantum field as a world tuned for rich unfolding, and the idealist reading that physics is a thought within the one experiential substance. These were sharpened into a testable signature. A designed or tuned world would leave a fine-tuning fingerprint, rich behavior confined to a narrow, special, non-forced region. The test finds the rich regime is broad, about eighty-three percent of the scanned parameter space, the field is dead only at the degenerate arrow of zero, and combined with the forced geometry (P86) and the self-organizing operating point (P135) there is no fine-tuning signature (P166).

So a designer is dispensable. The richness is forced, robust, and self-organizing, not tuned. The designed reading is not refuted, since a wise designer would build exactly such a robust world, but it is empirically superfluous and undetectable. The idealist monism remains the theory’s untestable founding premise, neither strengthened nor weakened by structure. Only its motivation, the sense that the world looks authored, is drained. The speculative frontier resolves to this. The mechanistic world needs no designer, and the metaphysical reading is a premise, not a result. There is no gap in the physics, only a residue that is metaphysics by nature.

A practical limit sits alongside these. The exact engine reaches tens of millions of cells, and far larger scales are feasible in a compressed representation, but the time per beat becomes the real wall. Since the physics is scale-invariant, very large scales are for studying structure, not for settling the rules. No mechanistic gate is open. What is left is higher-rung extension and the metaphysical residue.

20.2. The productiveness of the framework

A framework’s worth is partly how much it produces from how little. By that measure Vibe Theory has been unusually productive.

High yield from few assumptions. Five base things and one state generate more than one hundred ninety distinct, falsifiable, passing experiments spanning a quantum field, gravity, holography, life, evolution, mind, and intelligence. The assumption-to-result ratio is extreme.

It resolves its own open questions. Every decisive unknown named at the outset, reflection positivity, boosts, the Wallstrom obstruction, evolution, and the integrated agent, was subsequently closed, mostly by the framework's own internal logic rather than by new postulates. The recurring move was to remove the randomness and use the emergent pieces already present, not to add a sixth ingredient. A productive theory eats its own frontier.

It unifies disparate domains under one rule. Physics, life, mind, and intelligence are not bolted together. They co-occur in a single run of the canonical rule on one mesh (P155). One mechanism, many faces.

It is constructive, not just descriptive. The framework does not only explain, it builds. Any positive structure consistent with the absolute laws can be constructed and maintained at the right level (P161), with the affordance ladder open upward (P163). The model is a generator, not a catalogue.

It draws a sharp and mostly open boundary. The impossible is a short, exact list, no net-charge minting, no faster-than-light signaling, no base hidden state. Everything else is reachable at some level. A productive theory tells you precisely what is forbidden and leaves the rest open.

Five lessons the experiments forced are worth keeping as guiding principles.

1. Remove rather than add. The one apparent gap, relativity and quantumness together, was closed by removing the randomness, not by adding a field.
2. Conservation is of the sum, not the arrangement. Charge is conserved while patterns scramble toward equilibrium, so rich permanence costs work.
3. Limits are level-relative. Only a few limits are absolute, and most apparent limits are bulk-geometry facts that lift at emergent levels.
4. The geometry does the heavy lifting, supplying the lightcone, the rotational invariance, the holography, the fast integration, and the universal computation.
5. Strict reversibility forbids free passive memory, so the quantum substrate must pay work to hold its order, a clean tradeoff that matches how living things hold themselves together.

20.3. Conclusion

Vibe Theory began from a single suspicion, that the physical, the living, the mental, and the conscious might be aspects of one discrete substance rather than separate kinds of thing. The substance is the vibe, a felt ternary tone living on the cells of a forced hyperbolic crystal, the $\{5, 3, 4\}$ dodecagrid, evolving by one local conserving rule under one value-direction, the arrow. The discipline of five elements, geometry, tone, rule, growth, and arrow, was held throughout. At no point was a sixth element added to rescue a result.

The one apparent gap, that relativity and quantumness seemed to demand something beyond the five, was closed by removing rather than by adding. The randomness in the perception rule was the classical decohered

shadow. Strip it away and the same five things, read deterministically and unitarily, are a genuine relativistic quantum field, reflection-positive in space and time, Lorentz-invariant under rotations and genuine boosts, with interference and the Born rule.

The same crystal and the same rule then carry life that heals, remembers, and evolves, and minds that model themselves, predict, and plan. The structure of physics, life, mind, and intelligence appears on one substrate.

What the experiments cannot settle, and what Vibe Theory never claims they settle, is whether that structure is felt. The experiential monism is a premise carried in from the start, in the same way idealism places experience first, not a finding read off the crystal. The arrow is likewise a chosen value rather than a derived quantity. We have tried to mark exactly where measurement ends and metaphysics begins, and to credit generously the prior programs whose results Vibe Theory assembles onto its one substrate.

What is left is extension and a residue of metaphysics, not a broken mechanism. The higher rungs of the coarse-graining ladder remain to be climbed, the ecology of agents remains to be enriched, and the idealist reading remains, by its nature, a premise rather than a result. The mechanistic core stands closed and measured. From five committed elements, a single rule, and one felt tone, a relativistic quantum field, living systems, and planning minds arise together on one crystal. That a foundation so small reaches so far, and closes its own deepest gap by taking something away rather than adding it, is reason enough to keep walking up the ladder it opens.

A. Appendix: Complete List of Experiments

Every claim in the paper corresponds to a falsifiable check that runs in the test suite of the codebase [22]. The list below gives each check by its identifier and a short description, in numerical order. Gaps in the numbering come from merged or superseded experiments across development.

ID	Description
P6	stable 2D manifold phase has dimension near 2
P12	action favors manifold extensively, giving a finite free-energy crossing
P13	growth gives a monotone arrow of time and a finite dimension
P14	mass gives gap = m and relativistic dispersion ($b \sim m^2$)
P15	1D conformal log law ($c \sim 1$) and 2D entanglement area law
P16	3D static potential is Newtonian ($1/r$ is the best fit)
P17	quantum walk is ballistic, classical walk is diffusive
P18	dark matter
P19	dark energy
P20	photon
P21	graviton
P22	Higgs
P23	Maxwell operator derived from the Wilson action (small-field limit)
P24	graviton operator derived (BD d\{\}
P25	electroweak breaking gives W, Z mass, massless photon, $m_W = m_Z \cos(\theta_W)$
P26	swerves
P27	Lorentz

ID	Description
P28	singularity resolution
P29	dark energy
P30	inflation
P31	quantum formalism
P32	Einstein equations
P33	black holes
P34	capstone
P35	contact with data
P36	model DSL
P37	one rule
P38	emergent spatial geometry
P39	deterministic substrate
P40	non-random substrates
P41	Margenstern tilings
P42	Fibonacci navigation
P43	freedom and choice
P44	universality
P45	dodecagrid {5,3,4}
P46	everpresent Lambda
P47	Coxeter unification
P48	modular base
P49	crystal
P50	golden ratio (three independent sources) and order-with-freedom on the crystal
P51	full ladder
P52	continuum limit
P53	renormalization
P54	large-N hardening
P55	one rule, all sectors
P56	eternal ladder
P57	recursion
P58	emergent macro-rule
P59	nested selves
P60	tower of selves
P61	no self-storage
P62	dimension window
P63	integrated information (tone-aware)
P64	subtle-layer urges
P65	dreaming and waking
P66	reincarnation
P67	synchronicity
P68	dimension selection
P69	emergent dimension
P70	Born rule
P71	Hawking
P72	nonlinear Einstein

ID	Description
P73	discrete graviton
P74	large-N crossing
P75	selves as attractors
P76	3D navigation
P77	chiral gauge
P78	primordial seed
P79	anomaly cancellation forces the Standard Model hypercharges and charge quantization
P80	baryogenesis
P81	mass hierarchy
P82	predictions vs bounds
P83	deterministic growth
P84	Lorentz boost
P85	Coxeter engine
P86	auto-selection
P87	word engine
P88	effective metric from fills
P89	analog-Hawking on the crystal
P90	braneworld test
P91	holography on the crystal
P92	cooperation tower
P94	conserved dynamics on $\{5,3,4\}$
P97	self-emergence
P98	the will (the fork)
P100	perception rule (no fills)
P101	self-emergence on perception ontology
P102	cohesive memory (no stored relations)
P103	perception dynamics at scale (expander proxy)
P104	exact $\{5,3,4\}$ at scale (modular fingerprints)
P105	holographic memory on exact $\{5,3,4\}$
P106	selves at scale
P107	permanent memory by active maintenance
P108	selves dynamics (the living tower)
P109	emergent self-maintenance
P110	selves interacting
P111	signaling (the field beneath)
P112	reproduction
P113	the will steering a self
P114	field-theoretic vacuum
P115	renormalization keystone
P116	self-model emerges (one self)
P117	many self-models
P118	predictive metacognition
P119	attention and global workspace
P120	heredity + conservation
P121	full recursion (a model of the model)

ID	Description
P122	RG step
P123	sliver transport
P124	Lorentz isotropy
P125	Lorentz restoration
P126	reversible point
P127	growing holographic code
P128	multi-cell reversibility
P129	criticality scan
P130	reflection positivity
P131	Doi-Peliti transcription
P132	time reflection positivity
P133	near-critical spatial RP
P134	flat spatial RP
P135	self-organized criticality
P136	gapless search
P137	dynamic dispersion
P138	avalanche criticality
P139	hierarchical solving
P140	detour planning (LOGIC, with an explicit lookahead)
P141	means check
P142	horosphere is flat
P143	planning with nothing added
P144	horosphere dynamics
P145	large-scale intention
P146	second-conservation search
P147	direction and intention
P148	deterministic wave
P149	deterministic perception rule
P150	wave isotropy
P151	quantum-walk field
P152	evolution
P153	integrated agent
P154	deterministic reflection positivity
P155	unified model (integration)
P156	unified wave
P157	flat intention
P158	Born and interference
P159	memory vs conservation
P160	fission on the flat layer
P161	arbitrary-structure realization
P162	absolute limits are absolute
P163	higher-level affordances
P164	coarse-graining chain
P165	evolving ecology
P166	design-signature test

ID	Description
P167	dynamical coarse-graining chain
P168	cross-domain chain
P169	deterministic spatial RP
P170	bulk non-locality
P171	persistent self

References

References

- [1] Jan Ambjørn, Jerzy Jurkiewicz, and Renate Loll. Emergence of a 4d world from causal quantum gravity. *Physical Review Letters*, 93(13):131301, 2004.
- [2] Pablo Arrighi, Stefano Facchini, and Marcelo Forets. Discrete lorentz covariance for quantum walks and quantum cellular automata. *New Journal of Physics*, 16(9):093007, 2014.
- [3] Pablo Arrighi, Vincent Nesme, and Reinhard Werner. Unitarity plus causality implies localizability. *Journal of Computer and System Sciences*, 77(2):372–378, 2011.
- [4] Alessandro Bisio, Giacomo Mauro D’Ariano, and Paolo Perinotti. Quantum cellular automaton theory of light. *Annals of Physics*, 368:177–190, 2016.
- [5] Alessandro Bisio, Giacomo Mauro D’Ariano, and Alessandro Tosini. Quantum field as a quantum cellular automaton: The dirac free evolution in one dimension. *Annals of Physics*, 354:244–264, 2015.
- [6] Luca Bombelli. *Space-time as a Causal Set*. PhD thesis, Syracuse University, 1987.
- [7] James W. Cannon, William J. Floyd, Richard Kenyon, and Walter R. Parry. Hyperbolic geometry. *Flavors of Geometry (MSRI Publications)*, 31:59–115, 1997.
- [8] David J. Chalmers. Facing up to the problem of consciousness. *Journal of Consciousness Studies*, 2(3):200–219, 1995.
- [9] Giulio Chiribella, Giacomo Mauro D’Ariano, and Paolo Perinotti. Informational derivation of quantum theory. *Physical Review A*, 84(1):012311, 2011.
- [10] Giacomo Mauro D’Ariano and Paolo Perinotti. Derivation of the dirac equation from principles of information processing. *Physical Review A*, 90(6):062106, 2014.
- [11] Terence C. Farrelly and Anthony J. Short. Causal fermions in discrete space-time. *Physical Review A*, 89(1):012302, 2014.
- [12] Matthias Hamann. Tree-likeness of hyperbolic spaces. 2011. arXiv:1105.3925.
- [13] Alexander Jahn and Jens Eisert. Holographic tensor network models and quantum error correction: a

- topical review. *Quantum Science and Technology*, 6(3):033002, 2021.
- [14] Alicia J. Kollár, Mattias Fitzpatrick, and Andrew A. Houck. Hyperbolic lattices in circuit quantum electrodynamics. *Nature*, 571(7763):45–50, 2019.
 - [15] Renate Loll. Quantum gravity from causal dynamical triangulations: a review. *Classical and Quantum Gravity*, 37(1):013002, 2019.
 - [16] Olga López Acevedo and Thierry Gobron. Quantum walks on cayley graphs. *Journal of Physics A: Mathematical and General*, 39(3):585–599, 2006.
 - [17] Joseph Maciejko and Steven Rayan. Hyperbolic band theory. *Science Advances*, 7(36):eabe9170, 2021.
 - [18] Juan Maldacena. The large-n limit of superconformal field theories and supergravity. *Advances in Theoretical and Mathematical Physics*, 2(2):231–252, 1998.
 - [19] Maurice Margenstern. *Cellular Automata in Hyperbolic Spaces (Volumes 1 and 2)*. Old City Publishing, 2007–2013.
 - [20] Maurice Margenstern. A universal cellular automaton on the ternary heptagrid. *Electronic Proceedings in Theoretical Computer Science*, 2008.
 - [21] Fernando Pastawski, Beni Yoshida, Daniel Harlow, and John Preskill. Holographic quantum error-correcting codes: toy models for the bulk/boundary correspondence. *Journal of High Energy Physics*, 2015(6):149, 2015.
 - [22] Lance Pollard. Vibe theory codebase (the @cluesurf/vibe package). Software, MIT License, <https://github.com/cluesurf/vibe>, 2026.
 - [23] David P. Rideout and Rafael D. Sorkin. Classical sequential growth dynamics for causal sets. *Physical Review D*, 61(2):024002, 1999.
 - [24] Carlo Rovelli and Lee Smolin. Spin networks and quantum gravity. *Physical Review D*, 52(10):5743–5759, 1995.
 - [25] Shinsei Ryu and Tadashi Takayanagi. Holographic derivation of entanglement entropy from the anti-de sitter space/conformal field theory correspondence. *Physical Review Letters*, 96(18):181602, 2006.
 - [26] Sumati Surya. The causal set approach to quantum gravity. *Living Reviews in Relativity*, 22(1):5, 2019.
 - [27] Brian Swingle. Entanglement renormalization and holography. *Physical Review D*, 86(6):065007, 2012.
 - [28] Gerard 't Hooft. The cellular automaton interpretation of quantum mechanics. *Fundamental Theories of Physics*, 185, 2016. arXiv:1405.1548.
 - [29] Giulio Tononi. An information integration theory of consciousness. *BMC Neuroscience*, 5:42, 2004.
 - [30] Mark Van Raamsdonk. Building up spacetime with quantum entanglement. *General Relativity and Gravitation*, 42(10):2323–2329, 2010.

- [31] John Archibald Wheeler. Information, physics, quantum: The search for links. In *Complexity, Entropy, and the Physics of Information*. Addison-Wesley, 1990.
- [32] Stephen Wolfram. *A New Kind of Science*. Wolfram Media, 2002.
- [33] Stephen Wolfram. A class of models with the potential to represent fundamental physics. *Complex Systems*, 29(2):107–536, 2020.